

AB-100 Practice Questions - Deploy

Agentic AI Business Solutions Architect | Deploy domain

Generated 2026-05-05

46 questions, with answers and explanations

Deploy AI-powered business solutions (40-45% of the AB-100 exam).

Question types

Drag and drop: 6

Hotspot: 10

Multi-select: 7

Multiple choice: 23

Topic coverage

Topic 1: 7

Topic 2: 1

Topic 3: 38

Source: [practicetest/extracted/](#). Each answer references Microsoft Learn. Use this PDF for offline review; the interactive version lives at [app/index.html](#).

Topic 1

7 questions

Topic 1 - Question 15

Drag and drop | Domain: Deploy

topic-1/q15

You are designing end-to-end test scenarios for a business solution that uses Microsoft Dynamics 365 Sales and Dynamics 365 Finance.

You need to ensure that the business solution meets the following test requirements:

- Properly exchanges data between the Dynamics 365 apps
- Aligns with defined user workflows and business processes

Which type of testing should you use for each requirement? To answer, drag the appropriate testing types to the correct requirements. Each testing type may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

Drag sources

- Drift
- Exploratory
- Integration
- Performance
- User acceptance

Drop targets

1. Properly exchanges data between the Dynamics 365 apps - drop here
2. Aligns with defined user workflows and business processes - drop here

Answer

Properly exchanges data between the Dynamics 365 apps = Integration; Aligns with defined user workflows and business processes = User acceptance

Correct answer: - Properly exchanges data between the Dynamics 365 apps: **Integration** - Aligns with defined user workflows and business processes: **User acceptance**

Why these answers are correct

Microsoft Learn's "Types of tests" article and the broader Dynamics 365 testing strategy guidance map test types directly to these requirements:

Integration testing is described as "testing how different components or systems work together, such as data flows, interfaces, and APIs." That is precisely the test type for verifying that Dynamics 365 Sales and Dynamics 365 Finance correctly exchange data.

User acceptance testing (UAT) is described as "testing the solution from the user's perspective" and is performed by business users who run real-life process simulations to confirm the solution supports business operations and goals. UAT is the documented test type for validating alignment with user workflows and business processes, and it's a prerequisite for go-live in Success by Design.

Why the other options are wrong (for the slot they don't fit)

Drift typically refers to model drift or configuration drift monitoring, not functional testing of data exchange or workflow alignment.

Exploratory testing finds issues through ad hoc investigation. It is not the structured method for confirming data exchange between two business apps or aligning a solution to defined workflows.

Performance testing measures speed, throughput, and scalability. It does not confirm functional data

exchange or workflow alignment.

Microsoft Learn references

- Types of tests - <https://learn.microsoft.com/dynamics365/guidance/implementation-guide/testing-strategy-test-types>
- Test your Dynamics 365 solution before deployment - <https://learn.microsoft.com/dynamics365/guidance/implementation-guide/testing-strategy>
- Perform user acceptance testing in finance and operations apps (training) - <https://learn.microsoft.com/training/modules/perform-uat-finance-operations/>

Topic 1 - Question 16

Multiple choice | Domain: Deploy

topic-1/q16

A company has a Microsoft 365 tenant in Canada and multiple Microsoft Power Platform environments in Canada and the United States.

The company plans to deploy a Microsoft Copilot Studio agent to the Canadian environment that will use:

- Microsoft Dataverse data stored in Canada
- A connector that connects to an Azure OpenAI instance in the United States

You need to ensure that the agent adheres to data residency and data movement policies before being deployed.

What should you do?

- A. Ensure that the data processed by Azure OpenAI is stored in the United States.
- B. From the Microsoft Purview portal, validate the Data loss prevention settings.
- C. Migrate the tenant to the United States.
- D. Ensure that cross-region data movement is enabled for the Canadian environment and connector dependencies.

Answer

D

Correct answer: D. Ensure that cross-region data movement is enabled for the Canadian environment and connector dependencies.

Why D is correct

The agent will live in a Canadian environment using Canadian Dataverse data, but it depends on an Azure OpenAI connector hosted in the United States. Microsoft Learn's "Move data across regions for Copilots, AI agents, and generative AI features" article is explicit:

"When the **Move data across regions** checkbox is selected, your inputs (prompts) and outputs (results) might move outside of your region to the location where the generative AI feature is hosted. ... Some regions may have limited capacity, or no capacity at all. To ensure availability of Copilots and generative AI features, we may need to move the data outside of the region for processing."

For the Canada region, the table shows that data may be processed in the United States for Azure OpenAI. The administrator must explicitly enable the "Move data across regions" setting in the Power Platform admin center for the environment, and the related connector dependencies must support that flow. Without it, the agent's cross-region path is blocked or unsupported, and the data residency policy is not addressed.

This is the documented Microsoft administrative step for adhering to data residency and data movement policies in this exact scenario before the agent is deployed.

Why the other options are wrong

A. Ensure that the data processed by Azure OpenAI is stored in the United States is not the residency choice you make at the agent layer. Microsoft Learn states explicitly that "Microsoft doesn't log, store, or retain any input or output data during this process. There is no persistence of the data." You don't configure US storage for processed prompts; you control whether data movement is allowed.

B. From the Microsoft Purview portal, validate the Data loss prevention settings is good practice for sensitive data, but DLP validation alone does not address the cross-region data movement consent requirement. The Power Platform admin center is the location for that consent, not Purview.

C. Migrate the tenant to the United States is disproportionate, breaks Canadian data residency, and is not a Microsoft-recommended response to a US-hosted connector dependency.

Microsoft Learn references

- Move data across regions for Copilots, AI agents, and generative AI features - <https://learn.microsoft.com/power-platform/admin/geographical-availability-copilot>
- Configure data movement across geographic locations for generative AI - <https://learn.microsoft.com/microsoft-copilot-studio/manage-data-movement-outside-us>
- Geographic data residency in Copilot Studio - <https://learn.microsoft.com/microsoft-copilot-studio/geo-data-residency>

Topic 1 - Question 17

Hotspot | Domain: Deploy

topic-1/q17

A company has a Microsoft Copilot Studio agent for customer support.

You are reviewing and validating the following prompts:

- A prompt that has instructions to "help the customer as best you can"
- A prompt that helps retrieve product information from a knowledge base

You need to ensure that the agent delivers consistent and accurate responses.

What should you do for each prompt? To answer, select the appropriate options in the answer area.

Answer area

- **A prompt that has instructions to "help the customer as best you can":**

- ☐ Add filler words to make the prompt sound more natural and conversational.
- ☐ Keep the prompt vague to enable model flexibility.
- ☐ Rewrite the prompt with clear and task-specific instructions.

- **A prompt that helps retrieve product information from a knowledge base:**

- ☐ Add several open-ended questions to give the model broader context.
- ☐ Use responses with only reference sources and limit the response scope.
- ☐ Remove the knowledge source so that the model responds freely with general product information.

Answer

Vague prompt = Rewrite the prompt with clear and task-specific instructions; Knowledge-base prompt = Use responses with only reference sources and limit the response scope

Correct answer: - A prompt that has instructions to "help the customer as best you can": **Rewrite the**

prompt with clear and task-specific instructions. - A prompt that helps retrieve product information from a knowledge base: **Use responses with only reference sources and limit the response scope.**

Why these answers are correct

Microsoft Learn's prompt engineering guidance ("Prompt engineering techniques," "Use prompts to make your agent or agent flow perform specific tasks," "Use prompt modification to provide custom instructions to your agent") states the same principle repeatedly: be specific, give task-specific instructions, and start with clear instructions. Vague prompts like "help the customer as best you can" produce inconsistent and unreliable answers because the model has no concrete task definition. The correct fix is to rewrite the prompt with clear, task-specific instructions.

For the knowledge-base retrieval prompt, the goal is grounded, accurate responses based on the source documents. Microsoft Learn's "Best practices for integrating and deploying Copilot Studio" calls out: "Avoid relying solely on model-generated content for critical decisions, always ground outputs in verifiable data." The right pattern for knowledge-base retrieval is to constrain the response to reference sources and limit the response scope. That reduces hallucination and keeps answers traceable through citations.

Why the other options are wrong

Add filler words to make the prompt sound more natural and conversational introduces noise without adding task clarity. Microsoft's guidance is the opposite: keep prompts brief and specific.

Keep the prompt vague to enable model flexibility produces inconsistent responses. Microsoft's prompt engineering guidance explicitly warns against vague or ambiguous language.

Add several open-ended questions to give the model broader context widens the response scope, which is the opposite of what's needed for accurate knowledge-base retrieval.

Remove the knowledge source so that the model responds freely with general product information removes grounding entirely, which is the worst option for accuracy and consistency.

Microsoft Learn references

- Prompt engineering techniques - <https://learn.microsoft.com/azure/ai-foundry/openai/concepts/prompt-engineering>
- Use prompts to make your agent or agent flow perform specific tasks - <https://learn.microsoft.com/microsoft-copilot-studio/nlu-prompt-node>
- Use prompt modification to provide custom instructions to your agent - <https://learn.microsoft.com/microsoft-copilot-studio/nlu-generative-answers-prompt-modification>
- Best practices for integrating and deploying Copilot Studio - <https://learn.microsoft.com/microsoft-copilot-studio/system-service-card-copilot-studio>

Topic 1 - Question 18

Hotspot | Domain: Deploy

topic-1/q18

You are designing a testing solution for Microsoft Copilot Studio agents.

You need to validate prompt engineering best practices to ensure that the agents generate accurate and contextually relevant responses.

Which prompt validation techniques and metrics should you include in the solution? To answer, select the appropriate options in the answer area.

Answer area

- Prompt validation techniques:

- ☐ Exclude domain-specific terminology from the prompts.
- ☐ Use prompts that have varied phrasing.

☐ Use only simple, one-word prompts.

- Metrics:

☐ The number of words generated per response

☐ Response relevance and accuracy

☐ The response generation time

Answer

Prompt validation = Use prompts that have varied phrasing; Metrics = Response relevance and accuracy

Correct answer: - Prompt validation techniques: **Use prompts that have varied phrasing.** - Metrics: **Response relevance and accuracy**

Why these answers are correct

For prompt validation, Microsoft's prompt engineering guidance and Copilot Studio testing guidance both stress that you should test against varied phrasings of the same intent. Real users won't all phrase a question the same way, so validating with varied phrasing is how you confirm the agent generalizes. Microsoft Learn's "Test your agent" and "Best practices for integrating and deploying Copilot Studio" sections call out simulating real-world scenarios and varied user inputs as the way to verify the agent generates accurate, contextually relevant responses.

For the metric, Microsoft Learn's evaluation overview and Copilot Studio analytics guidance describe response relevance and accuracy as the primary qualitative metrics for prompt engineering testing. The Copilot Studio Application Card explicitly lists evaluation as the practice of "assessing agent quality, performance, safety, and reliability," with response quality (relevance and accuracy) at the center. Counting words generated or generation time tells you nothing about whether the response was correct or contextually useful.

Why the other options are wrong

Exclude domain-specific terminology from the prompts would reduce realism and break domain-specific reasoning testing. Real prompts in production use domain terminology, so testing should include it, not exclude it.

Use only simple, one-word prompts is unrealistic and doesn't test how the agent handles real user input.

The number of words generated per response is not a quality signal. A long answer can be wrong; a short answer can be correct.

The response generation time is a performance metric, not a metric for prompt engineering accuracy. You measure latency separately from response correctness.

Microsoft Learn references

- Prompt engineering techniques - <https://learn.microsoft.com/azure/ai-foundry/openai/concepts/prompt-engineering>
- Test your agent - <https://learn.microsoft.com/microsoft-copilot-studio/authoring-test-bot>
- Application Card: Microsoft Copilot Studio (Evaluation) - <https://learn.microsoft.com/microsoft-copilot-studio/system-service-card-copilot-studio>
- Evaluate generative AI applications - <https://learn.microsoft.com/azure/ai-foundry/concepts/evaluation-approach-gen-ai>

Topic 1 - Question 19

Multiple choice | Domain: Deploy

topic-1/q19

A company has two Microsoft Power Platform environments named Dev1 and Prod1. A Microsoft Copilot Studio agent named Agent1 is built into a solution in the Dev1 environment.

You plan to deploy Agent1 to Prod1.

You need to make Agent1 available to the users in Prod1. The solution must minimize administrative effort.

What should you do?

- A. Share Agent1 with the users in Prod1.
- B. Export the solution as an unmanaged solution and import the solution into Prod1.
- C. Export the solution as a managed solution and import the solution into Prod1.
- D. Create a new Copilot Studio agent in Prod1 by replicating the configuration of Agent1.

Answer

C

Correct answer: C. Export the solution as a managed solution and import the solution into Prod1.

Why C is correct

Microsoft's documented application lifecycle management (ALM) pattern for moving Copilot Studio agents between environments is to package the agent in a solution, export the solution as **managed**, and import it into the target environment.

Microsoft Learn's "Export and import agents using solutions" and "Create and manage solutions in Copilot Studio" articles are explicit on this. Managed solutions are the recommended deployment artifact for production environments because: - They prevent direct edits in production (preserving the integrity of the deployment). - They support upgrade and update flows back from Dev1. - They are the artifact type expected by Power Platform pipelines.

This minimizes administrative effort because once the managed solution is imported into Prod1, the agent and all its solution components (topics, tools, environment variables, etc.) are deployed in one operation, and future updates flow as solution upgrades rather than manual configuration.

Why the other options are wrong

A. Share Agent1 with the users in Prod1 does not move the agent to Prod1. Sharing extends access within the source environment; it doesn't deploy the agent to a different environment.

B. Export the solution as an unmanaged solution and import the solution into Prod1 is technically possible but is not the recommended pattern for production. Unmanaged solutions allow direct edits in the target environment, are harder to maintain across environments, and Microsoft's guidance is explicit that production should use managed solutions.

D. Create a new Copilot Studio agent in Prod1 by replicating the configuration of Agent1 is the most administrative effort possible (full manual rebuild) and creates drift between environments. It directly violates the "minimize administrative effort" requirement.

Microsoft Learn references

- Export and import agents using solutions - <https://learn.microsoft.com/microsoft-copilot-studio/authoring-solutions-import-export>
- Create and manage solutions in Copilot Studio -

- <https://learn.microsoft.com/microsoft-copilot-studio/authoring-solutions-overview>
- Managed and unmanaged solutions - <https://learn.microsoft.com/power-platform/alm/solution-concepts-alm>
- Use application lifecycle management in Copilot Studio - <https://learn.microsoft.com/power-platform/release-plan/2024wave2/data-platform/use-application-life-cycle-management-copilot-studio>

Topic 1 - Question 20

Multiple choice | Domain: Deploy

topic-1/q20

A company has a Microsoft Power Platform environment that contains Microsoft Dataverse data.

You create a Microsoft Copilot Studio agent named Agent1 that processes the Dataverse data.

You discover that Agent1 fails to return relevant or accurate results.

You need to improve the quality and reliability of data grounding.

What should you do?

- A. Retrain Agent1.
- B. Verify and cleanse the Dataverse data.
- C. Use an adaptive card in Agent1.
- D. Add example user inputs to the training data of Agent1.

Answer

B

Correct answer: B. Verify and cleanse the Dataverse data.

Why B is correct

The agent fails to return relevant or accurate results, and the goal is to improve the quality and reliability of **data grounding**. Grounding quality is a function of the source data quality. Microsoft Learn's Power Platform Well-Architected guidance for Copilot Studio agents states this directly: "The principle of 'garbage in, garbage out' is particularly important for agents and emphasizes the need for high-quality data. Providing accurate information to the agent ensures reliable and correct responses."

Microsoft Learn's "Improve copilot responses from Microsoft Dataverse" article also confirms that the lever for accuracy when an agent grounds in Dataverse is the data itself (cleanliness, descriptions, synonyms, glossary terms). Verifying and cleansing the Dataverse data is the targeted action that improves grounding quality, which directly improves the agent's relevance and accuracy.

Why the other options are wrong

A. Retrain Agent1 is not how Copilot Studio agents work for grounded data. Copilot Studio agents don't get "retrained" on Dataverse data. They retrieve grounded answers from the data sources at runtime, so retraining isn't the lever for grounding quality.

C. Use an adaptive card in Agent1 changes how responses are presented to users. It doesn't change the data the agent grounds in.

D. Add example user inputs to the training data of Agent1 improves intent recognition for classic-orchestration trigger phrases. It doesn't fix grounding quality from Dataverse.

Microsoft Learn references

- Improve copilot responses from Microsoft Dataverse - <https://learn.microsoft.com/power-apps/maker/data-platform/data-platform-copilot>
- Custom contact center solution with Copilot Studio agent (Reliability: garbage in, garbage out) - <https://learn.microsoft.com/power-platform/architecture/solution-ideas/agent-custom-contact-center>
- Add Dataverse tables as a knowledge source - <https://learn.microsoft.com/microsoft-copilot-studio/knowledge-add-dataverse>
- Grounding data design for AI workloads on Azure - <https://learn.microsoft.com/azure/well-architected/ai/grounding-data-design>

Topic 1 - Question 25

Multiple choice | Domain: Deploy

topic-1/q25

A company has an AI agent that automates the review of customer feedback stored in a cloud database.

You plan to generate monthly reports from the agent's output to provide insights into customer sentiment and guide product development and marketing.

You need to ensure that the data ingested by the agent is clean and suitable for the intended use.

What should you do to prepare the data?

- A. Create a workflow in Microsoft Power Automate.
- B. Identify and address biased data.
- C. Create an agent flow in Microsoft Copilot Studio.
- D. Sort the database by customer last name.

Answer

B

Correct answer: B. Identify and address biased data.

Why B is correct

The objective is to ensure the data ingested by the agent is **clean and suitable for the intended use** before generating monthly insights into customer sentiment. Microsoft Learn's responsible AI guidance for Copilot Studio and the Customer Insights sentiment analysis documentation both call out the same risk: biased input data leads to biased outputs and decisions, especially in sentiment analysis (where digital-only feedback or skewed sources misrepresent the customer base).

The "Apply responsible AI principles" article for Copilot Studio recommends using "diverse and representative data" and using "bias detection and mitigation tools" to "detect biases in the AI models" and apply corrective actions. The "Analyze sentiment in customer feedback" Customer Insights documentation explicitly warns about potential bias in feedback data and the need to address it before analysis. Identifying and addressing biased data is the data-preparation step that makes the ingested data suitable for the intended use.

Why the other options are wrong

- A. **Create a workflow in Microsoft Power Automate** is an automation step, not a data quality / bias step. It moves data; it doesn't make it clean or unbiased.
- C. **Create an agent flow in Microsoft Copilot Studio** is also automation orchestration, not data preparation.
- D. **Sort the database by customer last name** has no impact on data quality, suitability, or bias.

Microsoft Learn references

- Apply responsible AI principles (Copilot Studio) -
<https://learn.microsoft.com/microsoft-copilot-studio/guidance/responsible-ai>
- Analyze sentiment in customer feedback (potential bias) -
<https://learn.microsoft.com/dynamics365/customer-insights/data/sentiment-analysis>
- Understanding Data Curation and Management for AI Projects -
<https://learn.microsoft.com/ai/playbook/capabilities/data-curation/>
- Design training data for AI workloads on Azure -
<https://learn.microsoft.com/azure/well-architected/ai/training-data-design>

Topic 2

1 questions

Topic 2 - Question 18

Multiple choice | Domain: Deploy

topic-2/q18

You need to design an application lifecycle management (ALM) process for a Microsoft Power Platform environment that contains a solution named Solution1.

Solution1 must include a custom connector for Copilot in Microsoft Dynamics 365 Customer Service. Solution1 must meet the following requirements:

- * Ensure that the custom connector can be deployed consistently across environments as part of the ALM process.
- * Allow the custom connector to be edited only in the development environment.

What should you include in the design?

- A. Add the custom connector to GitHub.
- B. Share the custom connector.
- C. Create the custom connector in the default solution.
- D. Add the custom connector to Solution1.

Answer

D

Correct answer: D. Add the custom connector to Solution1.

Why D is correct

Power Platform application lifecycle management (ALM) treats solutions as the unit of transport between environments. To deploy a custom connector consistently across development, test, and production environments, the connector must be inside the solution that is exported from development and imported into the downstream environments.

From Microsoft Learn:

"Use [solutions](#) as containers to transport artifacts and customizations across environments."

"When you create an agent in Copilot Studio, the agent is created within a Power Platform solution. You can create custom solutions to manage your agents across multiple environments, or for pipeline deployments and other application lifecycle management (ALM) scenarios."

"Don't customize outside of a development environment. Always work in the context of solutions... Export and deploy solutions as managed, unless setting up a development environment."

By adding the connector to Solution1 and exporting Solution1 as managed, the connector is deployed consistently across environments and is editable only in the development environment, which matches both stated requirements.

Why the other options are wrong

A. Add the custom connector to GitHub. Source-controlling the connector definition in GitHub doesn't deploy it through Power Platform ALM. Solution-based deployment is still required to package and import the connector into target environments.

B. Share the custom connector. Sharing grants user permissions. It doesn't move the component across environments or restrict edits to development only.

C. Create the custom connector in the default solution. Components in the default solution can't be exported. Microsoft's ALM golden rule is to always work in custom (unmanaged) solutions during development and never customize the default solution for ALM-bound components.

Microsoft Learn references

- Establish an application lifecycle management strategy (Copilot Studio ALM) - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/alm>
- Create and manage solutions in Copilot Studio - <https://learn.microsoft.com/microsoft-copilot-studio/authoring-solutions-overview>
- Export and import agents using solutions - <https://learn.microsoft.com/microsoft-copilot-studio/authoring-solutions-import-export>
- Power Platform solution concepts - <https://learn.microsoft.com/power-platform/alm/solution-concepts-alm>

Topic 3

38 questions

Topic 3 - Question 1

Multiple choice | Domain: Deploy | Case study: Fabrikam

topic-3/q01

Which tool should you recommend to address the sensitive information concerns in the sales process?

- A. the Analytics tab in Microsoft Copilot Studio
- B. Model Context Protocol (MCP)
- C. Application Insights
- D. Microsoft Foundry Tracing UI
- E. Monitoring in Microsoft Foundry

Answer

C

Correct answer: C. Application Insights

Why C is correct

The case study describes a single low-code Copilot Studio agent grounded in Dataverse. Two requirements steer the answer:

1. "Detailed telemetry must be logged for the first created AI agent to help troubleshoot and optimize the agent during the initial AI agent adoption process."
2. "Any sensitive information, such as user IDs and names, shared via the AI agent must be tracked for future auditing."

Application Insights is the integrated telemetry option for Copilot Studio agents. Microsoft Learn shows that Copilot Studio captures detailed telemetry (conversations, exceptions, tool usage, latency, custom events) into Application Insights, and the Copilot Studio dashboard workbook in Application Insights surfaces those signals for troubleshooting and optimization. That telemetry is queryable in Log Analytics, which gives auditors the durable, indexable trail needed to track who shared which sensitive values with the agent.

Why the other options are wrong

A. The Analytics tab in Microsoft Copilot Studio shows aggregate KPIs (engagement, resolution, escalation, CSAT, generated answer quality). It does not capture per-message telemetry suitable for auditing sensitive information at the user level.

B. Model Context Protocol (MCP) is an open protocol for connecting agents to external tools and data sources. It is not an audit or telemetry tool.

D. Microsoft Foundry Tracing UI traces Microsoft Foundry agents (prompt, hosted, workflow agents) using OpenTelemetry. The Fabrikam agent is built in Copilot Studio with Dataverse, so Foundry Tracing UI is not the integrated telemetry surface for it.

E. Monitoring in Microsoft Foundry also targets Foundry agents and Foundry-hosted models. It is not the path for a Copilot Studio agent.

Microsoft Learn references

- Capture telemetry with Application Insights (Copilot Studio) - <https://learn.microsoft.com/microsoft-copilot-studio/advanced-bot-framework-composer-capture-telemetry>
- Audit Copilot Studio activities in Microsoft Purview - <https://learn.microsoft.com/microsoft-copilot-studio/admin-logging-copilot-studio>
- Analytics overview for Copilot Studio - <https://learn.microsoft.com/microsoft-copilot-studio/analytics-overview>

Topic 3 - Question 3

Multiple choice | Domain: Deploy | Case study: Fabrikam
topic-3/q03

Which tool should you recommend to help secure funding for future AI agent development?

- A. Evaluations in Microsoft Foundry
- B. the Azure Cost Optimization workbook
- C. Azure Operator Insights
- D. the Analytics tab in Microsoft Copilot Studio
- E. Direct Preference Optimization (DPO)

Answer

D

Correct answer: D. the Analytics tab in Microsoft Copilot Studio

Why D is correct

The case study states: "Sales managers must report on the adoption of the AI agent to key Fabrikam stakeholders on a monthly basis," and "funding for future AI initiatives will depend on demonstrating an increase in AI adoption month over month."

The Analytics tab in Copilot Studio is the built-in surface for measuring agent adoption and usage. Microsoft Learn describes it as the place to view "summaries of agent usage, conversation volume, and engagement," along with KPIs (engagement rate, resolution rate, escalation rate, CSAT, savings, generated answer quality, themes). Sales managers can review monthly trend lines on the Analytics page to demonstrate adoption growth and justify continued AI investment.

Why the other options are wrong

A. Evaluations in Microsoft Foundry measures AI quality, safety, and groundedness for Foundry-hosted models and agents. It is not designed to surface adoption trends for a Copilot Studio agent.

B. The Azure Cost Optimization workbook tracks Azure resource spend. It does not measure agent adoption or usage trends.

C. Azure Operator Insights is for telecommunications operators analyzing network data. It is not relevant to Copilot Studio adoption tracking.

E. Direct Preference Optimization (DPO) is a fine-tuning technique for language models. It does not measure or report adoption.

Microsoft Learn references

- Analytics overview (Copilot Studio) - <https://learn.microsoft.com/microsoft-copilot-studio/analytics-overview>
- Analyze conversational agent effectiveness - <https://learn.microsoft.com/microsoft-copilot-studio/analytics-improve-agent-effectiveness>
- Measure and improve agent performance with KPIs and analytics - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/analytics>

Topic 3 - Question 5

Multiple choice | Domain: Deploy

topic-3/q05

You are designing an AI business solution that contains the following components:

A Microsoft Power Automate workflow

A Microsoft Copilot Studio agent

A Microsoft Dataverse database -

A Microsoft Power Apps app -

As part of the application lifecycle management (ALM) process, you plan to package the components, so that they can be deployed to other environments as a group.

You need to recommend a solution that supports versioning, dependencies, and deployments.

What should you include in the recommendation?

- A. GitHub Actions
- B. Azure DevOps
- C. Microsoft Power Platform solutions

Answer

C

Correct answer: C. Microsoft Power Platform solutions

Why C is correct

Power Platform solutions are the packaging mechanism for ALM across Power Apps, Power Automate, Microsoft Dataverse, and Microsoft Copilot Studio. Microsoft Learn defines solutions as "the mechanism for implementing ALM; you use them to distribute components across environments through export and import. A component represents an artifact used in your application... Anything that can be included in a solution is a component, such as tables, columns, canvas and model-driven apps, Power Automate flows, agents, charts, and plug-ins."

Solutions track component dependencies, support versioning, and can be exported as managed or unmanaged for deployment between environments. That matches all three requirements (versioning, dependencies, deployments) for the four components listed in the question.

Why the other options are wrong

A. GitHub Actions automates CI/CD and can deploy Power Platform artifacts, but it does not itself package the components or track dependencies. It typically operates on top of solutions.

B. Azure DevOps provides CI/CD pipelines and source control. Like GitHub Actions, it automates the deployment of solutions but is not the packaging mechanism itself.

Microsoft Learn references

- Overview of ALM with Microsoft Power Platform - <https://learn.microsoft.com/power-platform/alm/overview-alm>
- Solution concepts (ALM) - <https://learn.microsoft.com/power-platform/alm/solution-concepts-alm>
- Create and manage solutions in Copilot Studio -

<https://learn.microsoft.com/microsoft-copilot-studio/authoring-solutions-overview>

- Dependency tracking for solution components -
<https://learn.microsoft.com/power-platform/alm/dependency-tracking-solution-components>

Topic 3 - Question 6

Multi-select | Domain: Deploy

topic-3/q06

A company has Microsoft 365 Copilot agents.

You need to design a security solution for the agents. The solution must meet the following requirements:

Identify and mitigate potential risks that relate to AI use.

Protect AI apps and the sensitive data processed or generated by the agents.

Support responsible AI governance by retaining and logging interactions, detecting policy violations, and investigating incidents.

Which two components should you include in the design? Each correct answer presents part of the solution.

- A. Microsoft Purview
- B. Azure AI Content Safety
- C. role-based access control (RBAC) in Microsoft Foundry
- D. Microsoft Defender

Select all that apply.

Answer

A, D

Correct answer: A. Microsoft Purview and D. Microsoft Defender

Why A and D are correct

The three requirements break cleanly across the two services Microsoft positions for AI security and governance.

Microsoft Defender (for Cloud / Defender XDR for AI) identifies and mitigates risks that relate to AI use. Microsoft Learn describes Defender for Cloud's AI security posture management and AI threat protection: discovering generative AI workloads and agents (including Copilot Studio agents through Defender for Cloud Apps), surfacing vulnerabilities and attack paths, and detecting threats such as prompt injection and abnormal execution against Microsoft Foundry workloads.

Microsoft Purview protects AI apps and the data they process and supports responsible AI governance. Microsoft Learn calls out DSPM for AI, DLP, sensitivity labels, audit, communication compliance, and insider risk management as the controls for monitoring AI interactions, retaining and logging prompts and responses, detecting policy violations, and investigating incidents.

Together they cover all three requirements: risk identification (Defender), data protection (Purview), and responsible AI governance through retention, logging, policy detection, and investigation (Purview).

Why the other options are wrong

B. Azure AI Content Safety filters harmful content in prompts and responses (violence, hate, sexual, self-harm, jailbreak detection). It is one piece of model safety, but it does not log interactions, detect

policy violations across the broader environment, or investigate incidents.

C. Role-based access control (RBAC) in Microsoft Foundry restricts who can configure or fine-tune Foundry resources. It is an identity control, not a risk-mitigation, data-protection, or incident-investigation platform for Microsoft 365 Copilot agents.

Microsoft Learn references

- Microsoft Purview data security and compliance protections for generative AI apps - <https://learn.microsoft.com/purview/ai-microsoft-purview>
- DSPM for AI in Microsoft Purview - <https://learn.microsoft.com/purview/data-security-posture-management-learn-about>
- Microsoft Defender for Cloud AI security and threat protection - <https://learn.microsoft.com/azure/defender-for-cloud/defender-for-cloud-introduction#ai-security-and-threat-protection>
- Discover AI agents and assess security posture using Microsoft Defender - <https://learn.microsoft.com/defender-xdr/security-for-ai/ai-agent-inventory>

Topic 3 - Question 7

Multiple choice | Domain: Deploy

topic-3/q07

You are creating validation criteria for a custom generative AI model that produces business reports based on internal enterprise data.

You need to assess whether the model's outputs are appropriate and meaningful for the business reports.

Which metric should you use?

- A. the number of active users interacting with the model
- B. alignment of the output to domain-specific tasks
- C. the average system resource usage during inference
- D. the model training duration

Answer

B

Correct answer: B. alignment of the output to domain-specific tasks

Why B is correct

The question asks whether the model's outputs are appropriate and meaningful for business reports based on internal enterprise data. That is a quality-of-output question, and Microsoft's evaluation guidance maps it to AI quality metrics that measure how well the response reflects the intended task and domain. Microsoft Foundry evaluators include task adherence, intent resolution, relevance, groundedness, and similar AI-assisted quality measures designed to test whether the output matches the domain task. Domain-specific alignment is the dimension that determines whether the output is appropriate and meaningful for the use case.

Why the other options are wrong

A. The number of active users interacting with the model is an adoption metric. It says nothing about whether the outputs themselves are appropriate or meaningful.

C. The average system resource usage during inference is an operational metric (compute, latency, cost). It does not measure output quality.

D. The model training duration is a process metric for fine-tuning. It does not measure whether the resulting outputs are appropriate for business reports.

Microsoft Learn references

- Built-in evaluators reference (Microsoft Foundry) - <https://learn.microsoft.com/azure/foundry/concepts/built-in-evaluators>
- Observability in generative AI (evaluation, monitoring, tracing) - <https://learn.microsoft.com/azure/foundry/concepts/observability>
- Evaluation metrics built-in (Azure AI Studio) - <https://learn.microsoft.com/azure/ai-studio/concepts/evaluation-metrics-built-in>

Topic 3 - Question 8

Multiple choice | Domain: Deploy

topic-3/q08

A company has Microsoft Foundry agents that generate responses by using Azure OpenAI resources. The agents are deployed to both the United States and Europe.

A company mandate states that the agents and their grounding data must adhere to data residency and movement regulations.

You need to recommend a governance solution for the agents.

What should you include in the recommendation?

- A. Microsoft Defender for Cloud
- B. Azure Policy
- C. Azure Monitor
- D. Microsoft Purview

Answer

B

Correct answer: B. Azure Policy

Why B is correct

The requirement is that agents and their grounding data must adhere to data residency and movement regulations. Azure Policy is the Azure service for enforcing where resources can be deployed. Microsoft Learn calls out the pattern explicitly: "If you have data residency requirements that require you to keep all your data in a single Azure region, you can enforce... by using an Azure Policy. You can also enforce a policy that the [resources] are not geo-replicated to other regions." Policy can restrict allowed locations for Azure OpenAI accounts, Foundry resources, storage, and other services that hold grounding data, and it continuously evaluates resources for compliance.

Why the other options are wrong

A. Microsoft Defender for Cloud assesses security posture and detects threats. It does not enforce regional deployment constraints.

C. Azure Monitor collects logs and metrics for diagnostics and alerts. It does not enforce data residency.

D. Microsoft Purview governs and protects data (DSPM for AI, DLP, sensitivity labels, audit). It is not the tool for enforcing the Azure region in which resources are deployed.

Microsoft Learn references

- Use Azure Policy to enforce residency requirements - <https://learn.microsoft.com/azure/cosmos-db/data-residency#use-azure-policy-to-enforce-the-residency-requirements>
- Move data across regions for Copilots and generative AI features - <https://learn.microsoft.com/power-platform/admin/geographical-availability-copilot>
- Azure Policy overview - <https://learn.microsoft.com/azure/governance/policy/overview>

Topic 3 - Question 9

Multiple choice | Domain: Deploy

topic-3/q09

A company has a Microsoft Copilot Studio agent that uses custom connectors to interact with enterprise APIs.

You need to recommend an application lifecycle management (ALM) process to ensure that the connectors are deployed consistently across development, test, and production environments and meet governance and traceability requirements.

What should you recommend?

- A. Deploy the APIs as Azure Functions.
- B. Manage the connectors as solution components and deploy the components by using ALM pipelines.
- C. Maintain connector definitions in environment variables.
- D. Export and import the connectors between the environments as unmanaged solutions.

Answer

B

Correct answer: B. Manage the connectors as solution components and deploy the components by using ALM pipelines.

Why B is correct

Custom connectors are solution-aware components in Power Platform. Microsoft's ALM guidance is to package the connector (along with the agent and any related artifacts) inside a Dataverse solution and promote it through an automated ALM pipeline (Pipelines in Power Platform, Azure DevOps with the Power Platform Build Tools, or GitHub Actions for Power Platform). This approach gives consistent deployment across development, test, and production, and it produces traceable deployment records that satisfy governance requirements.

Why the other options are wrong

A. Deploy the APIs as Azure Functions. Hosting the API differently does not address the ALM problem. The connector and the agent still need a packaging and promotion strategy.

C. Maintain connector definitions in environment variables. Environment variables hold values that change between environments (URLs, IDs, secrets). They do not move connector definitions between environments.

D. Export and import the connectors between the environments as unmanaged solutions. Production should receive managed solutions, not unmanaged ones. Unmanaged in production prevents clean upgrade and rollback paths and undermines traceability.

Microsoft Learn references

- Establish an application lifecycle management strategy (Copilot Studio) - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/alm>
- Pipelines in Power Platform - <https://learn.microsoft.com/power-platform/alm/pipelines>
- Power Platform Build Tools for Azure DevOps - <https://learn.microsoft.com/power-platform/alm/devops-build-tools>
- GitHub Actions for Power Platform - <https://learn.microsoft.com/power-platform/alm/devops-github-actions>

Topic 3 - Question 10

Hotspot | Domain: Deploy

topic-3/q10

HOTSPOT -

A company plans to implement an AI solution that will contain a Microsoft Copilot Studio agent and a Microsoft Foundry agent. The solution will be stored in a source code repository.

You need to recommend a deployment method for each agent. The solution must meet the following requirements:

A test environment must be used before a deployment to production.

Production must be isolated from development and testing.

The deployment must be repeatable and fully automated.

The solution must NOT require manual intervention.

Which deployment method should you recommend for each agent? To answer, select the appropriate options in the answer area.

Answer area

- Copilot Studio:

- ☐ Export from the source code repository and import to the target environment.
- ☐ Use a Bicep file.
- ☐ Use a Microsoft Power Platform deployment pipeline.

- Microsoft Foundry:

- ☐ Use a Bicep file.
- ☐ Use a Microsoft Power Platform deployment pipeline.
- ☐ Use an Azure DevOps pipeline.

Answer

Copilot Studio = Microsoft Power Platform deployment pipeline; Microsoft Foundry = Azure DevOps pipeline

Correct answer: - Copilot Studio: Use a Microsoft Power Platform deployment pipeline. - Microsoft Foundry: Use an Azure DevOps pipeline.

Why these selections are correct

Copilot Studio agents are Power Platform solution components. Microsoft's documented ALM path for Copilot Studio is to package the agent inside a solution and promote it across dev, test, and production with Pipelines in Power Platform. Pipelines are repeatable, automated, and require no manual intervention, and they natively understand Dataverse and Copilot Studio dependencies. Bicep does not deploy Copilot Studio agents.

Microsoft Foundry agents are Azure resources. Microsoft Learn provides Bicep templates and Azure CLI/PowerShell deployment patterns for Foundry. To make the deployment fully repeatable across dev, test, and production with isolation and no manual intervention, the template should run inside an Azure DevOps pipeline (or GitHub Actions equivalent) using the `AzureResourceManagerTemplateDeployment` task. Pipelines in Power Platform and Bicep on its own do not satisfy "repeatable and fully automated" end-to-end for a Foundry agent stored in source control. An Azure DevOps pipeline does, and it can invoke the Bicep template as part of its job.

Why the other options are wrong

Copilot Studio with Bicep file is not supported. Bicep cannot deploy Copilot Studio agents.

Copilot Studio export from source repo, import to target is a manual operation. It violates the "no manual intervention" rule.

Microsoft Foundry with Power Platform deployment pipeline is not supported. Power Platform pipelines target Dataverse solution components, not Azure resources.

Microsoft Foundry with Bicep file alone describes the artifact, not the automation. A Bicep file by itself does not deploy automatically; you need a pipeline orchestrating the deployment.

Microsoft Learn references

- Pipelines in Power Platform - <https://learn.microsoft.com/power-platform/alm/pipelines>
- Solution management in Microsoft Copilot Studio - <https://learn.microsoft.com/microsoft-copilot-studio/authoring-solutions-overview>
- Quickstart: Deploy a Microsoft Foundry resource by using a Bicep file - <https://learn.microsoft.com/azure/foundry/how-to/create-resource-template>
- Integrate Bicep with Azure Pipelines - <https://learn.microsoft.com/azure/azure-resource-manager/bicep/add-template-to-azure-pipelines>
- `AzureResourceManagerTemplateDeployment@3` task - <https://learn.microsoft.com/azure/devops/pipelines/tasks/reference/azure-resource-manager-template-deployment-v3>

Topic 3 - Question 11

Hotspot | Domain: Deploy

topic-3/q11

HOTSPOT -

A company has a Microsoft Copilot Studio prompt-and-response agent.

You need to ensure that the agent meets the following requirements:

Provides effective and relevant responses

Provides conversational outcomes

Which metric should you use for each requirement? To answer, select the appropriate options in the answer area.

Answer area

- Provides effective and relevant responses:

- ☐ Generated answer rate and quality
- ☐ Reactions
- ☐ Tool use

- Provides conversational outcomes:

- ☐ Satisfaction

- ☐ Tool use
- ☐ Topics by outcome

Answer

Effective and relevant responses = Generated answer rate and quality; Conversational outcomes = Topics by outcome

Correct answer: - Provides effective and relevant responses: Generated answer rate and quality - Provides conversational outcomes: Topics by outcome

Why these selections are correct

Generated answer rate and quality measures whether the agent answered the user's question and whether the answer is good. Microsoft Learn states that Copilot Studio "looks at a sample set of answered questions and analyzes different quality, including completeness, relevance, and level of groundedness of a response." That maps directly to "effective and relevant responses." Reactions show user thumbs up/down feedback only, and Tool use measures tool invocations rather than answer quality, so Generated answer rate and quality is the correct match.

Topics by outcome shows the relationship between topics and the outcomes of the conversations they led to (resolved, escalated, abandoned). Microsoft Learn describes the topic analytics pane as showing "a chart showing the aggregated breakdown by outcome for the selected period." Mapping conversational outcomes to the topics that produced them is precisely what this metric is for. Satisfaction tracks CSAT and sentiment but not outcomes, and Tool use measures tool calls.

Microsoft Learn references

- Analyze conversational agent effectiveness (Generated answer rate and quality, Conversation outcomes) - <https://learn.microsoft.com/microsoft-copilot-studio/analytics-improve-agent-effectiveness>
- Analyze topic usage in Copilot Studio - <https://learn.microsoft.com/microsoft-copilot-studio/analytics-topic-details>
- Measure agent outcomes - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/measuring-outcomes>

Topic 3 - Question 12

Multi-select | Domain: Deploy

topic-3/q12

A company extends Copilot in Microsoft Dynamics 365 Customer Service.

You need to recommend an automated application lifecycle management (ALM) process so that the Copilot components can be safely developed, tested, and promoted to production.

Which two actions should you include in the ALM process? Each correct answer presents part of the solution.

- A. Use an unmanaged solution in production.
- B. Rebuild the agents in each environment.
- C. Use Microsoft Power Platform pipelines.
- D. Include the components in a solution.
- E. Store the agent transcripts in source control.

Select all that apply.

Answer

C, D

Correct answer: C. Use Microsoft Power Platform pipelines and D. Include the components in a solution.

Why C and D are correct

Copilot extensions in Dynamics 365 Customer Service are Copilot Studio components, and Power Platform's prescribed ALM pattern packages them in a solution and promotes them through pipelines.

Microsoft Learn states: "When you create an agent in Copilot Studio, the agent is created within a Power Platform solution. You can create custom solutions to manage your agents across multiple environments, or for pipeline deployments and other application lifecycle management (ALM) scenarios." Pipelines in Power Platform "democratize ALM... by bringing ALM automation and continuous integration and continuous delivery (CI/CD) capabilities into the service."

Together, packaging in a solution (D) and promoting through pipelines (C) provide safe, repeatable development, test, and production movement.

Why the other options are wrong

A. Use an unmanaged solution in production. Microsoft's ALM golden rule is to deploy managed solutions to production. Unmanaged in production prevents clean uninstall and upgrade.

B. Rebuild the agents in each environment. Manual rebuilds break consistency and traceability and bypass ALM entirely.

E. Store the agent transcripts in source control. Transcripts are runtime conversation data, not ALM artifacts. They do not belong in a deployment pipeline.

Microsoft Learn references

- Establish an application lifecycle management strategy (Copilot Studio) - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/alm>
- Pipelines in Power Platform - <https://learn.microsoft.com/power-platform/alm/pipelines>
- Create and manage solutions in Copilot Studio - <https://learn.microsoft.com/microsoft-copilot-studio/authoring-solutions-overview>
- ALM basics with Microsoft Power Platform - <https://learn.microsoft.com/power-platform/alm/basics-alm>

Topic 3 - Question 13

Hotspot | Domain: Deploy

topic-3/q13

HOTSPOT -

You are designing a testing solution for a Microsoft Copilot Studio agent that integrates with Microsoft Dynamics 365 Customer Service and Dynamics 365 Sales.

You need to design end-to-end scenarios to test the agent's ability to perform the following actions:

Coordinate tasks and data interactions across both Dynamics 365 apps.

Interpret user input and provide contextually relevant outputs.

Which test scenario and metric should you include in the design? To answer, select the appropriate options in the answer area.

Answer area

- Test scenario:

☐ In each app, test isolated tasks without using workflows.

☐ Run task-based scenarios that involve both apps.

☐ Test visual consistency across both apps.

- Metric:

☐ Measure the initial prompt response time for each app.

☐ Track the average click rate across both apps.

☐ Track the successful completion of cross-app tasks.

Answer

Test scenario = Run task-based scenarios that involve both apps; Metric = Track the successful completion of cross-app tasks

Correct answer: - Test scenario: Run task-based scenarios that involve both apps. - Metric: Track the successful completion of cross-app tasks.

Why these selections are correct

The agent must coordinate tasks across Dynamics 365 Customer Service and Dynamics 365 Sales and interpret user input to produce contextually relevant outputs. End-to-end testing for that requires task-based scenarios that exercise both applications together. Isolated tests in each app would not validate cross-app coordination, and visual consistency tests do not validate functional behavior.

Microsoft's evaluation guidance for Copilot Studio agents focuses on real task execution. The Copilot Studio Kit and Power Platform pipelines run automated test cases that validate agent behavior end-to-end before promotion. Tracking successful completion of those cross-app tasks is the metric that reflects whether the agent actually performed the intended work. Initial prompt response time measures latency only, and click rate is not a meaningful agent quality metric.

Microsoft Learn references

- Automate testing and deployment of agents using pipelines in Power Platform - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/kit-automate-test-deploy>
- Copilot Studio Kit overview - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/kit-overview>
- Improve your Copilot Studio projects (analytics and outcomes) - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/improve-overview>
- Measure agent outcomes - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/measuring-outcomes>

Topic 3 - Question 14

Multiple choice | Domain: Deploy

topic-3/q14

A company has multiple AI models that support generation of sales transactions.

Each release of the models must be reviewed by a security and compliance team before being deployed to the production environment. The security and compliance team must have access to prior versions to properly determine potential exposures introduced.

You need to recommend a solution to evaluate the impact of each deployment to production. The solution must enhance business continuity.

What should you recommend?

- A. Create a central model registry that uses version history.
- B. Establish a promotion process by using a quality gate.
- C. Implement version control for all the AI system components.

D. Track model retirement schedules to prevent service disruptions.

Answer

A

Correct answer: A. Create a central model registry that uses version history.

Why A is correct

The requirement combines three things: a security and compliance review of each release, access to prior versions for impact analysis, and business continuity. A central model registry with version history is the canonical pattern for that. It tracks every published model version with metadata, lineage, and lifecycle stage; security and compliance teams can review the new version and compare it against prior versions, and the registry preserves rollback paths if a release introduces regressions or exposures, supporting business continuity.

Microsoft's GenAIOps guidance and Foundry Control Plane both center on registered model assets with versioning, evaluation results, and observability tied to each version. That is the structure that lets reviewers see what changed release over release and enables rapid rollback.

Why the other options are wrong

B. Establish a promotion process by using a quality gate. A quality gate enforces tests before promotion, but it does not by itself give the security team access to prior versions or guarantee business continuity if a deployed release fails.

C. Implement version control for all the AI system components. Source-level version control covers prompts, code, and config, but not the model artifacts and their evaluation history. A model registry is the purpose-built complement.

D. Track model retirement schedules to prevent service disruptions. Retirement tracking helps with vendor model deprecation, not the impact assessment of new releases.

Microsoft Learn references

- Generative AI operations for organizations with MLOps investments - <https://learn.microsoft.com/azure/architecture/ai-ml/guide/genaiops-for-mlops>
- Microsoft Foundry Control Plane (assets, observability, compliance) - <https://learn.microsoft.com/azure/foundry/control-plane/overview>
- Manage agents at scale in Microsoft Foundry Control Plane - <https://learn.microsoft.com/azure/foundry/control-plane/how-to-manage-agents>

Topic 3 - Question 15

Drag and drop | Domain: Deploy

topic-3/q15

DRAG DROP -

A company has an AI solution that uses a Microsoft Copilot Studio agent.

You need to monitor the agent's performance. The solution must meet the following requirements:

Monitor the agent's telemetry in near-real-time (NRT).

Download transcripts of full conversations.

Monitor the agent's usage and performance.

What should you use for each requirement? To answer, drag the appropriate options to the correct requirements. Each option may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

Drag sources

- Application Insights
- Copilot Studio
- Log Analytics
- Microsoft Power Apps

Drop targets

1. Monitor the agent's telemetry in NRT: - drop here
2. Download transcripts of full conversations: - drop here
3. Monitor the agent's usage and performance: - drop here

Answer

NRT telemetry = Application Insights; Download transcripts = Copilot Studio; Usage and performance = Copilot Studio

Correct answer: - Monitor the agent's telemetry in NRT: Application Insights - Download transcripts of full conversations: Copilot Studio - Monitor the agent's usage and performance: Copilot Studio

Why these selections are correct

Telemetry in near-real-time: Microsoft Learn shows that Copilot Studio sends detailed telemetry (conversations, tool usage, exceptions, latency, custom events) to Application Insights, which is designed for near-real-time analysis through Live Metrics, Log Analytics, and the Copilot Studio Dashboard workbook. Application Insights is the integrated NRT signal source.

Download transcripts of full conversations: Conversation transcripts are stored in Dataverse and accessible from Copilot Studio. Microsoft Learn states: "You can download conversation transcripts a few minutes after the conversation times out... You can download them in Dataverse using the Power Apps portal and as session chat transcripts using the Copilot Studio app." Copilot Studio is the surface for downloading transcripts.

Usage and performance: The Analytics page in Copilot Studio aggregates usage metrics (sessions, engagement, resolution, escalation, savings, generated answer quality, themes). It is the maker-facing view for usage and performance.

Log Analytics could host the data behind Application Insights, but the question lists Application Insights as a separate option, and Log Analytics is not the labeled NRT entry point. Microsoft Power Apps is unrelated to agent monitoring.

Microsoft Learn references

- Capture telemetry with Application Insights (Copilot Studio) - <https://learn.microsoft.com/microsoft-copilot-studio/advanced-bot-framework-composer-capture-telemetry>
- Download agent session transcripts - <https://learn.microsoft.com/microsoft-copilot-studio/analytics-sessions>
- Analytics overview (Copilot Studio) - <https://learn.microsoft.com/microsoft-copilot-studio/analytics-overview>
- Analyze conversational agent effectiveness - <https://learn.microsoft.com/microsoft-copilot-studio/analytics-improve-agent-effectiveness>

Topic 3 - Question 16

Hotspot | Domain: Deploy

topic-3/q16

HOTSPOT -

A company deploys a Microsoft Copilot Studio agent that integrates with a Microsoft Power Automate desktop flow.

You need to recommend a testing solution that meets the following requirements:

Test cases must validate the most recent changes to the agent before the agent is released.

The flow must be validated as part of the agent's orchestration.

What should you recommend for each requirement? To answer, select the appropriate options in the answer area.

Answer area

- Validate the most recent changes to the agent before release:

- ☐ Publish the agent to a channel and test the agent on live users.
- ☐ Run tests against the latest unpublished version of the agent.
- ☐ Run tests against the production version of the agent.

- Validate the flow as part of the agent's orchestration:

- ☐ Add the flow to the agent as a tool.
- ☐ Add the flow to a canvas app.
- ☐ Use the Power Automate for desktop console.

Answer

Recent changes = Run tests against the latest unpublished version of the agent; Validate the flow = Add the flow to the agent as a tool

Correct answer: - Validate the most recent changes to the agent before release: Run tests against the latest unpublished version of the agent. - Validate the flow as part of the agent's orchestration: Add the flow to the agent as a tool.

Why these selections are correct

Latest unpublished version: Microsoft's testing guidance for Copilot Studio is to test agents in the maker authoring environment before publishing. The Test panel and the Copilot Studio Kit run tests against the unpublished, in-progress version so changes are validated before they reach users. Publishing the agent to a channel and testing on live users skips validation, and the production version is what you are trying to gate against.

Flow as a tool: Agent flows are added to a Copilot Studio agent as tools so the orchestrator can invoke them as part of conversation logic. Microsoft Learn says that to add a flow to an agent it "must be a solution flow and contain the When an agent calls the flow trigger and the Respond to the agent action," and the agent invokes it through the orchestrator. Validating the flow as part of the agent's orchestration therefore requires adding the flow as a tool. A canvas app is unrelated, and the Power Automate for desktop console is for unattended desktop flow runs, not for orchestrated agent calls.

Microsoft Learn references

- Use agent flows with your agent - <https://learn.microsoft.com/microsoft-copilot-studio/advanced-flow>
- Call an agent flow - <https://learn.microsoft.com/microsoft-copilot-studio/advanced-use-flow>
- Agent flows in Microsoft Copilot Studio FAQ -

<https://learn.microsoft.com/microsoft-copilot-studio/flows-faqs>

- Automate testing and deployment of agents using pipelines in Power Platform - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/kit-automate-test-deploy>

Topic 3 - Question 17

Multi-select | Domain: Deploy

topic-3/q17

A company has a Microsoft Copilot Studio agent that provides answers based on a knowledge base for customer support.

Users report that, occasionally, the agent provides inaccurate answers.

You need to use metrics from the Analytics tab in Copilot Studio to identify the cause of the inaccuracies.

Which two options should you use? Each correct answer presents part of the solution.

- A. survey results
- B. session information and session outcomes
- C. topic usage and topics with low resolution
- D. engagement, resolution, and escalation rates
- E. quality of generated answers

Select all that apply.

Answer

C, E

Correct answer: C. topic usage and topics with low resolution and E. quality of generated answers

Why C and E are correct

The agent draws on a knowledge base, and users report inaccurate answers. Two Analytics signals point at the cause.

Topic usage and topics with low resolution identifies topics that route a high share of sessions into unresolved or escalated outcomes. Microsoft Learn lists "Topics with low resolution" among the optimization metrics to focus on monthly: "Resolution rate for an individual topic." A low resolution rate on a specific topic typically means the topic's logic, knowledge mapping, or grounding has gaps, which directly correlates with inaccurate answers.

Quality of generated answers measures whether the AI-generated responses are good. Microsoft Learn describes the Generated answer rate and quality section: "Copilot Studio looks at a sample set of answered questions and analyzes different quality, including completeness, relevance, and level of groundedness of a response. If the answer meets a set standard, Copilot Studio labels the answer as Good quality." The "Poor" reasons (irrelevant, incomplete, ungrounded) explain why answers are inaccurate and identify which knowledge sources are responsible.

Why the other options are wrong

A. Survey results capture user satisfaction (CSAT) but do not tell you why an answer was wrong.

B. Session information and session outcomes show whether sessions resolved, escalated, or were abandoned. They do not isolate inaccuracy in a specific knowledge-grounded answer.

D. Engagement, resolution, and escalation rates are top-line KPIs. They show that something is wrong but not which topic or knowledge source produced the inaccurate answer.

Microsoft Learn references

- Analyze conversational agent effectiveness (Generated answer rate and quality) - <https://learn.microsoft.com/microsoft-copilot-studio/analytics-improve-agent-effectiveness>
- Measure and improve agent performance with KPIs and analytics - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/analytics>
- Analyze topic usage in Copilot Studio - <https://learn.microsoft.com/microsoft-copilot-studio/analytics-topic-details>
- Topic escalation analysis - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/deflection-topic-escalation-analysis>

Topic 3 - Question 18

Multi-select | Domain: Deploy

topic-3/q18

A company uses a fine-tuned Microsoft Foundry model that requires frequent updates as new customer feedback becomes available.

You need to design an application lifecycle management (ALM) process that meets the following requirements:

Data changes must be tracked and versioned.

The model must be retrained consistently by using approved training data.

Which two actions should you include in the design? Each correct answer presents part of the solution.

- A. Associate the storage location to the fine-tuning job.
- B. Create a content filter.
- C. Store the training data in Azure Files.
- D. Upload the training data to Microsoft Foundry data files
- E. Store the training data in Azure Blob Storage that has version control enabled.

Select all that apply.

Answer

A, E

Correct answer: A. Associate the storage location to the fine-tuning job and E. Store the training data in Azure Blob Storage that has version control enabled.

Why A and E are correct

The two requirements are tracking and versioning training data, and consistent retraining with approved data.

Azure Blob Storage with versioning enabled is the durable, versioned data store Microsoft recommends for fine-tuning datasets. Microsoft Learn states: "For large data files, we recommend that you import from an Azure Blob store. Large files can become unstable when uploaded through multipart forms because the requests are atomic and can't be retried or resumed." Blob Storage's blob versioning automatically maintains previous versions of objects, so you can audit changes and roll back to an approved dataset version.

Associating the storage location to the fine-tuning job ensures every retraining run pulls from the same governed source. The fine-tuning job references the dataset in Blob Storage, the version is recorded with the job, and the team retrains consistently against approved data.

Why the other options are wrong

B. Create a content filter. Content filters control output safety. They do not version training data.

C. Store the training data in Azure Files. Azure Files is an SMB/NFS file share. It is not the recommended pattern for fine-tuning datasets and lacks the same versioning model. Microsoft Learn specifically recommends Blob Storage for large fine-tuning files.

D. Upload the training data to Microsoft Foundry data files. Direct uploads through Foundry are appropriate for small files, but they do not give you versioning or change tracking. The recommendation is Blob Storage with versioning.

Microsoft Learn references

- Customize a model with fine-tuning (upload training data) - <https://learn.microsoft.com/azure/ai-foundry/openai/how-to/fine-tuning>
- What is Azure Blob Storage? - <https://learn.microsoft.com/azure/storage/blobs/storage-blobs-overview>
- Blob versioning - <https://learn.microsoft.com/azure/storage/blobs/versioning-overview>
- Design training data for AI workloads on Azure - <https://learn.microsoft.com/azure/well-architected/ai/training-data-design>

Topic 3 - Question 19

Hotspot | Domain: Deploy

topic-3/q19

HOTSPOT -

A company deploys agents that generate responses by using Azure OpenAI resources. The agents are deployed to both the United States and Europe.

You need to recommend a governance solution that meets the following requirements:

Enforces the deployment of the resources to only approved Azure regions

Provides continuous compliance verification of the resources

What should you include in the recommendation for each requirement? To answer, select the appropriate options in the answer area.

Answer area

- Enforces the deployment of the resources to only approved regions:

- ☐ Azure Monitor
- ☐ Azure Policy
- ☐ Microsoft Defender for Cloud
- ☐ Microsoft Purview
- ☐ Microsoft Sentinel

- Provides continuous compliance verification of the resources:

- ☐ Azure Monitor
- ☐ Azure Policy
- ☐ Microsoft Defender for Cloud
- ☐ Microsoft Purview
- ☐ Microsoft Sentinel

Answer

Enforce regions = Azure Policy; Continuous compliance verification = Microsoft Defender for Cloud

Correct answer: - Enforces the deployment of the resources to only approved regions: Azure Policy - Provides continuous compliance verification of the resources: Microsoft Defender for Cloud

Why these selections are correct

Azure Policy is the Azure service for enforcing what can be deployed and where. The built-in `allowedLocations` policy and similar definitions block deployments to non-approved regions and audit existing resources. That is exactly the "enforce deployment to only approved regions" requirement.

Microsoft Defender for Cloud provides continuous regulatory compliance assessment through the Compliance dashboard. Defender for Cloud continuously evaluates resources against built-in or custom standards (NIST, ISO 27001, PCI DSS, custom standards), produces compliance scores, and surfaces non-compliant resources for remediation. Microsoft Learn describes this as cloud security posture management with continuous monitoring.

Azure Monitor handles logs and metrics, Microsoft Purview governs data, and Microsoft Sentinel is the SIEM for threat investigation. None of those is the continuous-compliance evaluator that Defender for Cloud is.

Microsoft Learn references

- Use Azure Policy to enforce residency - <https://learn.microsoft.com/azure/cosmos-db/data-residency#use-azure-policy-to-enforce-the-residency-requirements>
- Azure Policy overview - <https://learn.microsoft.com/azure/governance/policy/overview>
- Microsoft Defender for Cloud overview (CSPM and compliance) - <https://learn.microsoft.com/azure/defender-for-cloud/defender-for-cloud-introduction>
- Regulatory compliance in Defender for Cloud - <https://learn.microsoft.com/azure/defender-for-cloud/concept-regulatory-compliance>

Topic 3 - Question 20

Multiple choice | Domain: Deploy

topic-3/q20

A company has an AI solution that uses Azure OpenAI models.

You need to recommend a governance solution that monitors and audits changes to model configurations and data usage. The solution must minimize administrative effort.

What should you include in the recommendation?

- A. Azure Monitor
- B. Azure Stream Analytics
- C. Azure API Management
- D. Azure Policy
- E. Microsoft Purview

Answer

E

Correct answer: E. Microsoft Purview

Why E is correct

The requirement is governance that monitors and audits changes to model configurations and data usage with minimal administrative effort. Microsoft Purview is the unified governance and compliance platform Microsoft positions for AI workloads. Microsoft Learn describes Purview's native integration with Azure AI Foundry: "Azure Admins can turn on the setting for any given Azure subscription. There's no action your developers need to take. This setting enables data from all Azure AI based applications running in that subscription to be sent to Microsoft Purview to support governance and compliance outcomes." Purview Audit captures activity logs, DSPM for AI provides visibility into data interactions and policy violations, and the integration is configured at the subscription level, which minimizes administrative effort.

Why the other options are wrong

A. Azure Monitor collects logs and metrics for diagnostics. It can ingest activity but does not provide AI-specific governance, audit, and data usage tracking out of the box.

B. Azure Stream Analytics processes streaming data. It is not a governance tool.

C. Azure API Management front-loads APIs and enforces gateway policies. It does not audit AI model configuration changes.

D. Azure Policy enforces rules on resource configurations and detects non-compliance. It does not capture detailed audit trails for data usage or AI app interactions.

Microsoft Learn references

- Use Microsoft Purview to manage data security and compliance for Microsoft Foundry - <https://learn.microsoft.com/purview/ai-azure-foundry>
- Microsoft Purview data security and compliance protections for generative AI apps - <https://learn.microsoft.com/purview/ai-microsoft-purview>
- Audit Copilot Studio activities in Microsoft Purview - <https://learn.microsoft.com/microsoft-copilot-studio/admin-logging-copilot-studio>
- Enable Data Security for Azure AI with Microsoft Purview - <https://learn.microsoft.com/azure/defender-for-cloud/ai-onboarding#enable-data-security-for-azure-ai-with-microsoft-purview>

Topic 3 - Question 21

Hotspot | Domain: Deploy

topic-3/q21

HOTSPOT -

A company uses Azure OpenAI models that use grounding data from Microsoft Fabric for agents. The models are fine-tuned by using proprietary datasets.

You need to design a governance solution that meets the following requirements:

Restricts access to the grounding data to only assigned roles

Restricts model fine-tuning to only the AI engineering team

What should you include in the design? To answer, select the appropriate options in the answer area.

Answer area

- Restricts access to the grounding data:

- ☐ Azure AI Content Safety
- ☐ Azure Monitor alerts
- ☐ Azure Policy compliance rules
- ☐ Microsoft Purview access policies

- Restricts model fine-tuning:

- ☐ Azure Policy compliance rules
- ☐ Azure Resource Manager (ARM) resource locks
- ☐ Microsoft Entra Conditional Access
- ☐ Role-based access control (RBAC) in Microsoft Foundry

Answer

Grounding data = Microsoft Purview access policies; Fine-tuning = Role-based access control (RBAC) in Microsoft Foundry

Correct answer: - Restricts access to the grounding data: Microsoft Purview access policies - Restricts model fine-tuning: Role-based access control (RBAC) in Microsoft Foundry

Why these selections are correct

Grounding data in Microsoft Fabric is governed by Microsoft Purview. Purview access policies (and the Unified Catalog) control who can read or use the data assets that ground the AI models. Microsoft Learn describes Purview as the platform for "data security solutions" and "unified data governance solutions" and details data access governance for Fabric assets through Purview. Azure Policy enforces resource-level configuration, Azure Monitor alerts surface signals, and Content Safety filters output, none of which is the data access control plane.

Fine-tuning is controlled by RBAC in Microsoft Foundry. Foundry uses Azure RBAC to scope what each principal can do at the Foundry resource and project level. Roles such as Microsoft Foundry Developer, Microsoft Foundry Project Manager, and similar control whether a user can create or run fine-tuning jobs. Microsoft's RBAC for Foundry article enumerates these built-in roles. Conditional Access governs sign-in conditions for users, ARM resource locks prevent resource deletion or modification, and Azure Policy governs resource configuration. None of those scope a user's ability to launch a fine-tuning job the way RBAC does.

Microsoft Learn references

- Microsoft Purview overview - <https://learn.microsoft.com/purview/purview>
- Data governance with Microsoft Purview (Unified Catalog) - <https://learn.microsoft.com/purview/data-governance-overview>
- Role-based access control for Microsoft Foundry - <https://learn.microsoft.com/azure/foundry/concepts/rbac-foundry>
- Microsoft Fabric and Purview integration - <https://learn.microsoft.com/fabric/governance/microsoft-purview-fabric>

Topic 3 - Question 22

Drag and drop | Domain: Deploy

topic-3/q22

DRAG DROP -

A company has a Microsoft Copilot Studio agent that has been in production for three months.

The agent has received positive feedback from users.

You need to identify the number of questions unanswered by the agent and the number of abandoned sessions between the users and the agent.

Which Copilot Studio insights should you use? To answer, drag the appropriate insights to the correct requirements. Each insight may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

Drag sources

- Conversation outcomes
- Generated answer rate and quality
- Reactions
- Survey results

Drop targets

1. The number of unanswered questions: - drop here
2. The number of abandoned sessions: - drop here

Answer

Unanswered questions = Generated answer rate and quality; Abandoned sessions = Conversation outcomes

Correct answer: - The number of unanswered questions: Generated answer rate and quality - The number of abandoned sessions: Conversation outcomes

Why these selections are correct

Generated answer rate and quality explicitly tracks unanswered questions. Microsoft Learn states: "The Answer rate shows the number of answered and unanswered questions within the selected time period and the percentage change over time." It also breaks down quality categories for poor answers (irrelevant, incomplete, ungrounded). That is the metric for unanswered questions.

Conversation outcomes tracks the four session outcomes: Resolved, Escalated, Abandoned, and Unengaged. Microsoft Learn defines abandoned as "an engaged session that isn't resolved or escalated and is inactive after 60 minutes from the beginning of the session" and exposes the count and percentage on the Conversation outcomes chart. Reactions captures thumbs up/down, and Survey results captures CSAT, neither of which is the abandoned-session count.

Microsoft Learn references

- Analyze conversational agent effectiveness (Generated answer rate and quality, Conversation outcomes) - <https://learn.microsoft.com/microsoft-copilot-studio/analytics-improve-agent-effectiveness>
- Measure agent outcomes - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/measuring-outcomes>
- Deflection overview (key metrics including Abandon Rate) - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/deflection-overview>

Topic 3 - Question 23

Multiple choice | Domain: Deploy
topic-3/q23

You are evaluating a Microsoft Copilot Studio agent that supports Microsoft Dynamics 365 Customer Service representatives.

You need to recommend a testing solution that meets the following requirements:

Evaluates agent effectiveness during active sessions

Validates whether the agent delivers accurate and helpful responses

Provides measurable, actionable insights for continuous improvement

What should you recommend?

- A. Track resolution, deflection, and accuracy by using dashboards and use scripts to ensure consistent responses.
- B. Perform load testing to validate how the agent scales under a high chat volume.
- C. Review historical tickets to find agents that have the shortest resolution times.
- D. Measure uptime and page load times.

Answer

A

Correct answer: A. Track resolution, deflection, and accuracy by using dashboards and use scripts to ensure consistent responses.

Why A is correct

The three requirements (evaluate effectiveness during active sessions, validate accuracy and helpfulness, and produce measurable, actionable insights for continuous improvement) align with the Copilot Studio analytics + scripted-test pattern. The Analytics dashboard tracks resolution rate, escalation rate, and deflection in near-real-time, and the Generated answer rate and quality section measures whether responses are accurate. Scripted tests, run through the Test panel or the Copilot Studio Kit's automated testing pipeline, drive consistency by validating the same prompts produce the expected behavior across releases. Together, dashboards plus scripted regression tests deliver measurable signals and a continuous-improvement loop.

Why the other options are wrong

B. Perform load testing to validate how the agent scales under a high chat volume. Load testing measures throughput and latency, not accuracy or helpfulness.

C. Review historical tickets to find agents that have the shortest resolution times. Historical ticket review evaluates human reps, not the AI agent's effectiveness during active sessions.

D. Measure uptime and page load times. These are infrastructure metrics, not agent quality metrics.

Microsoft Learn references

- Analyze conversational agent effectiveness - <https://learn.microsoft.com/microsoft-copilot-studio/analytics-improve-agent-effectiveness>
- Deflection overview - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/deflection-overview>
- Automate testing and deployment of agents using pipelines in Power Platform - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/kit-automate-test-deploy>
- Improve your Copilot Studio projects - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/improve-overview>

Topic 3 - Question 24

Multiple choice | Domain: Deploy

topic-3/q24

A company uses multiple Microsoft Copilot Studio agents across different channels.

You need to recommend a monitoring solution that provides comprehensive telemetry data and performance insights for the agents.

What should you include in the recommendation?

- A. Application Insights
- B. Azure Advisor
- C. Azure DevOps

D. Microsoft Dynamics 365 Customer Voice**Answer****A****Correct answer: A. Application Insights****Why A is correct**

Application Insights is the Azure service Microsoft positions for comprehensive Copilot Studio agent telemetry. Microsoft Learn explains that Copilot Studio sends conversations, events, exceptions, latency, and tool usage to Application Insights and that the Copilot Studio Dashboard workbook surfaces single-pane operational metrics across these signals. For agents running across multiple channels, Application Insights captures the channel dimension on each event, allowing cross-channel performance comparisons through Azure Workbooks, Log Analytics queries, and dashboards.

Why the other options are wrong

B. Azure Advisor issues best-practice recommendations on Azure resources. It is not a telemetry surface for Copilot Studio agents.

C. Azure DevOps provides source control and pipelines for ALM. It does not collect agent telemetry.

D. Microsoft Dynamics 365 Customer Voice is a survey product. It captures user feedback, not comprehensive agent telemetry.

Microsoft Learn references

- Capture telemetry with Application Insights (Copilot Studio) - <https://learn.microsoft.com/microsoft-copilot-studio/advanced-bot-framework-composer-capture-telemetry>
- Configure telemetry to capture bot data - <https://learn.microsoft.com/microsoft-copilot-studio/advanced-bot-framework-composer-capture-telemetry>
- Combine web and agent analytics - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/pass-web-analytics-tracking-id-from-web-page>

Topic 3 - Question 25*Multi-select | Domain: Deploy**topic-3/q25*

A company has an AI solution built by using Microsoft Copilot Studio and Power Platform. The solution is used by the company's sales, marketing, and customer service teams.

You are performing a return on AI investment (ROAI) analysis to evaluate the impact of the solution.

You need to identify which measurable business drivers to include in the analysis.

Which two business drivers should you identify? Each correct answer presents part of the solution.

- A.** the reduced average case resolution time
- B.** market capitalization
- C.** economic market predictability
- D.** increased employee productivity
- E.** brand awareness

Select all that apply.

Answer

A, D

Correct answer: A. the reduced average case resolution time and D. increased employee productivity

Why A and D are correct

Return on AI investment analysis compares the cost of running the AI solution against measurable business outcomes that the solution influences. Microsoft's Copilot Studio guidance and Viva Insights ROI calculations both center on operational outcomes that the agent can directly affect.

Reduced average case resolution time is a direct, measurable outcome of an AI-assisted customer service flow. The Copilot Studio Savings calculator and Viva Insights agent-assisted impact calculations translate time saved into dollar value, which is exactly what ROAI measures.

Increased employee productivity is the second pillar Microsoft emphasizes. Generative AI in sales, marketing, and customer service offloads routine work and lets people spend time on higher-value activities. Productivity gains, measured in hours saved per employee per week or in higher throughput per team, feed directly into ROI calculations.

Why the other options are wrong

B. Market capitalization reflects investor sentiment about the entire company. The AI solution cannot be isolated as the cause.

C. Economic market predictability is a macroeconomic concept. It is not a measurable business driver attributable to a Copilot Studio solution.

E. Brand awareness depends on marketing reach and reputation, which is not a measurable business driver tied to Copilot Studio adoption.

Microsoft Learn references

- Analyze time and cost savings for agents (Copilot Studio Savings calculator) - <https://learn.microsoft.com/microsoft-copilot-studio/analytics-cost-savings>
- View aggregated analytics in Viva Insights (ROI calculations) - <https://learn.microsoft.com/microsoft-copilot-studio/analytics-viva-insights>
- Measure agent value using Copilot Studio Kit - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/kit-agent-value-dashboard>

Topic 3 - Question 26

Multiple choice | Domain: Deploy

topic-3/q26

A company has a Microsoft Foundry agent that summarizes customer feedback and recommends products to customers. The agent references data from multiple knowledge sources.

Users report that the agent response time is slow.

Telemetry data shows that the agent frequently reaches its token usage limit.

You need to recommend a solution to reduce token usage without degrading the quality of the generated responses.

What should you recommend?

- A. Chunk documents during indexing.
- B. Reduce the number of knowledge sources used by the agent.
- C. Reconfigure the prompts to limit the amount of retrieved content from the knowledge sources.
- D. Lower the maximum token usage limit for the responses.

Answer

C

Correct answer: C. Reconfigure the prompts to limit the amount of retrieved content from the knowledge sources.

Why C is correct

Token usage on a RAG-style agent is dominated by the retrieved content stuffed into the prompt context for grounding. When telemetry shows the agent is hitting its token cap, the highest-leverage fix is to reduce the volume of retrieved content per prompt while keeping the same knowledge sources in scope. Microsoft's optimization guidance points to refining prompts and retrieval to reduce ambiguity and tighten the context window. Limiting retrieved chunks per call drops token consumption while preserving response quality, because the agent still consults all the knowledge sources, just less of each per turn.

Why the other options are wrong

A. Chunk documents during indexing. Chunking is good practice for retrieval quality, but it does not by itself reduce the number of chunks pulled into the prompt. Without limiting how many chunks the prompt retrieves, token usage stays high.

B. Reduce the number of knowledge sources used by the agent. Cutting knowledge sources risks degrading response quality, which violates the "without degrading quality" requirement.

D. Lower the maximum token usage limit for the responses. Capping the response length might truncate the answer, which degrades quality. It does not address the input-side token consumption from retrieval.

Microsoft Learn references

- Optimize prompts and instructions (Copilot Studio) - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/optimize-prompts-overview>
- Knowledge sources summary (Copilot Studio) - <https://learn.microsoft.com/microsoft-copilot-studio/knowledge-copilot-studio>
- Observability in generative AI (token usage and latency) - <https://learn.microsoft.com/azure/foundry/concepts/observability>

Topic 3 - Question 27

Hotspot | Domain: Deploy

topic-3/q27

HOTSPOT -

A company has Microsoft Power Platform development, staging, and production environments. Each environment has its own Microsoft Dataverse tables and Azure AI Search index.

You are designing an application lifecycle management (AIM) process to deploy a Microsoft Copilot Studio agent between the environments.

The company has a Copilot Studio agent named Agent1 in development. Agent1 uses the following

grounding data sources:

- A Dataverse table named CustomerOrders
- An Azure AI Search index named customer-knowledge

You need to deploy Agent1 to production. The solution must ensure that the agent uses the production grounding data sources, minimizes downtime, and handles credentials and endpoints securely.

What should you include in the deployment package solution, and what should you reconfigure after the deployment? To answer, select the appropriate options in the answer area.

Answer area

- Include in the deployment package solution:

- ☐ Agent1 only
- ☐ The data sources only
- ☐ Agent1 and the data source connections
- ☐ Agent1 and references to the data sources
- ☐ Agent1, the data sources, and the data source connections

- Reconfigure after the deployment:

- ☐ The Dataverse connection only
- ☐ The Azure AI Search connection only
- ☐ The Dataverse and Azure AI Search connections
- ☐ The Agent1 configuration
- ☐ The environment variables

Answer

Include = Agent1 and references to the data sources; Reconfigure = The environment variables

Correct answer: - Include in the deployment package solution: Agent1 and references to the data sources - Reconfigure after the deployment: The environment variables

Why these selections are correct

Agent1 and references to the data sources: Each environment has its own Dataverse tables and its own Azure AI Search index, so the production data sources are not the same artifacts as the development data sources. The deployment package should contain the agent and references to the data sources (typically through environment variables), not the data themselves and not the data source connections, which are environment-specific. Microsoft's ALM golden rules recommend exporting agents and components as managed solutions and using environment variables for settings that change across environments.

Reconfigure the environment variables: After import, the environment variables that point at the Dataverse table and the Azure AI Search endpoint need to be set to the production values. The Dataverse and Azure AI Search connections also need to authenticate, but Microsoft's guidance is to drive that through environment variables and connection references that are bound after import. Reconfiguring environment variables minimizes downtime, handles credentials securely, and keeps the package portable across environments.

Microsoft Learn references

- Establish an application lifecycle management strategy (Copilot Studio) - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/alm>
- Export and import agents using solutions - <https://learn.microsoft.com/microsoft-copilot-studio/authoring-solutions-import-export>
- Environment variables in solutions - <https://learn.microsoft.com/power-apps/maker/data-platform/environmentvariables>
- Connection references -

<https://learn.microsoft.com/power-apps/maker/data-platform/create-connection-reference>

Topic 3 - Question 28

Multiple choice | Domain: Deploy

topic-3/q28

A company has a Microsoft Copilot Studio agent that uses generative AI to assist Microsoft Dynamics 365 Customer Service representatives.

The agent currently exhibits a low resolution rate and a high escalation rate.

You need to identify the issue.

What should you use?

- A. the Agent dashboard of Dynamics 365 Customer Service historical analytics
- B. the Insights tab from the Search & intelligence settings of the Microsoft 365 admin center
- C. the Copilot hub in the Power Platform admin center
- D. the Analytics tab in Copilot Studio

Answer

D

Correct answer: D. the Analytics tab in Copilot Studio

Why D is correct

Resolution rate and escalation rate are core Copilot Studio analytics metrics, surfaced on the agent's Analytics tab. Microsoft Learn details the metrics: Resolution Rate is "the percentage of engaged sessions that are resolved," Escalation Rate is "the percentage of engaged sessions that are escalated," and the Analytics page exposes both alongside Escalation Rate Drivers, which highlights which topics generate the most escalations. That is the right place to identify the issue behind a low resolution rate and a high escalation rate.

Why the other options are wrong

A. The Agent dashboard of Dynamics 365 Customer Service historical analytics focuses on human representative performance and overall conversation handling, not the Copilot Studio agent's topic-level resolution and escalation drivers.

B. The Insights tab from Search & intelligence settings of Microsoft 365 admin center covers Microsoft 365 Search analytics. It is unrelated to Copilot Studio agent performance.

C. The Copilot hub in the Power Platform admin center provides tenant-level Copilot governance views. It does not surface per-agent topic-level resolution or escalation drivers the way the Analytics tab does.

Microsoft Learn references

- Analyze conversational agent effectiveness - <https://learn.microsoft.com/microsoft-copilot-studio/analytics-improve-agent-effectiveness>
- Deflection overview (Resolution and Escalation rates) - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/deflection-overview>
- Topic escalation analysis - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/deflection-topic-escalation-analysis>
- Measure and improve agent performance with KPIs and analytics - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/analytics>

Topic 3 - Question 29

Multiple choice | Domain: Deploy

topic-3/q29

A company has a Microsoft Foundry generative AI model.

You need to evaluate the model's output to measure the overall quality and coherence of generated responses. The evaluation must use GPT-4o as a judge and return a numeric score for each output.

Which type of metric should you use?

- A. AI quality (NLP)
- B. risk and safety
- C. Groundedness
- D. AI quality (AI assisted)

Answer

D

Correct answer: D. AI quality (AI assisted)

Why D is correct

The question describes overall quality and coherence of generated responses, evaluated with GPT-4o as a judge, returning a numeric score per output. That is the AI-assisted quality category in Microsoft Foundry's evaluator framework. Microsoft Learn lists the AI-assisted quality evaluators as Coherence, Fluency, Relevance, Groundedness, Similarity, Retrieval, Response Completeness, and similar measures, all of which use an LLM judge (GPT-4.1, GPT-4o, or an o-series reasoning model) to score outputs on a numeric scale. Coherence specifically "measures the logical and orderly presentation of ideas in a response."

Why the other options are wrong

A. AI quality (NLP) uses traditional NLP overlap metrics such as F1, BLEU, GLEU, ROUGE, METEOR. They compare model output to a ground-truth reference using token overlap and do not require an LLM judge.

B. Risk and safety evaluates harmful content (violence, hate, sexual, self-harm, jailbreak susceptibility, protected materials). It is not a measure of overall quality and coherence.

C. Groundedness is one specific AI-assisted metric (whether the response is supported by the source data). It does not by itself measure overall quality and coherence; it is a component of the broader AI-assisted quality category.

Microsoft Learn references

- Built-in evaluators reference (Microsoft Foundry) - <https://learn.microsoft.com/azure/foundry/concepts/built-in-evaluators>
- General purpose evaluators (Coherence and Fluency) - <https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/general-purpose-evaluators>
- Evaluate your AI agents (model support for AI-assisted evaluators) - <https://learn.microsoft.com/azure/foundry-classic/how-to/develop/agent-evaluate-sdk>
- Observability in generative AI - <https://learn.microsoft.com/azure/foundry/concepts/observability>

Topic 3 - Question 30

Hotspot | Domain: Deploy

topic-3/q30

HOTSPOT -

You use Microsoft Copilot Studio analytics to analyze the performance of a deployed Copilot Studio agent.

You need to identify which performance metrics to use to measure the following:

- The percentage of engaged sessions that are escalated to a live customer service representative
- The number of agent queries that cause a knowledge source error

What should you identify for each requirement? To answer, select the appropriate options in the answer area.

Answer area

- **The percentage of engaged sessions that are escalated to a live customer service representative:**

- ☐ Answer quality
- ☐ Customer Satisfaction (CSAT) score
- ☐ Engagement rate
- ☐ Escalation rate

- **The number of agent queries that cause a knowledge source error:**

- ☐ Answer quality
- ☐ Engagement rate
- ☐ Escalation rate

Answer

Escalated sessions = Escalation rate; Knowledge source error = Answer quality

Correct answer: - Percentage of engaged sessions escalated to a live representative: Escalation rate - Number of agent queries that cause a knowledge source error: Answer quality

Why these selections are correct

Escalation rate is defined in Microsoft Learn as "the percentage of engaged sessions that are escalated. An escalated session is an engaged session escalated to a human representative." That is the literal metric for the first requirement.

Answer quality (part of the Generated answer rate and quality section) is where Copilot Studio surfaces knowledge-source errors. Microsoft Learn states the side panel shows "Errors... the percentage of queries that resulted in a knowledge-related error for each knowledge source type (for example, SharePoint)." The Answer quality breakdown is the metric that exposes knowledge source errors. CSAT score measures customer satisfaction, and engagement rate measures whether sessions engaged a topic, neither of which is the knowledge-error count.

Microsoft Learn references

- Deflection overview (Escalation Rate definition) - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/deflection-overview>
- Analyze conversational agent effectiveness (Generated answer rate and quality, knowledge source errors) - <https://learn.microsoft.com/microsoft-copilot-studio/analytics-improve-agent-effectiveness>
- Analytics overview (Copilot Studio) - <https://learn.microsoft.com/microsoft-copilot-studio/analytics-overview>

Topic 3 - Question 31

Multiple choice | Domain: Deploy

topic-3/q31

A company plans to deploy a Microsoft Foundry agent.

You need to recommend an application lifecycle management (ALM) process to ensure that the agent evaluates against baseline accuracy metrics before being deployed.

What should you recommend?

- A. Configure GitHub Actions for new agent versions.
- B. Deploy each new agent version directly to production.
- C. Use Observability in Foundry Control Plane with evaluation and drift monitoring.
- D. Enable Application Insights and use Azure Monitor.

Answer

C

Correct answer: C. Use Observability in Foundry Control Plane with evaluation and drift monitoring.

Why C is correct

The question requires baseline accuracy evaluation before each agent version is deployed. Microsoft Foundry Control Plane brings observability, evaluation, and drift monitoring into a single ALM-aligned surface for Foundry agents. Microsoft Learn describes the Control Plane: "Continuously evaluate agent performance, quality, and risk dimensions. Risk dimensions might include task adherence, intent resolution, tool call success, groundedness, sensitive data leakage, and exposure to jailbreak and cross-domain prompt injection attacks." The Agent Monitoring Dashboard supports continuous evaluation and integrates results into the deployment flow, which is the prescribed path for evaluating agents against baseline metrics before promotion.

Why the other options are wrong

A. Configure GitHub Actions for new agent versions. GitHub Actions can orchestrate deployments, but it does not by itself perform Foundry-specific evaluation against baseline accuracy metrics.

B. Deploy each new agent version directly to production. This skips the evaluation gate entirely.

D. Enable Application Insights and use Azure Monitor. Application Insights captures runtime telemetry. It does not run pre-deployment evaluations against baseline accuracy metrics the way Foundry observability and evaluation do.

Microsoft Learn references

- What is Microsoft Foundry Control Plane? - <https://learn.microsoft.com/azure/foundry/control-plane/overview>
- Monitor agents with the Agent Monitoring Dashboard - <https://learn.microsoft.com/azure/foundry/observability/how-to/how-to-monitor-agents-dashboard>
- Observability in generative AI (evaluation, monitoring, tracing) - <https://learn.microsoft.com/azure/foundry/concepts/observability>
- Built-in evaluators reference - <https://learn.microsoft.com/azure/foundry/concepts/built-in-evaluators>

Topic 3 - Question 32

Multiple choice | Domain: Deploy

topic-3/q32

You need to recommend a security solution for agents in a Microsoft Power Platform environment.

The agents must use only approved connectors and services. The solution must prevent the agents from accessing sensitive data.

What should you recommend?

- A. Configure Azure Monitor to capture connector activity logs.
- B. Enable a Microsoft Dataverse audit.
- C. Deploy data loss prevention (DLP) policies in Power Platform.
- D. Enable customer-managed keys in Microsoft Dataverse.

Answer

C

Correct answer: C. Deploy data loss prevention (DLP) policies in Power Platform.

Why C is correct

DLP policies in Power Platform classify connectors as Business, Non-business, or Blocked, and prevent agents and apps from combining connectors across categories. Microsoft Learn states: "Data policies, including data loss prevention (DLP) policies, act as guardrails to help prevent users from unintentionally exposing organizational data and to protect information security in the tenant. Data policies enforce rules for which connectors are enabled for each environment, and which connectors can be used together." That maps directly to "agents must use only approved connectors" and "prevent the agents from accessing sensitive data."

Why the other options are wrong

- A. Configure Azure Monitor to capture connector activity logs.** Logging captures evidence after the fact. It does not enforce which connectors agents can use.
- B. Enable a Microsoft Dataverse audit.** Auditing tracks data access in Dataverse but does not control which connectors an agent can call or block sensitive data flow.
- D. Enable customer-managed keys in Microsoft Dataverse.** CMK encrypts data at rest with a customer-managed key. It does not constrain connector usage.

Microsoft Learn references

- Data policies (Power Platform) - <https://learn.microsoft.com/power-platform/admin/wp-data-loss-prevention>
- Implement a data policy strategy - <https://learn.microsoft.com/power-platform/guidance/adoption/dlp-strategy>
- Power Platform DLP policies on Day 1 - <https://learn.microsoft.com/microsoft-365/community/power-platform-dlp-policies-you-should-be-considering-on-day-1>
- Advanced connector policies (preview) - <https://learn.microsoft.com/power-platform/admin/advanced-connector-policies>

Topic 3 - Question 34

Drag and drop | Domain: Deploy

topic-3/q34

DRAG DROP -

A company has a cloud-based AI solution that uses Azure OpenAI models.

You need to design a monitoring solution that meets the following requirements:

- Monitors performance metrics and operational health for the models
- Monitors AI apps and agents for compliance
- Uses Azure-native capabilities
- Minimizes development effort

What should you use for each requirement? To answer, drag the appropriate options to the correct requirements. Each option may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

Drag sources

- Azure API Management
- Azure Policy
- Microsoft Defender
- Azure Monitor
- Azure Stream Analytics
- Microsoft Purview

Drop targets

1. Monitors AI apps and agents for compliance: - drop here
2. Monitors performance metrics and operational health: - drop here

Answer

Compliance = Microsoft Purview; Performance/operational health = Azure Monitor

Correct answer: - Monitors AI apps and agents for compliance: Microsoft Purview - Monitors performance metrics and operational health: Azure Monitor

Why these selections are correct

Microsoft Purview is Microsoft's compliance and governance platform for AI. Microsoft Learn describes Purview's native integration with Microsoft Foundry and Azure OpenAI: subscription-level toggling sends prompts and responses to Purview for audit, communication compliance, DLP, DSPM for AI, and insider risk monitoring. That matches "monitor AI apps and agents for compliance" and uses Azure-native capabilities with minimal development effort.

Azure Monitor collects metrics, logs, and traces from Azure OpenAI and Foundry resources. It is the Azure-native observability platform for performance and operational health: latency, throughput, request volume, errors, dependencies. Application Insights (a feature of Azure Monitor) feeds the Foundry Agent Monitoring Dashboard.

Defender protects against threats and discovers AI workloads but is not the operational performance monitor. Azure Policy enforces configuration, API Management front-doors APIs, and Stream Analytics processes streaming data, none of which fits either requirement.

Microsoft Learn references

- Use Microsoft Purview to manage data security and compliance for Microsoft Foundry - <https://learn.microsoft.com/purview/ai-azure-foundry>
- Microsoft Purview data security and compliance protections for generative AI apps - <https://learn.microsoft.com/purview/ai-microsoft-purview>
- Monitor agents with the Agent Monitoring Dashboard (Azure Monitor / App Insights) - <https://learn.microsoft.com/azure/foundry/observability/how-to/how-to-monitor-agents-dashboard>
- Azure Monitor overview - <https://learn.microsoft.com/azure/azure-monitor/overview>

Topic 3 - Question 35

Multi-select | Domain: Deploy

topic-3/q35

A company has Microsoft Copilot Studio agents.

The company plans to deploy custom connectors across development, test and production environments.

You need to design an application lifecycle management (ALM) process to ensure consistency and prevent direct editing in production.

Which two actions should you include in the design? Each correct answer presents part of the solution.

- A. Deploy managed solutions to production.
- B. Deploy unmanaged solutions to production.
- C. Manually rebuild the agents in each environment.
- D. Move the agents between the environments by using data export and import.
- E. Include agents and connectors in a solution.

Select all that apply.

Answer

A, E

Correct answer: A. Deploy managed solutions to production and E. Include agents and connectors in a solution.

Why A and E are correct

Power Platform's ALM golden rules call for packaging components in a solution and deploying as managed in non-development environments. Microsoft Learn states: "Export and deploy solutions as managed, unless setting up a development environment." Managed solutions in production prevent direct edits, enforce consistency across environments, and provide clean uninstall and upgrade behavior. Including the agents and custom connectors in a single solution treats them as one ALM unit so the connector and agent versions stay in lockstep.

Why the other options are wrong

B. Deploy unmanaged solutions to production. Unmanaged solutions allow direct edits in production, which is exactly what the requirement forbids.

C. Manually rebuild the agents in each environment. Manual rebuilds break consistency and traceability and make connector versions drift across environments.

D. Move the agents between the environments by using data export and import. Data export and import bypasses solution-aware ALM, prevents dependency tracking, and undermines version control.

Microsoft Learn references

- Establish an application lifecycle management strategy (Copilot Studio, ALM golden rules) - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/alm>
- Organize your solutions - <https://learn.microsoft.com/power-platform/alm/organize-solutions>
- ALM basics with Microsoft Power Platform - <https://learn.microsoft.com/power-platform/alm/basics-alm>
- Pipelines in Power Platform - <https://learn.microsoft.com/power-platform/alm/pipelines>

Topic 3 - Question 36

Multiple choice | Domain: Deploy

topic-3/q36

A company uses multiple Microsoft Copilot Studio agents across different channels.

You need to recommend a monitoring solution that provides comprehensive telemetry data and performance insights for the agents.

What should you include in the recommendation?

- A. Application Insights
- B. Microsoft Dynamics 365 Customer Voice
- C. Log Analytics
- D. Microsoft Purview

Answer

A

Correct answer: A. Application Insights

Why A is correct

Application Insights is the integrated telemetry and monitoring service Microsoft positions for Copilot Studio agents. Microsoft Learn shows that Copilot Studio sends conversations, exceptions, latency, tool usage, and custom events into Application Insights. The Copilot Studio Dashboard workbook in Application Insights surfaces total conversations, latency, exceptions, tool usage, and topic analytics in a single view, which provides comprehensive telemetry data and performance insights across multiple channels.

Why the other options are wrong

B. Microsoft Dynamics 365 Customer Voice is a survey product. It captures user feedback, not comprehensive agent telemetry.

C. Log Analytics is the query layer that sits behind Application Insights. The Application Insights workspace is the entry point Microsoft documents for Copilot Studio telemetry. Log Analytics alone is not the labeled telemetry destination for Copilot Studio agents.

D. Microsoft Purview governs and protects data and audits activity for compliance. It is not the comprehensive performance telemetry surface for Copilot Studio agents.

Microsoft Learn references

- Capture telemetry with Application Insights (Copilot Studio) - <https://learn.microsoft.com/microsoft-copilot-studio/advanced-bot-framework-composer-capture-telemetry>
- Configure telemetry to capture bot data - <https://learn.microsoft.com/microsoft-copilot-studio/advanced-bot-framework-composer-capture-telemetry>
- Combine web and agent analytics -

<https://learn.microsoft.com/microsoft-copilot-studio/guidance/pass-web-analytics-tracking-id-from-web-page>

Topic 3 - Question 37

Multiple choice | Domain: Deploy

topic-3/q37

A company has an AI solution named Solution1 that is deployed to the production environment. Solution1 uses an Azure OpenAI model to generate marketing emails for existing customers.

During an internal review, you identify that Solution1 creates different emails depending on the customers' traits.

You need to recommend a strategy to mitigate the bias. The strategy must adhere to Microsoft responsible AI principles.

What should you recommend?

- A. Modify Solution1 to randomly generate emails for different traits.
- B. Modify the system instructions of Solution1.
- C. Retrain the model by using a larger dataset.
- D. Modify the contents of the training dataset.

Answer

D

Correct answer: D. Modify the contents of the training dataset.

Why D is correct

The bias is producing different emails based on customer traits, which is a fairness problem rooted in the data the model learned from. Microsoft's Responsible AI principles include fairness, and the practical mitigation guidance for biased outputs starts with the dataset: identify under-represented or skewed groups in the training data, then rebalance, augment, or curate the dataset to reduce bias before retraining. Microsoft's Responsible AI dashboard, Fairlearn, and the Responsible AI scorecard all support this data-centered fairness assessment and mitigation flow.

Why the other options are wrong

A. Modify Solution1 to randomly generate emails for different traits. Randomization masks the issue but does not address the underlying bias and could degrade quality for everyone.

B. Modify the system instructions of Solution1. Prompt instructions can influence behavior, but they do not fix bias baked into the underlying model. The model still reflects the training data distribution.

C. Retrain the model by using a larger dataset. Larger does not equal less biased. Adding more skewed data can amplify existing bias. The correct fairness intervention is to modify the dataset's contents to address the imbalance, not just its size.

Microsoft Learn references

- Microsoft Responsible AI standard and principles - <https://www.microsoft.com/ai/responsible-ai>
- Assess AI systems by using the Responsible AI dashboard - <https://learn.microsoft.com/azure/machine-learning/concept-responsible-ai-dashboard>
- What is Responsible AI? - <https://learn.microsoft.com/azure/machine-learning/concept-responsible-ai>
- Fairness assessment in Responsible AI - <https://learn.microsoft.com/azure/machine-learning/concept-fairness-ml>

Topic 3 - Question 39

Multi-select | Domain: Deploy | Case study: Contoso
topic-3/q39

Which two components for the custom AI agent should you include in the application lifecycle management (ALM) process? Each correct answer presents part of the solution.

- A. an Azure package
- B. a ZIP package
- C. a Microsoft Power Platform solution
- D. a Cloud Scale Unit (CSU) package
- E. an X++model

Select all that apply.

Answer

B, C

Correct answer: B. a ZIP package and C. a Microsoft Power Platform solution

Why B and C are correct

The Contoso custom AI agent is a low-code Copilot Studio agent built on Dynamics 365 Supply Chain Management data. Microsoft Learn defines Power Platform solutions as the ALM container for Copilot Studio agents and connectors: "When you create an agent in Copilot Studio, the agent is created within a Power Platform solution." Solutions are exported as compressed ZIP packages for transport between environments. That is the supported ALM artifact for the Contoso agent.

In Microsoft Learn: "When you export a solution from Dataverse, it's exported as a single compressed file." That compressed file is the ZIP package, and the underlying construct is the Power Platform solution. Both items together are the components included in the ALM process.

Why the other options are wrong

A. an Azure package is not a Power Platform ALM artifact. Azure deployments use ARM, Bicep, or container artifacts, none of which apply to a Copilot Studio agent.

D. a Cloud Scale Unit (CSU) package belongs to Dynamics 365 Finance and Operations apps, not Copilot Studio agents.

E. an X++ model is a Finance and Operations developer artifact. It is not used to package Copilot Studio agents.

Microsoft Learn references

- Solution management in Microsoft Copilot Studio - <https://learn.microsoft.com/microsoft-copilot-studio/authoring-solutions-overview>
- Establish an application lifecycle management strategy (Copilot Studio) - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/alm>
- Modernize applications with Power Platform (solutions exported as compressed file) - <https://learn.microsoft.com/power-platform/guidance/adoption/application-modernization>
- Export and import agents using solutions - <https://learn.microsoft.com/microsoft-copilot-studio/authoring-solutions-import-export>

Topic 3 - Question 40

Drag and drop | Domain: Deploy | Case study: Contoso

topic-3/q40

DRAG DROP -

Which tools should you recommend to assist the CISO and the CIO with their specific responsibilities? To answer, drag the appropriate tools to the correct executives. Each tool may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

Drag sources

- Azure Blob Storage
- Azure Resource Graph Explorer
- Copilot Studio
- Microsoft Purview

Drop targets

1. CISO: - drop here
2. CIO: - drop here

Answer

CISO = Azure Resource Graph Explorer; CIO = Microsoft Purview

Correct answer: - CISO: Azure Resource Graph Explorer - CIO: Microsoft Purview

Why these selections are correct

CISO is responsible for "discover and inventory AI resources for auditing." Azure Resource Graph Explorer queries every Azure resource in the tenant across subscriptions, including Foundry projects, Azure OpenAI accounts, AI Search, Storage, and related grounding resources. Microsoft Learn positions Resource Graph as the inventory and exploration tool for resources at scale. Microsoft Purview, Copilot Studio, and Azure Blob Storage do not surface a tenant-wide inventory of AI resources the way Resource Graph does, which makes it the right tool for the CISO's audit mandate.

CIO is responsible for ensuring "appropriate security labels are assigned to the data used by the AI agents." Microsoft Purview is the platform that defines, applies, and audits sensitivity labels across Microsoft 365 and Azure data sources, including data that grounds AI agents. Microsoft Learn details Purview's information protection, sensitivity labels, and DSPM for AI as the controls for labeling and governing AI-relevant data.

Microsoft Learn references

- What is Azure Resource Graph? - <https://learn.microsoft.com/azure/governance/resource-graph/overview>
- Microsoft Purview overview - <https://learn.microsoft.com/purview/purview>
- Sensitivity labels in Microsoft Purview - <https://learn.microsoft.com/purview/sensitivity-labels>
- Microsoft Purview data security and compliance protections for generative AI apps - <https://learn.microsoft.com/purview/ai-microsoft-purview>

Topic 3 - Question 41

Multiple choice | Domain: Deploy | Case study: Contoso
topic-3/q41

What should you recommend to assist the CEO with their specific responsibilities?

- A. the Microsoft Service Trust Portal
- B. Microsoft Foundry Tools
- C. Microsoft Purview
- D. the Responsible AI dashboard
- E. Compliance Center

Answer

D

Correct answer: D. the Responsible AI dashboard

Why D is correct

The CEO's responsibility in the Contoso case study is to ensure all solutions adhere to industry-standard responsible AI practices, with quarterly assessment that "must verify reliability, interpretability, fairness, and compliance." Microsoft Learn describes the Responsible AI dashboard as a single interface that brings together model performance and fairness assessment, data exploration, model interpretability, error analysis, counterfactual analysis, and causal inference. The dashboard explicitly supports "model reliability, interpretability, fairness, and compliance," which matches every term in the requirement, and the Responsible AI scorecard packages these results into a shareable PDF for stakeholders and audit reviews.

Why the other options are wrong

A. The Microsoft Service Trust Portal publishes Microsoft's audit reports, certifications, and compliance documents about Microsoft cloud services. It does not assess Contoso's own AI solutions for reliability, interpretability, and fairness.

B. Microsoft Foundry Tools is a generic label and is not the named tool that produces a quarterly responsible AI assessment for the CEO.

C. Microsoft Purview governs data and audits AI interactions for compliance. It does not provide model interpretability and fairness assessments the way the Responsible AI dashboard does.

E. Compliance Center (the legacy Microsoft 365 Compliance Center) is a tenant compliance management portal. It does not perform AI model assessments for reliability and fairness.

Microsoft Learn references

- Assess AI systems by using the Responsible AI dashboard - <https://learn.microsoft.com/azure/machine-learning/concept-responsible-ai-dashboard>
- What is Responsible AI? - <https://learn.microsoft.com/azure/machine-learning/concept-responsible-ai>
- Generate Responsible AI insights in the studio UI - <https://learn.microsoft.com/azure/machine-learning/how-to-responsible-ai-insights-ui>
- Responsible AI scorecard - <https://learn.microsoft.com/azure/machine-learning/concept-responsible-ai-scorecard>

Topic 3 - Question 42

Drag and drop | Domain: Deploy | Case study: Contoso

topic-3/q42

DRAG DROP -

Which Copilot Studio analytics metrics should you recommend to assist the executives with their specific responsibilities? To answer, drag the appropriate metrics to the correct executives. Each metric may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

Drag sources

- Effectiveness
- Tool use
- Satisfaction
- Use

Drop targets

1. CFO concerns about Copilot Studio credit usage: - drop here
2. CTO concerns about poor feedback on AI agent responses: - drop here

Answer

CFO = Use; CTO = Satisfaction

Correct answer: - CFO concerns about Copilot Studio credit usage: Use - CTO concerns about poor feedback on AI agent responses: Satisfaction

Why these selections are correct

CFO and the Use metric: The case study states the CFO "wants to identify how many interactions with the AI agents are abandoned on a given day as compared to resolved conversations. Too many abandoned sessions might indicate that Copilot Studio credits are being used inefficiently by end users." The Copilot Studio Analytics page exposes a Use section that surfaces session and run volume and outcomes, and the Conversation outcomes view inside Use shows resolved versus abandoned counts. That is the metric category for tracking credit usage efficiency tied to session outcomes.

CTO and the Satisfaction metric: The case study says the CTO "wants to track user feedback on the quality of the AI agent responses during user interactions with the agents. Consistently poor feedback will trigger an escalated reengineering discussion." The Analytics page's Satisfaction (Effectiveness) section captures CSAT, sentiment, reactions (thumbs up/down), and user comments. That is the surface that reflects user feedback on response quality.

Effectiveness in Analytics aggregates user reactions and survey results, so it overlaps with Satisfaction, but the named Satisfaction metric is the cleanest match for "track user feedback on the quality of agent responses." Tool use measures tool invocations and is unrelated to either question.

Microsoft Learn references

- Analytics overview (Use, Satisfaction, Effectiveness sections) - <https://learn.microsoft.com/microsoft-copilot-studio/analytics-overview>
- Analyze conversational agent effectiveness (Conversation outcomes, Effectiveness, Satisfaction) - <https://learn.microsoft.com/microsoft-copilot-studio/analytics-improve-agent-effectiveness>
- Deflection overview (Resolution, Escalation, and Abandon rates) - <https://learn.microsoft.com/microsoft-copilot-studio/guidance/deflection-overview>
- Analyze time and cost savings for agents - <https://learn.microsoft.com/microsoft-copilot-studio/analytics-cost-savings>