AB-100 Practice Questions - Plan

Agentic AI Business Solutions Architect | Plan domain

Generated 2026-05-05

20 questions, with answers and explanations

Plan AI-powered business solutions (25-30% of the AB-100 exam).

Question types

Drag and drop: 1

Hotspot: 3

Multi-select: 2

Multiple choice: 14

Topic coverage

Topic 1: 14

Topic 2: 4

Topic 3: 2

Source: practicetest/extracted/. Each answer references Microsoft Learn. Use this PDF for offline review; the interactive version lives at app/index.html.

Topic 1

14 questions

Topic 1 - Question 1

Hotspot | Domain: Plan | Case study: Fabrikam

topic-1/q01

Which framework should you use to meet the AI agent requirements for the sales cycle enablement? To answer, select the appropriate options in the answer area.

Answer area

- For Microsoft Copilot Studio best practices:

[ ] the ALM Accelerator for Microsoft Power Platform
[ ] Microsoft Cloud Adoption Framework for Azure
[ ] Microsoft Power Platform Well-Architected framework
[ ] Success by Design

- For conversational user experiences:

[ ] the ALM Accelerator for Microsoft Power Platform
[ ] Microsoft Cloud Adoption Framework for Azure
[ ] Microsoft Power Platform Well-Architected framework
[ ] Success by Design

Answer

Microsoft Power Platform Well-Architected framework; Success by Design

Correct answer: - For Microsoft Copilot Studio best practices: **Microsoft Power Platform Well-Architected framework** - For conversational user experiences: **Success by Design**

Why these answers are correct

Fabrikam's case study calls out two distinct sets of guidance the team needs:

1. Microsoft-recommended best practices for the Copilot Studio agent itself (the low-code agent built on Dataverse).
2. Microsoft recommendations for the agent's conversational user experience (the headset-driven AI agent for the sales executives).

Microsoft Power Platform Well-Architected framework is the published Microsoft framework for designing, planning, and implementing modern application workloads with Power Platform, including Copilot Studio agents. It is organized around five pillars (Reliability, Security, Operational Excellence, Performance Efficiency, Experience Optimization) and is what Microsoft Learn directs makers to when they want best-practice design guidance for Copilot Studio.

Success by Design includes Microsoft's prescriptive guidance for conversational user experiences in Power Platform and Dynamics 365 implementations. The Power Platform Well-Architected Experience Optimization pillar (which is part of the broader Microsoft guidance set) publishes the article "Recommendations for designing conversational user experiences" for conversational AI design. In this exam family, "conversational user experiences" maps to Success by Design's experience guidance for Dynamics 365 and Power Platform projects, which is what Fabrikam's stated requirement ("the final AI agent must follow Microsoft recommendations for a conversational user experience") points to.

Why the other options are wrong

ALM Accelerator for Microsoft Power Platform is a tool for application lifecycle management (source control, pipelines, deployments). It is not a best-practice or conversational-design framework.

Microsoft Cloud Adoption Framework for Azure is the framework for adopting Azure infrastructure (used in Question 2 for the infrastructure migration), not Copilot Studio agent design or conversational UX.

Microsoft Learn references

- Microsoft Power Platform Well-Architected - https://learn.microsoft.com/power-platform/well-architected/
- Recommendations for designing conversational user experiences - https://learn.microsoft.com/power-platform/well-architected/experience-optimization/conversation-design
- Power Platform and Copilot Studio Architecture Center - https://learn.microsoft.com/power-platform/architecture/architecture-center-overview
- Introduction to Success by Design - https://learn.microsoft.com/dynamics365/guidance/implementation-guide/success-by-design

Topic 1 - Question 2

Multiple choice | Domain: Plan | Case study: Fabrikam

topic-1/q02

Which framework should you use for the infrastructure migration?

A. Microsoft Cloud Adoption Framework for Azure

B. Success by Design

C. Microsoft Power Platform Center of Excellence (CoE)

D. Microsoft Power Platform Project Setup Wizard

Answer

A

Correct answer: A. Microsoft Cloud Adoption Framework for Azure

Why A is correct

The case study sets two requirements that together point directly at the Cloud Adoption Framework (CAF):

1. "Azure must be used for all future infrastructure workloads."
2. "The company must follow Microsoft-recommended methodologies for infrastructure migration to the cloud."

The Microsoft Cloud Adoption Framework for Azure is Microsoft's prescriptive, end-to-end methodology for adopting Azure. It explicitly covers infrastructure migration through the CAF Migrate methodology (Plan, Prepare, Execute, Optimize, Decommission), and it is the framework Microsoft Learn points enterprises to when they need to "evaluate their existing IT estate and determine the best migration strategy for each workload."

From Microsoft Learn:

"The Cloud Adoption Framework is about adopting cloud technologies and integrating them into your business. It provides proven practices to govern, secure, and manage all workloads. It provides structured processes for migrating, modernizing, and building cloud-native solutions."

"If you're new to Azure, start here. The Cloud Adoption Framework guides you through organizational preparation. It describes Azure enrollment setup, platform landing zone configuration, and migration plan prioritization. Complete these foundational steps before you move workloads."

That maps cleanly to Fabrikam's stated need: a Microsoft-recommended methodology for migrating on-premises infrastructure to Azure.

Why the other options are wrong

B. Success by Design is Microsoft's framework for **Dynamics 365 implementation projects**, not infrastructure migration. Microsoft Learn describes it as "prescriptive guidance (approaches and recommended practices) for successfully architecting, building, testing, and deploying Dynamics 365 solutions." Fabrikam will use Success by Design when it implements Dynamics 365 Sales (a separate workstream in this case study), but it is not the framework for moving infrastructure to Azure.

C. Microsoft Power Platform Center of Excellence (CoE) is a governance and adoption toolkit for Power Platform environments (apps, flows, makers, environments). It is not an infrastructure-migration framework.

D. Microsoft Power Platform Project Setup Wizard scaffolds a Power Platform implementation project. It is not a methodology for migrating on-premises infrastructure to Azure.

Microsoft Learn references

- What is the Microsoft Cloud Adoption Framework? - https://learn.microsoft.com/azure/cloud-adoption-framework/overview
- Migrate workloads to Azure (CAF as the recommended migration starting point) - https://learn.microsoft.com/azure/migration/migrate-to-azure
- CAF Migrate - Plan your migration - https://learn.microsoft.com/azure/cloud-adoption-framework/migrate/plan-migration
- CAF Plan - Prepare your organization for the cloud - https://learn.microsoft.com/azure/cloud-adoption-framework/plan/prepare-organization-for-cloud
- Introduction to Success by Design (scope: Dynamics 365 implementations) - https://learn.microsoft.com/dynamics365/guidance/implementation-guide/success-by-design

Topic 1 - Question 7

Hotspot | Domain: Plan

topic-1/q07

You need to design a shared prompt library that will be used across multiple business units. The solution must meet the following requirements:

- Ensure consistent AI responses with reusable formats.
- Support governance and version control.
- Minimize administrative effort.
- Minimize ongoing costs.

What should you recommend for each requirement? To answer, select the appropriate options in the answer area.

Answer area

- Ensure consistent AI responses:

[ ] Delegate department-specific prompt templates.
[ ] Define standardized prompt templates.
[ ] Maintain a prompt history.

- Support governance and version control:

[ ] Define standardized prompt templates.
[ ] Store prompts in a Git repository.
[ ] Categorize prompts by business function.

Answer

Ensure consistent AI responses = Define standardized prompt templates; Support governance and version control = Store prompts in a Git repository

Correct answer: - Ensure consistent AI responses: **Define standardized prompt templates.** - Support governance and version control: **Store prompts in a Git repository.**

Why these answers are correct

Standardized prompt templates are the documented Microsoft pattern for getting consistent, reusable AI responses across business units. Microsoft Learn's "Get started with prompt library" article describes prompt libraries as "a collection of predesigned prompts, serving as templates to expedite the creation of AI prompts," with the explicit goal of accelerating development "and ensures best practices are followed in prompt engineering." Standardization at the template level is what produces consistent responses with reusable formats.

Git repository storage is the standard answer for prompt governance and version control. Microsoft Learn's "AI apps and agents publication templates" article describes a prompt lifecycle template (Draft, Review, Approved, Deployed) and recommends documenting changes in source control or a wiki for traceability. Git is Microsoft's recommended version control system (Azure Repos, GitHub) and provides branching, pull request review, and history, which are the core requirements for governance and version control. Categorization by business function does not provide version control, and standardized templates by themselves do not provide change history or approval workflows.

This combination also minimizes administrative effort and ongoing cost. Templates are authored once and reused, and Git is a well-understood developer practice that the organization likely already uses.

Why the other options are wrong

Delegate department-specific prompt templates would produce divergent prompts across business units, defeating the goal of consistent responses.

Maintain a prompt history captures usage but does not enforce a consistent response format.

Define standardized prompt templates (offered for the second slot) addresses consistency, not version control. Without Git or another source-control system, you have no diff history, branching, or rollback.

Categorize prompts by business function improves discoverability but offers no governance or versioning.

Microsoft Learn references

- Get started with prompt library - https://learn.microsoft.com/microsoft-copilot-studio/prompt-library
- AI apps and agents publication templates (prompt lifecycle) - https://learn.microsoft.com/partner-center/marketplace-offers/artificial-intelligence-templates#prompt-lifecycle-template
- What is Azure Repos / Git - https://learn.microsoft.com/azure/devops/repos/get-started/what-is-repos
- Git workflow - https://learn.microsoft.com/azure/devops/repos/git/gitworkflow

Topic 1 - Question 8

Drag and drop | Domain: Plan

topic-1/q08

A company has a Microsoft Foundry project that uses a single agent and a single prompt to complete a series of tasks.

The agent encounters the following issues:

- It frequently produces incomplete results.
- It struggles with domain-specific reasoning.
- Agent response times are remarkably slow.

You need to recommend a solution to improve the overall performance and accuracy of the agent.

What should you include in the recommendation? To answer, drag the appropriate actions to the correct requirements. Each action may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

Drag sources

- Add a grounding data source.
- Add a prebuilt connector.
- Move to a multi-agent architecture.
- Upgrade to a larger generative AI model.

Drop targets

1. To improve performance - drop here
2. To improve accuracy - drop here

Answer

Performance = Move to a multi-agent architecture; Accuracy = Add a grounding data source

Correct answer: - To improve performance: **Move to a multi-agent architecture.** - To improve accuracy: **Add a grounding data source.**

Why these answers are correct

The question describes three symptoms in a single-agent, single-prompt design: - Frequently produces incomplete results - Struggles with domain-specific reasoning - Slow response times

Microsoft Learn's "Transparency Note for Azure Agent Service" calls this out directly: "When a single agent's system message consistently struggles to handle the complexity, breadth, or depth of a task, such as frequently producing incomplete results, hitting reasoning bottlenecks, or requiring extensive domain-specific knowledge, the system may benefit from transitioning to a multi-agent architecture." Splitting the workload into specialized subtasks reduces cognitive overload on any single agent, which addresses both the slow response times and the incomplete results. That maps to the **performance** improvement.

For **accuracy**, Microsoft Learn's grounding data design guidance and the Azure AI Foundry agent transparency note both point to grounding as the lever for accuracy. Adding a grounding data source gives the model authoritative reference content to cite at inference time, which directly addresses domain-specific reasoning gaps and reduces hallucinated or incomplete answers. Grounding is described in Microsoft Learn as the way to "augment generative AI models by using context data during inference" so the model "can answer the user query and form the answer according to the expectations."

Why the other options are wrong (for the slot they don't fit)

Add a prebuilt connector integrates a specific external system but doesn't by itself fix domain reasoning or speed.

Upgrade to a larger generative AI model can sometimes improve accuracy and reasoning, but it usually increases latency and cost rather than improving performance, and Microsoft's published guidance for the symptoms listed (incomplete results, domain reasoning, slow response) points to multi-agent decomposition and grounding before model upgrades.

Microsoft Learn references

- Transparency Note for Azure Agent Service (single vs. multi-agent guidance) - https://learn.microsoft.com/azure/foundry/responsible-ai/agents/transparency-note
- Single agent or multiple agents - https://learn.microsoft.com/azure/cloud-adoption-framework/ai-agents/single-agent-multiple-agents
- Grounding data design for AI workloads on Azure - https://learn.microsoft.com/azure/well-architected/ai/grounding-data-design
- Knowledge sources in Copilot Studio - https://learn.microsoft.com/microsoft-copilot-studio/knowledge-copilot-studio

Topic 1 - Question 10

Multiple choice | Domain: Plan

topic-1/q10

A company plans to deploy a Microsoft Copilot Studio agent that will analyze historical business data to predict customer behavior.

The data is currently stored in an Azure SQL database, flat files, APIs, and logs.

You need to organize the data into a format that can be used as a knowledge source in Copilot Studio.

What should you include in the solution?

A. Azure AI Search
B. Azure Data Lake Storage
C. Azure Cosmos DB
D. Azure Translator in Foundry Tools

Answer

B

Correct answer: B. Azure Data Lake Storage

Why B is correct

The requirement is to consolidate data that currently lives in multiple, mixed sources (Azure SQL, flat files, APIs, logs) into a single organized format that can serve as a knowledge source in Copilot Studio for predictive analysis on historical business data.

Azure Data Lake Storage is the Microsoft-recommended landing place for that kind of mixed-source historical data. Microsoft Learn's grounding data design guidance describes consolidating structured, semi-structured, and unstructured data into a centralized data lake (or warehouse) so that it can be efficiently indexed, transformed, and queried for AI workloads. Azure Data Lake Storage also integrates with Microsoft Fabric, Azure Synapse, and downstream tools that can prepare the data for use as a knowledge source. Power Platform / Copilot Studio can then connect to that consolidated data (for example via a Fabric Data Agent connection or through Dataverse-mediated patterns) for grounding.

For predictive analysis over historical business data drawn from heterogeneous sources, the lake is the consolidation layer. After data lands there, it can be processed and surfaced to Copilot Studio.

Why the other options are wrong

A. Azure AI Search is a retrieval/index service, not a consolidation layer for raw mixed sources. It is a downstream component used for grounding once data is prepared. It doesn't ingest and organize SQL, flat files, APIs, and logs by itself.

C. Azure Cosmos DB is an operational NoSQL database optimized for low-latency transactional workloads. It is not the recommended target for consolidating diverse historical analytics data.

D. Azure Translator in Foundry Tools is a translation service. It has nothing to do with consolidating historical business data.

Microsoft Learn references

- Grounding data design for AI workloads on Azure - https://learn.microsoft.com/azure/well-architected/ai/grounding-data-design
- Azure Data Lake Storage introduction - https://learn.microsoft.com/azure/storage/blobs/data-lake-storage-introduction
- Microsoft Fabric Lakehouse overview -

https://learn.microsoft.com/fabric/data-engineering/lakehouse-overview
- Knowledge sources in Copilot Studio - https://learn.microsoft.com/microsoft-copilot-studio/knowledge-copilot-studio

Topic 1 - Question 11

Multiple choice | Domain: Plan

topic-1/q11

A retail company plans to deploy Microsoft Copilot Studio agents to support:

- Microsoft Dynamics 365 Commerce scenarios.
- A Microsoft Power Apps inventory management solution.

You need to recommend a solution to organize product catalog data as a consistent source for multiple AI systems.

What should you recommend?

A. Let each agent scrape product details from Microsoft SharePoint Online libraries.
B. Store the product catalog data in a separate custom table for each agent.
C. Configure prompts to pull product details from the PDFs of external vendors.
D. Centralize the product catalog data in Microsoft Dataverse and expose the data to both agents.

Answer

D

Correct answer: D. Centralize the product catalog data in Microsoft Dataverse and expose the data to both agents.

Why D is correct

The scenario has two AI systems (a Copilot Studio agent for Dynamics 365 Commerce, and a Power Apps inventory management solution) that both need a consistent view of product catalog data. Microsoft's documented pattern is to centralize that kind of master data in Microsoft Dataverse and use it as a single source of truth.

Microsoft Learn ("Why choose Microsoft Dataverse?") describes Dataverse as the platform's central data store, with built-in security, native integration with Power Apps, Power Automate, Dynamics 365, and Copilot Studio, and AI-ready features (Dataverse search, glossary terms, synonyms) that support generative AI grounding. The Power Platform Well-Architected guidance for Copilot Studio agents recommends consolidating data on a centralized platform such as Dataverse to streamline access and management for multiple agents.

Centralization also avoids data drift and inconsistent answers across the two AI systems, which is the core of "consistent source for multiple AI systems."

Why the other options are wrong

A. Let each agent scrape product details from Microsoft SharePoint Online libraries produces inconsistent results, depends on document layout, scales poorly, and violates the "consistent source" requirement.

B. Store the product catalog data in a separate custom table for each agent duplicates data, creates synchronization problems, and directly contradicts the requirement for a single consistent source.

C. Configure prompts to pull product details from the PDFs of external vendors routes catalog data through external documents, creates lag and accuracy problems, and is not a centralized authoritative source.

Microsoft Learn references

- Why choose Microsoft Dataverse? - https://learn.microsoft.com/power-apps/maker/data-platform/why-dataverse-overview
- Add Dataverse tables as a knowledge source - https://learn.microsoft.com/microsoft-copilot-studio/knowledge-add-dataverse
- Power Platform Well-Architected: Experience Optimization (consolidate data on Dataverse) - https://learn.microsoft.com/power-platform/architecture/solution-ideas/agent-custom-contact-center

Topic 1 - Question 12

Multiple choice | Domain: Plan

topic-1/q12

A company has a portfolio of AI initiatives at different stages of development.

You need to recommend a structured approach to evaluating the return on AI investment (ROAI) across all the initiatives. The solution must balance immediate results with long-term values and strategic innovations.

What should you include in the recommendation?

A. a simple cost and benefit analysis
B. a horizon-based framework
C. the internal rate of return (IRR) function
D. a prioritization grid

Answer

B

Correct answer: B. a horizon-based framework

Why B is correct

The requirement is to evaluate ROAI across a portfolio of AI initiatives at different stages of development, balancing immediate results, long-term value, and strategic innovations. That maps directly to a horizon-based framework.

Microsoft Learn's Power Platform "Define roles and responsibilities" article describes horizon mapping as a technique to "visualize and prioritize tasks across different time horizons" with three horizons: Horizon 1 (immediate), Horizon 2 (near-term), and Horizon 3 (long-term and strategic). The "Pillar 1: AI strategy and experience" guidance similarly describes sequencing initiatives "across near, mid, and longer term horizons" to balance short-term wins with strategic innovation. That is the exact balance the question asks for.

A horizon-based framework lets the organization classify each AI initiative by where it sits on the value timeline, then evaluate ROAI in a way that doesn't penalize long-horizon strategic bets and doesn't over-credit short-term efficiency wins.

Why the other options are wrong

A. a simple cost and benefit analysis captures one snapshot of cost vs. expected benefit. It doesn't structure a portfolio across time horizons or distinguish strategic innovation from immediate results.

C. the internal rate of return (IRR) function is a single financial metric. It can be one input but doesn't, by itself, structure a portfolio evaluation across short, medium, and long-term horizons.

D. a prioritization grid is useful for ranking individual initiatives by impact vs. effort but doesn't capture the time-horizon dimension Microsoft uses to balance immediate results against strategic innovation.

Microsoft Learn references

- Define roles and responsibilities (Horizon mapping) - https://learn.microsoft.com/power-platform/guidance/adoption/roles
- Pillar 1: AI strategy and experience (multi-horizon strategy) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/maturity-model-strategy
- Plan for AI adoption (prioritize use cases / ROAI) - https://learn.microsoft.com/azure/cloud-adoption-framework/ai/plan

Topic 1 - Question 13

Multiple choice | Domain: Plan

topic-1/q13

You need to recommend a Microsoft Power Platform business solution that consolidates data from multiple internal and external data sources. The solution must meet the following requirements:

- Provide the data as a centralized source for multiple AI systems, including Microsoft Copilot Studio agents, Dynamics 365 applications, and external AI models.
- Support built-in data classification and protection policies.
- Provide data for grounding and analytics.

What should you include in the recommendation?

A. Microsoft Dataverse
B. Azure Data Lake Storage
C. a Microsoft Power BI semantic model
D. Azure Cosmos DB

Answer

A

Correct answer: A. Microsoft Dataverse

Why A is correct

The requirement set is: 1. Centralized source for multiple AI systems (Copilot Studio agents, Dynamics 365 apps, external AI models). 2. Built-in data classification and protection policies. 3. Data for grounding and analytics.

Microsoft Dataverse is the only option that meets all three:

- It is the standard data backbone for Copilot Studio agents and Dynamics 365 apps. Microsoft Learn's "Why choose Microsoft Dataverse?" article describes it as the central data store with native integration to Power Apps, Power Automate, Dynamics 365, and Copilot Studio. External AI systems can also reach Dataverse data via APIs, virtual tables, MCP servers, and Microsoft Graph.
- It includes built-in data classification, security roles, row/column-level security, and data policies. Microsoft Purview integrates with Dataverse for sensitivity labels, DLP for Dataverse, and data classification.
- It supports grounding (it is a documented Copilot Studio knowledge source) and analytics (managed data lake export to Azure Data Lake / Fabric, Power BI integration).

This is the documented Microsoft Power Platform business solution choice for the requirements stated.

Why the other options are wrong

B. Azure Data Lake Storage is a data store for analytics but lacks the built-in business data classification, protection policies, and Power Platform-native integrations required to be a centralized source for Copilot Studio, Dynamics 365, and external AI models out of the box.

C. a Microsoft Power BI semantic model is a reporting layer over data, designed for analytics. It is not the centralized authoritative business data store, and it does not provide built-in data classification and protection policies in the way required across multiple AI systems.

D. Azure Cosmos DB is a globally distributed NoSQL database. It is not the Power Platform data backbone and lacks the integrated classification, sensitivity, and AI-ready features described for Dataverse.

Microsoft Learn references

- Why choose Microsoft Dataverse? - https://learn.microsoft.com/power-apps/maker/data-platform/why-dataverse-overview
- Microsoft Purview data classification (Dataverse) - https://learn.microsoft.com/purview/ai-microsoft-purview
- Add Dataverse tables as a knowledge source - https://learn.microsoft.com/microsoft-copilot-studio/knowledge-add-dataverse
- Step 2: Protect sensitive info in grounding data (Dataverse) - https://learn.microsoft.com/purview/deploymentmodels/depmod-sc-agents-step2

Topic 1 - Question 14

Multiple choice | Domain: Plan

topic-1/q14

A company plans to deploy an AI-based customer service app that will autonomously manage interactions, escalate complex cases, and learn from historical ticket data.

You need to perform a return on AI investment (ROAI) analysis of the app deployment. The solution must ensure that the analysis is accurate.

What should you do first?

A. Establish the AI performance metrics.
B. Conduct an AI market benchmarking study.
C. Model the customer experience.
D. Identify and quantify all the development, deployment, and operating costs.

Answer

D

Correct answer: D. Identify and quantify all the development, deployment, and operating costs.

Why D is correct

ROAI is, at its core, benefits relative to costs. To make the analysis accurate, you have to establish the cost baseline first. Microsoft Learn's "Plan to manage costs for Azure OpenAI" and "Quantify business value" guidance both put cost identification at the start of any ROI / ROAI analysis. The Azure Cloud Adoption Framework's "Business plan for AI agents" and "Plan for AI adoption" articles also describe baseline costs (development, deployment, operating) as the starting point for return-on-investment analysis, then comparing those costs against measurable benefits.

Without identifying and quantifying full costs (people, model usage, infrastructure, integration, support, retraining, governance), any subsequent metric or benchmark is built on an incomplete denominator. That is why the **first** step is identifying and quantifying all costs.

Why the other options are wrong

A. Establish the AI performance metrics is necessary later in the analysis, but you can't relate performance to ROAI until you know the cost basis you are measuring performance against.

B. Conduct an AI market benchmarking study is useful context but doesn't tell you anything about your own deployment's economics. ROAI accuracy depends on your costs and your benefits, not industry averages.

C. Model the customer experience informs benefits estimation but is not the first step. You need a cost baseline first.

Microsoft Learn references

- Business plan for AI agents (define success metrics, decision gates) - https://learn.microsoft.com/azure/cloud-adoption-framework/ai-agents/business-strategy-plan
- Plan to manage costs for Azure OpenAI - https://learn.microsoft.com/azure/ai-foundry/openai/how-to/manage-costs
- Quantify business value (FinOps) - https://learn.microsoft.com/cloud-computing/finops/framework/quantify/quantify-business-value
- Plan for AI adoption - https://learn.microsoft.com/azure/cloud-adoption-framework/ai/plan

Topic 1 - Question 22

Multiple choice | Domain: Plan

topic-1/q22

A manufacturing company wants to deploy an agent that will automate supplier invoice processing.

You are designing a solution to evaluate the financial implications of the deployment. The company is especially concerned about budget overruns.

You need to ensure that the solution considers the total cost of ownership (TCO), the expected savings from using automation, and whether to extend the existing AI capabilities.

What should you include in the design?

A. a break-even analysis only
B. adopting prebuilt agents to reduce the deployment time
C. training a custom model
D. a return on AI investment (ROAI) analysis

Answer

D

Correct answer: D. a return on AI investment (ROAI) analysis

Why D is correct

The design has to evaluate three things together: total cost of ownership (TCO), expected savings from the automation, and whether to extend existing AI capabilities. That is the exact scope of an ROAI analysis. Microsoft Learn's "Business plan for AI agents" and "Plan for AI adoption" articles describe ROAI as the structured comparison of all costs against measurable benefits, including decisions about whether to expand or modify the AI investment based on the data.

The Copilot Studio "Analyze agent return on investment" feature also lets organizations track time and money savings per agent run, feeding directly into an ROAI view. ROAI captures both the TCO side (development, deployment, operating, governance, model usage) and the benefits side (savings, productivity, cycle-time reductions), and it provides the basis for the extension decision.

Why the other options are wrong

A. a break-even analysis only answers when costs are recouped. It is one input to ROAI but doesn't cover the whole scope (TCO + ongoing savings + extension decision). The word "only" makes it incomplete for the requirements stated.

B. adopting prebuilt agents to reduce the deployment time is a tactical choice, not a financial evaluation framework. It doesn't analyze TCO or savings.

C. training a custom model is an implementation choice and increases costs. It doesn't evaluate the financial implications of the deployment.

Microsoft Learn references

- Business plan for AI agents - https://learn.microsoft.com/azure/cloud-adoption-framework/ai-agents/business-strategy-plan
- Plan for AI adoption - https://learn.microsoft.com/azure/cloud-adoption-framework/ai/plan
- Analyze agent return on investment - https://learn.microsoft.com/power-platform/release-plan/2025wave1/microsoft-copilot-studio/analyze-agent-return-investment
- Quantify business value (FinOps) - https://learn.microsoft.com/cloud-computing/finops/framework/quantify/quantify-business-value

Topic 1 - Question 23

Multiple choice | Domain: Plan

topic-1/q23

A company has a Microsoft Power Platform solution that contains the following components:

- Microsoft Dataverse tables
- A Microsoft Power BI workspace named WS1
- A canvas app named App1 that uses Dataverse
- A Power BI semantic model that connects to Dataverse by using DirectQuery

You plan to use generative AI to provide answers to queries based on a subset of corporate data.

You need to ensure that the data is available as a grounding data source for AI systems.

What should you do?

A. Populate a Dataverse table.
B. Share WS1.
C. Endorse the semantic model.
D. Export the semantic model.

Answer

C

Correct answer: C. Endorse the semantic model.

Why C is correct

The objective is to make the data available as a grounding source for AI systems that answer queries against a subset of corporate data. The data already lives in Dataverse, the canvas app already uses it, and a Power BI semantic model already connects to Dataverse via DirectQuery.

Microsoft Learn's "Endorse your content" and "Add a Power BI semantic model as a data source to Fabric data agent" articles explain that semantic models endorsed (promoted or certified) in Power BI are clearly identified, prioritized in searches, and discoverable. Endorsement is the documented way to mark a semantic model as authoritative and ready for reuse, including for AI grounding scenarios such as

Power BI Copilot, Fabric Data Agents, and Microsoft 365 Copilot. The semantic model best practices article for data agents calls out that endorsement combined with Prep for AI is what makes a semantic model a reliable grounding source for AI systems.

Endorsing the semantic model surfaces it in the OneLake catalog and in AI tools as a trusted grounding source, which directly satisfies the requirement to "make the data available as a grounding data source for AI systems."

Why the other options are wrong

A. Populate a Dataverse table doesn't add anything new for grounding. The Dataverse data is already there and is already queried by the existing semantic model.

B. Share WS1 controls workspace access but does not endorse or expose the model as a discoverable, trusted AI grounding source.

D. Export the semantic model moves it out of the Power BI service, which removes the integrated AI grounding pathway. Exported model files are not a grounding source for live AI queries.

Microsoft Learn references

- Endorse your content (Power BI promotion and certification) - https://learn.microsoft.com/power-bi/collaborate-share/service-endorse-content
- Promote and certify Power BI content with endorsement - https://learn.microsoft.com/power-bi/collaborate-share/service-endorsement-overview
- Add a Power BI semantic model as a data source to Fabric data agent - https://learn.microsoft.com/fabric/data-science/data-agent-semantic-model
- Semantic model best practices for data agent - https://learn.microsoft.com/fabric/data-science/semantic-model-best-practices

Topic 1 - Question 26

Multiple choice | Domain: Plan

topic-1/q26

A company is designing a Microsoft Power Platform solution to reduce the manual steps of a business process by deploying an existing AI model.

You need to calculate the return on AI investment (ROAI) by identifying the metadata and telemetry of the solution.

What should you use?

A. Microsoft Power Platform admin center
B. Success by Design
C. the Business value toolkit
D. Microsoft Cloud Adoption Framework for Azure

Answer

C

Correct answer: C. the Business value toolkit

Why C is correct

The requirement is to calculate ROAI for a Power Platform solution by identifying the **metadata and telemetry** of the solution. That is the exact purpose of the Business value toolkit.

Microsoft Learn's "Choose the best methods and tools to measure business value" article maps use cases to recommended methods and explicitly states: for "Everybody" (large numbers of users, smallest number of solutions), the identification method is "Metadata and telemetry" and the recommended methodology is the **Business value toolkit**. The "Capture and communicate value with the Business value toolkit" article describes the toolkit as part of the Center of Excellence starter kit, designed to capture and communicate the value of Power Platform solutions, including AI-driven solutions, using a structured storytelling process aligned to organizational objectives.

The Business value toolkit is also the documented Power Platform tool for ROI/ROAI measurement when the input is solution metadata and telemetry (which is what the question states).

Why the other options are wrong

A. Microsoft Power Platform admin center provides administrative monitoring and capacity views. It is not the ROAI calculation tool.

B. Success by Design is the prescriptive implementation methodology for Dynamics 365. It is not the ROAI / business-value measurement tool for Power Platform solutions.

D. Microsoft Cloud Adoption Framework for Azure is the framework for adopting Azure infrastructure (used for the broader migration in the Fabrikam case study). It is not the Power Platform business-value measurement tool.

Microsoft Learn references

- Choose the best methods and tools to measure business value - https://learn.microsoft.com/power-platform/guidance/adoption/business-value-methods
- Capture and communicate value with the Business value toolkit - https://learn.microsoft.com/power-platform/guidance/coe/business-value-toolkit
- Measure value and realize return on investment - https://learn.microsoft.com/power-platform/guidance/adoption/common-vision/realize-value
- Measure and communicate the business value of Power Platform solutions - https://learn.microsoft.com/power-platform/guidance/adoption/business-value

Topic 1 - Question 28

Multiple choice | Domain: Plan

topic-1/q28

A company has an Azure environment that supports multiple business units.

The company plans to implement an AI solution that will perform sentiment analysis on customer product reviews.

You need to evaluate the potential cost of the solution to support return on AI investment (ROAI) analysis.

What should you use?

A. Cost Management + Billing
B. Microsoft Fabric SKU Estimator
C. Total Cost of Ownership (TCO) Calculator
D. Azure Reservations

Answer

C

Correct answer: C. Total Cost of Ownership (TCO) Calculator

Why C is correct

The requirement is to **evaluate the potential cost** of an Azure AI solution to support an ROAI analysis. That is a forward-looking estimate of cost before deployment, and it is exactly the use case for the Microsoft TCO Calculator.

Microsoft Learn's "Make an inventory and collect data" article describes the TCO Calculator as the tool to "compare the costs of running your workloads on-premises versus on Azure" and to gain "insights into potential cost savings and efficiencies." The "Describe cost management in Azure" learning module teaches the TCO Calculator alongside the Pricing Calculator as the planning tools you use **before** deploying workloads, which is the input ROAI analysis needs.

For an organization evaluating the cost of a new sentiment-analysis solution before deployment, the TCO Calculator (or Pricing Calculator) is the documented planning tool, and TCO is the option offered.

Why the other options are wrong

A. Cost Management + Billing analyzes **actual** Azure spending after deployment. It is not a planning/estimation tool for a solution that hasn't been built yet.

B. Microsoft Fabric SKU Estimator estimates Microsoft Fabric capacity costs specifically. It is too narrow for evaluating a broader AI sentiment-analysis solution that may use Azure AI services beyond Fabric.

D. Azure Reservations is a purchasing option that gives discounts for committed usage. It is not an estimation tool for evaluating potential cost.

Microsoft Learn references

- Make an inventory and collect data (TCO Calculator) - https://learn.microsoft.com/azure/app-modernization-guidance/assess/make-an-inventory-and-collect-data
- Compare the Pricing and Total Cost of Ownership calculators (training) - https://learn.microsoft.com/training/modules/describe-cost-management-azure/3-compare-pricing-total-cost-of-ownership-calculators
- Plan to manage costs for Azure OpenAI - https://learn.microsoft.com/azure/ai-foundry/openai/how-to/manage-costs
- Estimate costs with the Azure pricing calculator - https://learn.microsoft.com/azure/cost-management-billing/costs/pricing-calculator

Topic 1 - Question 31

Multi-select | Domain: Plan | Case study: Contoso

topic-1/q31

Which two components in the custom AI agent design should the CFO evaluate in the quarterly agent analysis? Each correct answer presents part of the solution.

A. the GPT models used for the agent
B. the average characters in a chat message
C. the agent orchestration method
D. the average session time per agent

Select all that apply.

Answer

A, C

Correct answers: A. the GPT models used for the agent, and C. the agent orchestration method

Why these answers are correct

The Contoso case study states: "The CFO will analyze all the AI solutions quarterly to compare the estimated ROI against actual measured efficiencies and adoption. The CFO will use the Copilot Studio agent usage estimator to perform this analysis."

Microsoft Learn's Copilot Studio billing/licensing pages describe the inputs to the Copilot Studio agent usage estimator directly: "Forecast your agent's Copilot Credits volume using the Microsoft Copilot Studio agent usage estimator. Create estimates and potential consumption impacts by selecting from agent type, traffic, **orchestration**, knowledge, and tools."

The estimator (and Copilot Credits cost in general) is sensitive to two design choices that the CFO would evaluate:

A. The GPT models used for the agent. Copilot Studio's models page and the Application Card describe how different language models (GPT-4o, GPT-4.1, GPT-5, third-party models) drive different consumption and capability profiles. Model choice directly affects Copilot Credit consumption per interaction, which is a primary cost driver in the estimator and therefore central to the CFO's quarterly ROI analysis.

C. The agent orchestration method. Orchestration (generative vs. classic) is an explicit input to the agent usage estimator. Generative orchestration consumes more Copilot Credits per interaction because the model is used to plan and select tools, topics, and knowledge dynamically. Classic orchestration is more deterministic and typically consumes fewer credits. The CFO must consider this in the quarterly analysis.

Why the other options are wrong

B. The average characters in a chat message is not an input to the agent usage estimator and is not a primary Copilot Credit cost driver. Copilot Credits are billed by activity type and orchestration model usage, not character count.

D. The average session time per agent is a usage observation metric, not an input to the agent usage estimator. The CFO's tool of choice is the estimator, which forecasts based on agent type, traffic, orchestration, knowledge, and tools. Session time is closer to engagement analytics than ROAI cost forecasting.

Microsoft Learn references

- Copilot Studio licensing (agent usage estimator) - https://learn.microsoft.com/microsoft-copilot-studio/billing-licensing
- Microsoft Copilot Studio agent usage estimator - https://learn.microsoft.com/microsoft-copilot-studio/agent-usage-estimator
- Billing rates and management (Copilot Credits) - https://learn.microsoft.com/microsoft-copilot-studio/requirements-messages-management
- Select a primary AI model for your agent - https://learn.microsoft.com/microsoft-copilot-studio/authoring-select-agent-model
- Orchestrate agent behavior with generative AI - https://learn.microsoft.com/microsoft-copilot-studio/advanced-generative-actions

Topic 2

4 questions

Topic 2 - Question 15

Multiple choice | Domain: Plan

topic-2/q15

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.

After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.

A company uses Microsoft 365 and Dynamics 365.

You need to recommend a solution to automatically summarize email threads, generate suggested replies in Microsoft Outlook, and provide meeting preparation summaries that include relevant customer relationship management (CRM) data.

Solution: You recommend Microsoft 365 Copilot for Sales.

Does this meet the goal?

A. Yes
B. No

Answer

A

Correct answer: A. Yes

Why Yes is correct

The requirement is to automatically summarize email threads, generate suggested replies in Microsoft Outlook, and provide meeting preparation summaries that include CRM data. Microsoft 365 Copilot for Sales is built specifically for that combination of capabilities.

From Microsoft Learn:

"Microsoft 365 Copilot for Sales is an AI assistant designed for sellers. The Copilot for Sales app in Outlook gives you recommendations and information to help you stay connected to your customers, reduce data entry, and personalize your engagements."

"With Copilot, you can get related information from your CRM system, such as contact details, account history, and opportunities... You get an overview of recent interactions with your customers, including email summaries, meeting notes, and action items. Use AI capabilities to draft emails, summarize conversations, and generate follow-ups."

"Sales agent integrates Microsoft 365 Copilot in Outlook with your CRM system to help you build customer relationships and close deals faster. Microsoft 365 Copilot summarizes your emails, helping you quickly grasp key details. Sales agent enriches those summaries with sales information from your Dynamics 365 or Salesforce CRM."

That covers all three goals: email-thread summaries, suggested replies in Outlook, and CRM-grounded meeting preparation summaries.

Microsoft Learn references

- Microsoft Outlook experiences for Microsoft 365 Copilot for Sales - https://learn.microsoft.com/copilot/release-plan/2026wave1/copilot-sales/outlook-experiences
- Boost sales efficiency with CRM-enriched email summaries in Outlook - https://learn.microsoft.com/microsoft-sales-copilot/email-summary-premium
- Agents, Copilot, and AI capabilities in Dynamics 365 apps (Sales in Microsoft 365 Copilot) - https://learn.microsoft.com/dynamics365/copilot/ai-get-started
- Copilot in Dynamics 365 Sales overview - https://learn.microsoft.com/dynamics365/sales/copilot-overview

Topic 2 - Question 16

Multiple choice | Domain: Plan

topic-2/q16

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.

After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.

A company uses Microsoft 365 and Dynamics 365.

You need to recommend a solution to automatically summarize email threads, generate suggested replies in Microsoft Outlook, and provide meeting preparation summaries that include relevant customer relationship management (CRM) data.

Solution: You recommend a classic Microsoft Dataverse workflow.

Does this meet the goal?

A. Yes
B. No

Answer

B

Correct answer: B. No

Why No is correct

A classic Microsoft Dataverse workflow is a deterministic, rule-based process designed to automate record updates and business logic in Dataverse tables. It is not a generative AI capability.

The stated requirements are to summarize email threads, generate suggested replies in Outlook, and produce meeting preparation summaries that include CRM data. Those are all generative AI features delivered by Microsoft 365 Copilot and Microsoft 365 Copilot for Sales (with the Sales agent), not by classic Dataverse workflows.

From Microsoft Learn:

"Microsoft 365 Copilot for Sales is an AI assistant designed for sellers... With Copilot, you can get related information from your CRM system... You get an overview of recent interactions with your customers, including email summaries, meeting notes, and action items. Use AI capabilities to draft emails, summarize conversations, and generate follow-ups."

A classic Dataverse workflow cannot summarize emails, generate suggested replies in Outlook, or produce CRM-enriched meeting summaries. The proposed solution does not meet any of the stated goals.

Microsoft Learn references

- Microsoft Outlook experiences for Microsoft 365 Copilot for Sales - https://learn.microsoft.com/copilot/release-plan/2026wave1/copilot-sales/outlook-experiences
- Boost sales efficiency with CRM-enriched email summaries in Outlook - https://learn.microsoft.com/microsoft-sales-copilot/email-summary-premium
- Copilot in Dynamics 365 Sales overview - https://learn.microsoft.com/dynamics365/sales/copilot-overview
- Dataverse classic workflows (background processes, not generative AI) - https://learn.microsoft.com/power-automate/workflow-processes

Topic 2 - Question 17

Multiple choice | Domain: Plan

topic-2/q17

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.

After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.

A company uses Microsoft 365 and Dynamics 365.

You need to recommend a solution to automatically summarize email threads, generate suggested replies in Microsoft Outlook, and provide meeting preparation summaries that include relevant customer relationship management (CRM) data.

Solution: You recommend a Microsoft 365 Copilot agent template.

Does this meet the goal?

A. Yes

B. No

Answer

B

Correct answer: B. No

Why No is correct

A Microsoft 365 Copilot agent template is a starting point for building a custom declarative or custom-engine agent for Microsoft 365 Copilot. It is not the productized Outlook and CRM integration that delivers email summaries, suggested replies in Outlook, and CRM-grounded meeting preparation summaries out of the box.

The Microsoft-recommended product for the stated goals is Microsoft 365 Copilot for Sales, which integrates Copilot in Outlook with Dynamics 365 or Salesforce and provides those exact capabilities natively.

From Microsoft Learn:

"Microsoft 365 Copilot for Sales is an AI assistant that empowers sellers with insights, recommendations, actions, and up-to-date CRM data. It works directly within the Microsoft 365 applications sellers already use daily. Copilot for Sales is a Copilot-first application that works with both Dynamics 365 Sales and Salesforce."

"Sales agent integrates Microsoft 365 Copilot in Outlook with your CRM system to help you build customer relationships and close deals faster. Microsoft 365 Copilot summarizes your emails, helping you quickly grasp key details. Sales agent enriches those summaries with sales information from your Dynamics 365 or Salesforce CRM."

A generic Microsoft 365 Copilot agent template would still require substantial custom build work for Outlook email summaries, suggested replies, and CRM-enriched meeting prep. It is not the right fit and does not meet the goal as designed.

Microsoft Learn references

- Microsoft Outlook experiences for Microsoft 365 Copilot for Sales - https://learn.microsoft.com/copilot/release-plan/2026wave1/copilot-sales/outlook-experiences
- Boost sales efficiency with CRM-enriched email summaries in Outlook - https://learn.microsoft.com/microsoft-sales-copilot/email-summary-premium
- Copilot in Dynamics 365 Sales overview - https://learn.microsoft.com/dynamics365/sales/copilot-overview
- Microsoft 365 Copilot agent templates (declarative agents) - https://learn.microsoft.com/microsoft-365/copilot/extensibility/overview-declarative-agent

Topic 2 - Question 28

Multiple choice | Domain: Plan | Case study: Contoso

topic-2/q28

What should you recommend to assist the CTO with the prebuilt agent selection process?

A. Agent management
B. Copilot Studio
C. Lifecycle Services (LCS)
D. Immersive Home

Answer

A

Correct answer: A. Agent management

Why A is correct

The Contoso case study (questions 28 and 29) says the CTO "wants to view available agent templates to identify which agent will add the most business value" for Dynamics 365 Supply Chain Management. The Microsoft-documented experience for discovering, previewing, configuring, and activating prebuilt Dynamics 365 finance and operations agents is the agent management feature, accessed through Immersive Home in finance and operations apps.

From Microsoft Learn:

"By using the agent management feature in finance and operations apps, you can use autonomous, AI-powered agents to perform predefined tasks in your business ecosystem. Users can discover, configure, and manage agents that automate routine operational tasks."

"Use the following tabs on the Agent page to discover and activate new agents: Library - Browse available agent templates. Preview agent capabilities. Select and configure new agents for deployment. Manage - View agents that are currently active. Edit existing agent configurations. Activate or deactivate individual agent tasks."

The Library tab in Agent management is exactly where the CTO previews available prebuilt agent templates and selects the one that adds the most business value.

Why the other options are wrong

B. Copilot Studio is for building and customizing custom AI agents. The CTO is selecting a prebuilt agent from Dynamics 365 Supply Chain Management; the discovery and activation experience is Agent management, not Copilot Studio.

C. Lifecycle Services (LCS) is the legacy environment management portal for finance and operations apps. Microsoft is moving environment management to the Power Platform admin center, and LCS is not where the CTO browses prebuilt agent templates. (From Microsoft Learn: "Starting February 2026, new customers can't create projects in Microsoft Dynamics Lifecycle Services for Microsoft Dynamics 365 Finance, Microsoft Dynamics 365 Human Resources, Microsoft Dynamics 365 Supply Chain Management, and Microsoft Dynamics 365 Project Operations.")

D. Immersive Home is the navigation hub that gives access to agent-related pages, but the documented feature for selecting prebuilt agent templates is Agent management. Immersive Home is the entry point, not the selection feature itself.

Microsoft Learn references

- Enable agent management (production ready preview) - https://learn.microsoft.com/dynamics365/fin-ops-core/fin-ops/copilot/agent-mgmt
- Agent management release notes - https://learn.microsoft.com/dynamics365/release-plan/2024wave2/finance-supply-chain/dynamics365-finance/agent-management
- Agents, Copilot, and AI capabilities in Dynamics 365 apps - https://learn.microsoft.com/dynamics365/copilot/ai-get-started
- Immersive Home (production ready preview) - https://learn.microsoft.com/dynamics365/fin-ops-core/fin-ops/copilot/immersive-home

Topic 3

2 questions

Topic 3 - Question 2

Hotspot | Domain: Plan | Case study: Fabrikam

topic-3/q02

HOTSPOT -

Which existing tool and data should you use to gather the required metrics for stakeholder signoff for the AI agents? To answer, select the appropriate options in the answer area.

Answer area

- Tool:

[ ] Microsoft Foundry
[ ] Azure Resource Monitor (ARM)
[ ] Dynamics 365 Sales
[ ] Microsoft Copilot Studio

- Data required for the tool:

[ ] the cumulative time spent on the task over the past year
[ ] the current cost to complete the tasks per instance
[ ] the current time to complete the task today per instance
[ ] the current cost of the Dynamics 365 Sales licenses

Answer

Microsoft Copilot Studio; the current cost to complete the tasks per instance

Correct answer: - **Tool: Microsoft Copilot Studio** - **Data required for the tool: the current cost to complete the tasks per instance**

Why these selections are correct

The case study requires "the return on investment (ROI) of switching from the current process to the future process" for stakeholder sign-off, and the agent in question is the low-code Copilot Studio agent grounded in Dataverse. Copilot Studio includes a built-in **Savings calculator** on the Analytics page that estimates ROI for the agent. The calculator compares the agent's runs against the existing manual method on a per-run or per-tool basis and reports time and money saved over the selected period.

Microsoft Learn says: "Use the Savings calculator to estimate your savings per run or to calculate the savings of one or more of the tools your agent uses." To compute that comparison, you enter "the estimated cost required for another method to complete the task of the agent." That maps directly to the **current cost to complete the tasks per instance** baseline.

Microsoft Foundry, Azure Resource Monitor, and Dynamics 365 Sales do not provide a comparable built-in ROI calculator for the Copilot Studio agent in the case study. Cumulative time over the past year, current time per instance, and Dynamics 365 Sales license cost do not match the calculator's expected input, which is the per-instance cost of the existing method to complete the task.

Microsoft Learn references

- Analyze time and cost savings for agents (Savings calculator) - https://learn.microsoft.com/microsoft-copilot-studio/analytics-cost-savings
- Estimate savings per run for an agent - https://learn.microsoft.com/microsoft-copilot-studio/analytics-cost-savings#estimate-savings-per-run-for-an-agent
- Analyze agent return on investment (release plan) - https://learn.microsoft.com/power-platform/release-plan/2025wave1/microsoft-copilot-studio/analyze-agent-return-investment

Topic 3 - Question 33

Multi-select | Domain: Plan

topic-3/q33

A company has an AI business solution that uses Microsoft Copilot Studio agents.

You need to recommend prompt best practices to improve the effectiveness of agent interactions.

Which two actions should you include in the recommendation? Each correct answer presents part of the solution.

A. Track the duration of the average user session.
B. Analyze the prompt length distribution.
C. Regularly test and refine the prompts based on user input.
D. Use clear and specific instructions in the prompts.
E. Measure system resource usage during prompt processing.

Select all that apply.

Answer

C, D

Correct answer: C. Regularly test and refine the prompts based on user input and D. Use clear and specific instructions in the prompts.

Why C and D are correct

Microsoft's prompt and instruction guidance for Copilot Studio centers on two practices: write clear, specific instructions, and iterate based on real interactions. Microsoft Learn's "Optimize prompts and instructions" guide says to "refine prompts, tool descriptions, and instructions to improve clarity, reduce ambiguity, and guide the orchestrator toward more accurate and consistent responses." That covers both clarity (D) and the iterative refinement loop (C). Real user input and feedback signals (reactions, low-quality answers, themes) drive continuous prompt improvement.

Why the other options are wrong

A. Track the duration of the average user session. Session duration is an engagement metric. It does not improve prompt effectiveness directly.

B. Analyze the prompt length distribution. Length analysis is descriptive. It does not produce a better prompt without the clarity and iteration practices.

E. Measure system resource usage during prompt processing. Resource usage is an operational concern, not a prompt-quality practice.

Microsoft Learn references

- Optimize prompts and instructions (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/optimize-prompts-overview
- Improve your Copilot Studio projects (iterative improvement loop) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/improve-overview
- Review the improve checklist - https://learn.microsoft.com/microsoft-copilot-studio/guidance/improve-checklist