AB-100 Practice Questions

Agentic AI Business Solutions Architect

Generated 2026-05-05

103 questions, with answers and explanations

Question types

Drag and drop: 11

Hotspot: 22

Multi-select: 11

Multiple choice: 59

AB-100 domain coverage

Plan: 20 (19.4%)

Design: 37 (35.9%)

Deploy: 46 (44.7%)

Source: practicetest/extracted/. Each answer references Microsoft Learn. Use this PDF for offline review; the interactive version lives at app/index.html.

Topic 1

32 questions

Topic 1 - Question 1

Hotspot | Domain: Plan | Case study: Fabrikam
topic-1/q01

Which framework should you use to meet the AI agent requirements for the sales cycle enablement? To answer, select the appropriate options in the answer area.

Answer area

- For Microsoft Copilot Studio best practices:
 [ ] the ALM Accelerator for Microsoft Power Platform
 [ ] Microsoft Cloud Adoption Framework for Azure
 [ ] Microsoft Power Platform Well-Architected framework
 [ ] Success by Design

- For conversational user experiences:
 [ ] the ALM Accelerator for Microsoft Power Platform
 [ ] Microsoft Cloud Adoption Framework for Azure
 [ ] Microsoft Power Platform Well-Architected framework
 [ ] Success by Design

Answer

Microsoft Power Platform Well-Architected framework; Success by Design

Correct answer: - For Microsoft Copilot Studio best practices: **Microsoft Power Platform Well-Architected framework** - For conversational user experiences: **Success by Design**

Why these answers are correct

Fabrikam's case study calls out two distinct sets of guidance the team needs:

1. Microsoft-recommended best practices for the Copilot Studio agent itself (the low-code agent built on Dataverse).
2. Microsoft recommendations for the agent's conversational user experience (the headset-driven AI agent for the sales executives).

Microsoft Power Platform Well-Architected framework is the published Microsoft framework for designing, planning, and implementing modern application workloads with Power Platform, including Copilot Studio agents. It is organized around five pillars (Reliability, Security, Operational Excellence, Performance Efficiency, Experience Optimization) and is what Microsoft Learn directs makers to when they want best-practice design guidance for Copilot Studio.

Success by Design includes Microsoft's prescriptive guidance for conversational user experiences in Power Platform and Dynamics 365 implementations. The Power Platform Well-Architected Experience Optimization pillar (which is part of the broader Microsoft guidance set) publishes the article "Recommendations for designing conversational user experiences" for conversational AI design. In this exam family, "conversational user experiences" maps to Success by Design's experience guidance for Dynamics 365 and Power Platform projects, which is what Fabrikam's stated requirement ("the final AI agent must follow Microsoft recommendations for a conversational user experience") points to.

Why the other options are wrong

ALM Accelerator for Microsoft Power Platform is a tool for application lifecycle management (source control, pipelines, deployments). It is not a best-practice or conversational-design framework.

Microsoft Cloud Adoption Framework for Azure is the framework for adopting Azure infrastructure (used in Question 2 for the infrastructure migration), not Copilot Studio agent design or conversational UX.

Microsoft Learn references

- Microsoft Power Platform Well-Architected - https://learn.microsoft.com/power-platform/well-architected/
- Recommendations for designing conversational user experiences - https://learn.microsoft.com/power-platform/well-architected/experience-optimization/conversation-design
- Power Platform and Copilot Studio Architecture Center - https://learn.microsoft.com/power-platform/architecture/architecture-center-overview
- Introduction to Success by Design - https://learn.microsoft.com/dynamics365/guidance/implementation-guide/success-by-design

Topic 1 - Question 2

Multiple choice | Domain: Plan | Case study: Fabrikam

topic-1/q02

Which framework should you use for the infrastructure migration?

A. Microsoft Cloud Adoption Framework for Azure
B. Success by Design
C. Microsoft Power Platform Center of Excellence (CoE)
D. Microsoft Power Platform Project Setup Wizard

Answer

A

Correct answer: A. Microsoft Cloud Adoption Framework for Azure

Why A is correct

The case study sets two requirements that together point directly at the Cloud Adoption Framework (CAF):

1. "Azure must be used for all future infrastructure workloads."
2. "The company must follow Microsoft-recommended methodologies for infrastructure migration to the cloud."

The Microsoft Cloud Adoption Framework for Azure is Microsoft's prescriptive, end-to-end methodology for adopting Azure. It explicitly covers infrastructure migration through the CAF Migrate methodology (Plan, Prepare, Execute, Optimize, Decommission), and it is the framework Microsoft Learn points enterprises to when they need to "evaluate their existing IT estate and determine the best migration strategy for each workload."

From Microsoft Learn:

"The Cloud Adoption Framework is about adopting cloud technologies and integrating them into your business. It provides proven practices to govern, secure, and manage all workloads. It provides structured processes for migrating, modernizing, and building cloud-native solutions."

"If you're new to Azure, start here. The Cloud Adoption Framework guides you through organizational preparation. It describes Azure enrollment setup, platform landing zone configuration, and migration plan prioritization. Complete these foundational steps before you move workloads."

That maps cleanly to Fabrikam's stated need: a Microsoft-recommended methodology for migrating on-premises infrastructure to Azure.

Why the other options are wrong

B. Success by Design is Microsoft's framework for **Dynamics 365 implementation projects**, not infrastructure migration. Microsoft Learn describes it as "prescriptive guidance (approaches and recommended practices) for successfully architecting, building, testing, and deploying Dynamics 365 solutions." Fabrikam will use Success by Design when it implements Dynamics 365 Sales (a separate workstream in this case study), but it is not the framework for moving infrastructure to Azure.

C. Microsoft Power Platform Center of Excellence (CoE) is a governance and adoption toolkit for Power Platform environments (apps, flows, makers, environments). It is not an infrastructure-migration framework.

D. Microsoft Power Platform Project Setup Wizard scaffolds a Power Platform implementation project. It is not a methodology for migrating on-premises infrastructure to Azure.

Microsoft Learn references

- What is the Microsoft Cloud Adoption Framework? - https://learn.microsoft.com/azure/cloud-adoption-framework/overview
- Migrate workloads to Azure (CAF as the recommended migration starting point) - https://learn.microsoft.com/azure/migration/migrate-to-azure
- CAF Migrate - Plan your migration - https://learn.microsoft.com/azure/cloud-adoption-framework/migrate/plan-migration
- CAF Plan - Prepare your organization for the cloud - https://learn.microsoft.com/azure/cloud-adoption-framework/plan/prepare-organization-for-cloud
- Introduction to Success by Design (scope: Dynamics 365 implementations) - https://learn.microsoft.com/dynamics365/guidance/implementation-guide/success-by-design

Topic 1 - Question 3

Multiple choice | Domain: Design

topic-1/q03

A company uses Microsoft Dynamics 365 Sales to manage leads that are stored in a Microsoft Dataverse table named Lead and use non-standard terminology and custom columns.

You need to configure business terms in the Lead table so that Microsoft Copilot controls can summarize the leads efficiently. The solution must minimize administrative effort.

How should you configure the business terms?

A. Combine all the fields into one custom field.
B. Map the field display names as business terms.
C. Add the schema names as business terms.
D. Create new business terms for each field.

Answer

B

Correct answer: B. Map the field display names as business terms.

Why B is correct

For Copilot to interpret data in a Dataverse table that uses non-standard terminology and custom columns, it needs natural-language hints that connect business vocabulary to actual columns. Microsoft Learn's guidance on adding Dataverse tables as a knowledge source describes glossary terms and synonyms as the way to give the AI grounding context for column meaning. The lowest-effort approach is to map the existing field display names (which already describe the column in human-readable form) as the business terms Copilot uses when summarizing leads.

This minimizes administrative effort because: - Display names already exist on every column. - Mapping them as business terms avoids creating new metadata from scratch. - Copilot Studio's Dataverse knowledge source uses this mapping to recognize user requests and return responses grounded in the table's columns.

Why the other options are wrong

A. Combine all the fields into one custom field destroys data structure, breaks reporting, and is not how Copilot summarization works. Copilot summarizes by reading multiple columns, not one merged blob.

C. Add the schema names as business terms is the wrong source. Schema names (like `cr1f3_LeadValue`) are technical identifiers, not business terms. They make summaries less readable, not more.

D. Create new business terms for each field works but is more administrative effort than mapping existing display names. The question explicitly asks for the solution that minimizes administrative effort.

Microsoft Learn references

- Add Dataverse tables as a knowledge source - https://learn.microsoft.com/microsoft-copilot-studio/knowledge-add-dataverse
- Improve copilot responses from Microsoft Dataverse - https://learn.microsoft.com/power-apps/maker/data-platform/data-platform-copilot
- Synonyms and glossary terms (Dataverse knowledge) - https://learn.microsoft.com/microsoft-copilot-studio/knowledge-add-dataverse#synonyms-and-glossary-terms

Topic 1 - Question 4

Drag and drop | Domain: Design

topic-1/q04

You are designing two Microsoft Copilot Studio agents named Agent1 and Agent2. Each agent must meet the following requirements:

- Each agent must use a standard model.
- Each agent must NOT use generative orchestration.
- Agent1 must support simple and short phrases for a given topic.
- Agent2 must integrate with Microsoft Dynamics 365 Contact Center voice channel.

You need to recommend language models for the agents.

What should you recommend for each agent? To answer, drag the appropriate language models to the correct agents. Each language model may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

Drag sources

- Azure Language in Foundry Tools
- Azure OpenAI
- Conversational language understanding (CLU)
- Natural language understanding (NLU)
- Natural language understanding + (NLU+)

Drop targets

1. Agent1 - drop here
2. Agent2 - drop here

Answer

Agent1 = Natural language understanding (NLU); Agent2 = Natural language understanding + (NLU+)

Correct answer: - Agent1: **Natural language understanding (NLU)** - Agent2: **Natural language understanding + (NLU+)**

Why these answers are correct

The requirements are: - Standard model (no Azure dependency, no bring-your-own model) - No generative orchestration (so this is classic orchestration with the built-in NLU options) - Agent1: simple and short phrases per topic - Agent2: integration with Dynamics 365 Contact Center voice channel

Microsoft Learn's "Natural language understanding (NLU) overview" and "Design effective language understanding" articles describe the classic orchestration options:

NLU is the original built-in model. It is described as the option to use "if you want an easier programmable design or have simpler orchestration needs" and supports "5 to 20 short phrases per topic." That matches Agent1's requirement for simple, short phrases.

NLU+ is the high-accuracy enterprise option. Microsoft Learn states explicitly: "The NLU+ option is available when you manage your voice or chat channels with a Dynamics 365 Contact Center license," and "if you have a voice-enabled agent, your NLU+ training data is also used to optimize your speech recognition capabilities." That maps directly to Agent2's Dynamics 365 Contact Center voice channel requirement.

Why the other options are wrong

Azure Language in Foundry Tools is not one of the Copilot Studio language model options for an agent's NLU configuration; the Copilot Studio classic orchestration options are NLU, NLU+, and Azure CLU.

Azure OpenAI is associated with generative orchestration and generative answers, which the requirements explicitly exclude. It is also not a "standard model" per the question.

Conversational language understanding (CLU) requires an Azure subscription, ongoing model maintenance in Azure, and is described as the "bring-your-own NLU" option. It is not the standard model the requirements call for.

Microsoft Learn references

- Natural language understanding (NLU) overview - https://learn.microsoft.com/microsoft-copilot-studio/nlu-overview
- Design effective language understanding - https://learn.microsoft.com/microsoft-copilot-studio/guidance/language-understanding
- Configure NLU+ - https://learn.microsoft.com/microsoft-copilot-studio/nlu-plus-configure
- Voice-enabled agent overview - https://learn.microsoft.com/microsoft-copilot-studio/voice-overview

Topic 1 - Question 5

Multiple choice | Domain: Design

topic-1/q05

A company uses Microsoft Dynamics 365 finance and operations apps.

The company plans to use Microsoft Copilot in-app help and guidance to generate responses for internal business processes.

You need to add an additional knowledge source for the business processes. The solution must NOT add new topics to the Copilot agent for the finance and operations apps.

Which knowledge source should you add?

A. Microsoft Dataverse
B. a public website
C. Azure AI Search

D. a file upload

Answer

D

Correct answer: D. a file upload

Why D is correct

Microsoft Learn's guidance for extending Copilot in finance and operations apps ("Add knowledge to generative help and guidance with Copilot") states directly that you add additional knowledge by **uploading files** to the *Copilot for finance and operations apps* agent in Copilot Studio: "You can add knowledge to the Copilot help and guidance feature by uploading files to Copilot Studio (in file formats such as PDF, RTF, or Word)."

The article also explains why other knowledge source types are different: "If you want to add other types of knowledge (such as from SharePoint or other data sources), then you must add your own topic in the Copilot for Finance and Operations apps agent in Copilot Studio." That confirms file upload is the option that satisfies the constraint of not adding new topics.

The conversational boosting topic (out-of-the-box) handles uploaded files automatically, so no new topic is required. That meets the "must NOT add new topics" requirement and adds a usable knowledge source for internal business processes.

Why the other options are wrong

A. Microsoft Dataverse is supported as a knowledge source for the agent, but using Dataverse virtual entities published from finance and operations apps is explicitly called out on Microsoft Learn as not currently supported. Adding native Dataverse tables also typically involves more configuration effort and is geared toward grounded data, not generic business-process documentation.

B. a public website is a Bing/web grounding option. It is not how internal business process documentation is typically added, and it doesn't satisfy "internal business processes" since it pulls from public web data.

C. Azure AI Search is not the documented method to add knowledge to the Copilot for finance and operations apps agent without authoring a new topic. Adding Azure AI Search as a generative-answers source per Microsoft Learn requires custom topic configuration, which violates the "must NOT add new topics" requirement.

Microsoft Learn references

- Add knowledge to generative help and guidance with Copilot - https://learn.microsoft.com/dynamics365/fin-ops-core/dev-itpro/copilot/extend-copilot-generative-help
- Upload files as a knowledge source (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/knowledge-add-file-upload
- Generative help and guidance with Copilot - https://learn.microsoft.com/dynamics365/fin-ops-core/fin-ops/copilot/copilot-generative-help

Topic 1 - Question 6

Multiple choice | Domain: Design

topic-1/q06

A company has an AI business solution.

You need to extend the solution so that Microsoft 365 Copilot can invoke external logic hosted in Azure services.

What should you include in the solution?

A. Microsoft Copilot Studio skills
B. Microsoft Power Platform connectors
C. custom engine agents

Answer

A

Correct answer: A. Microsoft Copilot Studio skills

Why A is correct

The requirement is that **Microsoft 365 Copilot must invoke external logic hosted in Azure services**. Microsoft Copilot Studio skills are designed for exactly this scenario. Microsoft Learn's "Use skills in Copilot Studio" article describes skills as a way to extend an agent by embedding bots or services built with pro-code tools (including the Microsoft 365 Agents SDK and Bot Framework, hosted in Azure) into the agent. A skill becomes a callable capability the agent can invoke from a topic, and the Copilot Studio agent can then be published as a Microsoft 365 Copilot agent that surfaces that skill to Microsoft 365 Copilot.

This pattern lets Microsoft 365 Copilot reach Azure-hosted code (a registered skill) through a Copilot Studio agent, which is how Microsoft documents extending Microsoft 365 Copilot with external logic from Azure services.

Why the other options are wrong

B. Microsoft Power Platform connectors integrate with REST APIs and SaaS services. They are useful for data and action calls, but the option that specifically describes extending Copilot agents with externally hosted bot/agent logic (the typical pattern for "external logic in Azure services") is skills. Connectors are a more generic integration mechanism and not the canonical answer for invoking external Azure-hosted logic from Microsoft 365 Copilot.

C. Custom engine agents are built using the Microsoft 365 Agents SDK (or similar) and replace Microsoft 365 Copilot's built-in orchestration with your own engine. They are a different extensibility model and aren't the mechanism Microsoft Learn describes for letting Microsoft 365 Copilot invoke external Azure-hosted logic from within the standard Copilot orchestration.

Microsoft Learn references

- Use skills in Copilot Studio - https://learn.microsoft.com/microsoft-copilot-studio/advanced-use-skills
- Microsoft 365 Copilot extensibility overview - https://learn.microsoft.com/microsoft-365-copilot/extensibility/
- Skills (Bot Framework / Microsoft 365 Agents SDK) - https://learn.microsoft.com/azure/bot-service/skill-pva

Topic 1 - Question 7

Hotspot | Domain: Plan

topic-1/q07

You need to design a shared prompt library that will be used across multiple business units. The solution must meet the following requirements:

- Ensure consistent AI responses with reusable formats.
- Support governance and version control.
- Minimize administrative effort.
- Minimize ongoing costs.

What should you recommend for each requirement? To answer, select the appropriate options in the answer area.

Answer area

- Ensure consistent AI responses:

[ ] Delegate department-specific prompt templates.
[ ] Define standardized prompt templates.
[ ] Maintain a prompt history.

- Support governance and version control:

[ ] Define standardized prompt templates.
[ ] Store prompts in a Git repository.
[ ] Categorize prompts by business function.

Answer

Ensure consistent AI responses = Define standardized prompt templates; Support governance and version control = Store prompts in a Git repository

Correct answer: - Ensure consistent AI responses: **Define standardized prompt templates.** - Support governance and version control: **Store prompts in a Git repository.**

Why these answers are correct

Standardized prompt templates are the documented Microsoft pattern for getting consistent, reusable AI responses across business units. Microsoft Learn's "Get started with prompt library" article describes prompt libraries as "a collection of predesigned prompts, serving as templates to expedite the creation of AI prompts," with the explicit goal of accelerating development "and ensures best practices are followed in prompt engineering." Standardization at the template level is what produces consistent responses with reusable formats.

Git repository storage is the standard answer for prompt governance and version control. Microsoft Learn's "AI apps and agents publication templates" article describes a prompt lifecycle template (Draft, Review, Approved, Deployed) and recommends documenting changes in source control or a wiki for traceability. Git is Microsoft's recommended version control system (Azure Repos, GitHub) and provides branching, pull request review, and history, which are the core requirements for governance and version control. Categorization by business function does not provide version control, and standardized templates by themselves do not provide change history or approval workflows.

This combination also minimizes administrative effort and ongoing cost. Templates are authored once and reused, and Git is a well-understood developer practice that the organization likely already uses.

Why the other options are wrong

Delegate department-specific prompt templates would produce divergent prompts across business units, defeating the goal of consistent responses.

Maintain a prompt history captures usage but does not enforce a consistent response format.

Define standardized prompt templates (offered for the second slot) addresses consistency, not version control. Without Git or another source-control system, you have no diff history, branching, or rollback.

Categorize prompts by business function improves discoverability but offers no governance or versioning.

Microsoft Learn references

- Get started with prompt library - https://learn.microsoft.com/microsoft-copilot-studio/prompt-library
- AI apps and agents publication templates (prompt lifecycle) - https://learn.microsoft.com/partner-center/marketplace-offers/artificial-intelligence-templates#prompt-lifecycle-template
- What is Azure Repos / Git - https://learn.microsoft.com/azure/devops/repos/get-started/what-is-repos
- Git workflow - https://learn.microsoft.com/azure/devops/repos/git/gitworkflow

Topic 1 - Question 8

Drag and drop | Domain: Plan

topic-1/q08

A company has a Microsoft Foundry project that uses a single agent and a single prompt to complete a series of tasks.

The agent encounters the following issues:

- It frequently produces incomplete results.
- It struggles with domain-specific reasoning.
- Agent response times are remarkably slow.

You need to recommend a solution to improve the overall performance and accuracy of the agent.

What should you include in the recommendation? To answer, drag the appropriate actions to the correct requirements. Each action may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

Drag sources

- Add a grounding data source.
- Add a prebuilt connector.
- Move to a multi-agent architecture.
- Upgrade to a larger generative AI model.

Drop targets

1. To improve performance - drop here
2. To improve accuracy - drop here

Answer

Performance = Move to a multi-agent architecture; Accuracy = Add a grounding data source

Correct answer: - To improve performance: **Move to a multi-agent architecture.** - To improve accuracy: **Add a grounding data source.**

Why these answers are correct

The question describes three symptoms in a single-agent, single-prompt design: - Frequently produces incomplete results - Struggles with domain-specific reasoning - Slow response times

Microsoft Learn's "Transparency Note for Azure Agent Service" calls this out directly: "When a single agent's system message consistently struggles to handle the complexity, breadth, or depth of a task, such as frequently producing incomplete results, hitting reasoning bottlenecks, or requiring extensive domain-specific knowledge, the system may benefit from transitioning to a multi-agent architecture." Splitting the workload into specialized subtasks reduces cognitive overload on any single agent, which addresses both the slow response times and the incomplete results. That maps to the **performance** improvement.

For **accuracy**, Microsoft Learn's grounding data design guidance and the Azure AI Foundry agent transparency note both point to grounding as the lever for accuracy. Adding a grounding data source gives the model authoritative reference content to cite at inference time, which directly addresses domain-specific reasoning gaps and reduces hallucinated or incomplete answers. Grounding is described in Microsoft Learn as the way to "augment generative AI models by using context data during inference" so the model "can answer the user query and form the answer according to the expectations."

Why the other options are wrong (for the slot they don't fit)

Add a prebuilt connector integrates a specific external system but doesn't by itself fix domain reasoning or speed.

Upgrade to a larger generative AI model can sometimes improve accuracy and reasoning, but it usually increases latency and cost rather than improving performance, and Microsoft's published guidance for the symptoms listed (incomplete results, domain reasoning, slow response) points to multi-agent decomposition and grounding before model upgrades.

Microsoft Learn references

- Transparency Note for Azure Agent Service (single vs. multi-agent guidance) - https://learn.microsoft.com/azure/foundry/responsible-ai/agents/transparency-note
- Single agent or multiple agents - https://learn.microsoft.com/azure/cloud-adoption-framework/ai-agents/single-agent-multiple-agents
- Grounding data design for AI workloads on Azure - https://learn.microsoft.com/azure/well-architected/ai/grounding-data-design
- Knowledge sources in Copilot Studio - https://learn.microsoft.com/microsoft-copilot-studio/knowledge-copilot-studio

Topic 1 - Question 9

Multiple choice | Domain: Design

topic-1/q09

A financial services company uses Microsoft Dynamics 365 Finance.

Currently, the company's support staff manually reviews customer transaction histories to detect potential fraud cases before escalating the cases.

You need to recommend an automation solution for the review process. The solution must ensure that escalations reach a human analyst for final decision making. What should you recommend?

A. Deploy an autonomous agent that closes non-fraud cases automatically.
B. Use Microsoft 365 Copilot in Word to automatically finalize fraud detection policies.
C. Configure a task agent to generate fraud risk scores for the human analyst to review.
D. Export the data to a data lake for analysis in Microsoft Power BI.

Answer

C

Correct answer: C. Configure a task agent to generate fraud risk scores for the human analyst to review.

Why C is correct

The requirement has two parts: 1. Automate the review process for customer transaction histories. 2. Ensure escalations reach a **human analyst for final decision making**.

A task agent that generates fraud risk scores fits both parts. The agent takes a defined task (analyze the transaction history, output a risk score), the human analyst reviews the score, and the analyst still owns the final escalation decision. This pattern is the human-in-the-loop design Microsoft Learn recommends for sensitive financial scenarios. The Azure Cloud Adoption Framework's AI agent guidance describes task agents as agents that "perform specific tasks within defined workflows" with knowledge plus action tools, and recommends keeping autonomous decision-making out of high-stakes financial processes.

The Microsoft Copilot Studio responsible AI guidance and the Copilot Studio anomaly detection solution idea both stress that for finance fraud-style work, the agent should surface findings (with severity scores and recommendations) and route them to a human reviewer rather than auto-resolving cases.

Why the other options are wrong

A. Deploy an autonomous agent that closes non-fraud cases automatically removes the human from cases the agent decides are not fraud. That violates the requirement that escalations reach a human analyst for final decision making, and it conflicts with Microsoft's responsible AI guidance for high-impact financial actions.

B. Use Microsoft 365 Copilot in Word to automatically finalize fraud detection policies is unrelated. Word Copilot drafts documents; it doesn't review transactions or generate risk scores.

D. Export the data to a data lake for analysis in Microsoft Power BI gives the analyst dashboards but no automation of the review process. Power BI is reporting, not an agent.

Microsoft Learn references

- AI agent types (task agents) - https://learn.microsoft.com/azure/cloud-adoption-framework/ai-agents/
- Build anomaly detection with Copilot Studio and Fabric - https://learn.microsoft.com/power-platform/architecture/solution-ideas/agent-anomaly-detection
- Responsible AI principles for Copilot Studio - https://learn.microsoft.com/microsoft-copilot-studio/guidance/responsible-ai
- Best practices for integrating and deploying Copilot Studio (human oversight) - https://learn.microsoft.com/microsoft-copilot-studio/system-service-card-copilot-studio

Topic 1 - Question 10

Multiple choice | Domain: Plan

topic-1/q10

A company plans to deploy a Microsoft Copilot Studio agent that will analyze historical business data to predict customer behavior.

The data is currently stored in an Azure SQL database, flat files, APIs, and logs.

You need to organize the data into a format that can be used as a knowledge source in Copilot Studio.

What should you include in the solution?

A. Azure AI Search
B. Azure Data Lake Storage
C. Azure Cosmos DB
D. Azure Translator in Foundry Tools

Answer

B

Correct answer: B. Azure Data Lake Storage

Why B is correct

The requirement is to consolidate data that currently lives in multiple, mixed sources (Azure SQL, flat files, APIs, logs) into a single organized format that can serve as a knowledge source in Copilot Studio for predictive analysis on historical business data.

Azure Data Lake Storage is the Microsoft-recommended landing place for that kind of mixed-source historical data. Microsoft Learn's grounding data design guidance describes consolidating structured, semi-structured, and unstructured data into a centralized data lake (or warehouse) so that it can be efficiently indexed, transformed, and queried for AI workloads. Azure Data Lake Storage also integrates with Microsoft Fabric, Azure Synapse, and downstream tools that can prepare the data for use as a knowledge source. Power Platform / Copilot Studio can then connect to that consolidated data (for example via a Fabric Data Agent connection or through Dataverse-mediated patterns) for grounding.

For predictive analysis over historical business data drawn from heterogeneous sources, the lake is the consolidation layer. After data lands there, it can be processed and surfaced to Copilot Studio.

Why the other options are wrong

A. Azure AI Search is a retrieval/index service, not a consolidation layer for raw mixed sources. It is a downstream component used for grounding once data is prepared. It doesn't ingest and organize SQL, flat files, APIs, and logs by itself.

C. Azure Cosmos DB is an operational NoSQL database optimized for low-latency transactional workloads. It is not the recommended target for consolidating diverse historical analytics data.

D. Azure Translator in Foundry Tools is a translation service. It has nothing to do with consolidating historical business data.

Microsoft Learn references

- Grounding data design for AI workloads on Azure - https://learn.microsoft.com/azure/well-architected/ai/grounding-data-design
- Azure Data Lake Storage introduction - https://learn.microsoft.com/azure/storage/blobs/data-lake-storage-introduction
- Microsoft Fabric Lakehouse overview - https://learn.microsoft.com/fabric/data-engineering/lakehouse-overview
- Knowledge sources in Copilot Studio - https://learn.microsoft.com/microsoft-copilot-studio/knowledge-copilot-studio

Topic 1 - Question 11

Multiple choice | Domain: Plan

topic-1/q11

A retail company plans to deploy Microsoft Copilot Studio agents to support:

- Microsoft Dynamics 365 Commerce scenarios.
- A Microsoft Power Apps inventory management solution.

You need to recommend a solution to organize product catalog data as a consistent source for multiple AI systems.

What should you recommend?

A. Let each agent scrape product details from Microsoft SharePoint Online libraries.
B. Store the product catalog data in a separate custom table for each agent.
C. Configure prompts to pull product details from the PDFs of external vendors.
D. Centralize the product catalog data in Microsoft Dataverse and expose the data to both agents.

Answer

D

Correct answer: D. Centralize the product catalog data in Microsoft Dataverse and expose the data to both agents.

Why D is correct

The scenario has two AI systems (a Copilot Studio agent for Dynamics 365 Commerce, and a Power Apps inventory management solution) that both need a consistent view of product catalog data. Microsoft's documented pattern is to centralize that kind of master data in Microsoft Dataverse and use it as a single source of truth.

Microsoft Learn ("Why choose Microsoft Dataverse?") describes Dataverse as the platform's central data store, with built-in security, native integration with Power Apps, Power Automate, Dynamics 365, and Copilot Studio, and AI-ready features (Dataverse search, glossary terms, synonyms) that support generative AI grounding. The Power Platform Well-Architected guidance for Copilot Studio agents recommends consolidating data on a centralized platform such as Dataverse to streamline access and management for multiple agents.

Centralization also avoids data drift and inconsistent answers across the two AI systems, which is the core of "consistent source for multiple AI systems."

Why the other options are wrong

A. Let each agent scrape product details from Microsoft SharePoint Online libraries produces inconsistent results, depends on document layout, scales poorly, and violates the "consistent source" requirement.

B. Store the product catalog data in a separate custom table for each agent duplicates data, creates synchronization problems, and directly contradicts the requirement for a single consistent source.

C. Configure prompts to pull product details from the PDFs of external vendors routes catalog data through external documents, creates lag and accuracy problems, and is not a centralized authoritative source.

Microsoft Learn references

- Why choose Microsoft Dataverse? - https://learn.microsoft.com/power-apps/maker/data-platform/why-dataverse-overview
- Add Dataverse tables as a knowledge source - https://learn.microsoft.com/microsoft-copilot-studio/knowledge-add-dataverse
- Power Platform Well-Architected: Experience Optimization (consolidate data on Dataverse) - https://learn.microsoft.com/power-platform/architecture/solution-ideas/agent-custom-contact-center

Topic 1 - Question 12

Multiple choice | Domain: Plan

topic-1/q12

A company has a portfolio of AI initiatives at different stages of development.

You need to recommend a structured approach to evaluating the return on AI investment (ROAI) across all the initiatives. The solution must balance immediate results with long-term values and strategic innovations.

What should you include in the recommendation?

A. a simple cost and benefit analysis

B. a horizon-based framework
C. the internal rate of return (IRR) function
D. a prioritization grid

Answer

B

Correct answer: B. a horizon-based framework

Why B is correct

The requirement is to evaluate ROAI across a portfolio of AI initiatives at different stages of development, balancing immediate results, long-term value, and strategic innovations. That maps directly to a horizon-based framework.

Microsoft Learn's Power Platform "Define roles and responsibilities" article describes horizon mapping as a technique to "visualize and prioritize tasks across different time horizons" with three horizons: Horizon 1 (immediate), Horizon 2 (near-term), and Horizon 3 (long-term and strategic). The "Pillar 1: AI strategy and experience" guidance similarly describes sequencing initiatives "across near, mid, and longer term horizons" to balance short-term wins with strategic innovation. That is the exact balance the question asks for.

A horizon-based framework lets the organization classify each AI initiative by where it sits on the value timeline, then evaluate ROAI in a way that doesn't penalize long-horizon strategic bets and doesn't over-credit short-term efficiency wins.

Why the other options are wrong

A. a simple cost and benefit analysis captures one snapshot of cost vs. expected benefit. It doesn't structure a portfolio across time horizons or distinguish strategic innovation from immediate results.

C. the internal rate of return (IRR) function is a single financial metric. It can be one input but doesn't, by itself, structure a portfolio evaluation across short, medium, and long-term horizons.

D. a prioritization grid is useful for ranking individual initiatives by impact vs. effort but doesn't capture the time-horizon dimension Microsoft uses to balance immediate results against strategic innovation.

Microsoft Learn references

- Define roles and responsibilities (Horizon mapping) - https://learn.microsoft.com/power-platform/guidance/adoption/roles
- Pillar 1: AI strategy and experience (multi-horizon strategy) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/maturity-model-strategy
- Plan for AI adoption (prioritize use cases / ROAI) - https://learn.microsoft.com/azure/cloud-adoption-framework/ai/plan

Topic 1 - Question 13

Multiple choice | Domain: Plan

topic-1/q13

You need to recommend a Microsoft Power Platform business solution that consolidates data from multiple internal and external data sources. The solution must meet the following requirements:

- Provide the data as a centralized source for multiple AI systems, including Microsoft Copilot Studio agents, Dynamics 365 applications, and external AI models.
- Support built-in data classification and protection policies.
- Provide data for grounding and analytics.

What should you include in the recommendation?

A. Microsoft Dataverse
B. Azure Data Lake Storage
C. a Microsoft Power BI semantic model
D. Azure Cosmos DB

Answer

A

Correct answer: A. Microsoft Dataverse

Why A is correct

The requirement set is: 1. Centralized source for multiple AI systems (Copilot Studio agents, Dynamics 365 apps, external AI models). 2. Built-in data classification and protection policies. 3. Data for grounding and analytics.

Microsoft Dataverse is the only option that meets all three:

- It is the standard data backbone for Copilot Studio agents and Dynamics 365 apps. Microsoft Learn's "Why choose Microsoft Dataverse?" article describes it as the central data store with native integration to Power Apps, Power Automate, Dynamics 365, and Copilot Studio. External AI systems can also reach Dataverse data via APIs, virtual tables, MCP servers, and Microsoft Graph.
- It includes built-in data classification, security roles, row/column-level security, and data policies. Microsoft Purview integrates with Dataverse for sensitivity labels, DLP for Dataverse, and data classification.
- It supports grounding (it is a documented Copilot Studio knowledge source) and analytics (managed data lake export to Azure Data Lake / Fabric, Power BI integration).

This is the documented Microsoft Power Platform business solution choice for the requirements stated.

Why the other options are wrong

B. Azure Data Lake Storage is a data store for analytics but lacks the built-in business data classification, protection policies, and Power Platform-native integrations required to be a centralized source for Copilot Studio, Dynamics 365, and external AI models out of the box.

C. a Microsoft Power BI semantic model is a reporting layer over data, designed for analytics. It is not the centralized authoritative business data store, and it does not provide built-in data classification and protection policies in the way required across multiple AI systems.

D. Azure Cosmos DB is a globally distributed NoSQL database. It is not the Power Platform data backbone and lacks the integrated classification, sensitivity, and AI-ready features described for Dataverse.

Microsoft Learn references

- Why choose Microsoft Dataverse? - https://learn.microsoft.com/power-apps/maker/data-platform/why-dataverse-overview
- Microsoft Purview data classification (Dataverse) - https://learn.microsoft.com/purview/ai-microsoft-purview
- Add Dataverse tables as a knowledge source - https://learn.microsoft.com/microsoft-copilot-studio/knowledge-add-dataverse
- Step 2: Protect sensitive info in grounding data (Dataverse) - https://learn.microsoft.com/purview/deploymentmodels/depmod-sc-agents-step2

Topic 1 - Question 14

Multiple choice | Domain: Plan

topic-1/q14

A company plans to deploy an AI-based customer service app that will autonomously manage interactions, escalate complex cases, and learn from historical ticket data.

You need to perform a return on AI investment (ROAI) analysis of the app deployment. The solution must ensure that the analysis is accurate.

What should you do first?

A. Establish the AI performance metrics.
B. Conduct an AI market benchmarking study.
C. Model the customer experience.
D. Identify and quantify all the development, deployment, and operating costs.

Answer

D

Correct answer: D. Identify and quantify all the development, deployment, and operating costs.

Why D is correct

ROAI is, at its core, benefits relative to costs. To make the analysis accurate, you have to establish the cost baseline first. Microsoft Learn's "Plan to manage costs for Azure OpenAI" and "Quantify business value" guidance both put cost identification at the start of any ROI / ROAI analysis. The Azure Cloud Adoption Framework's "Business plan for AI agents" and "Plan for AI adoption" articles also describe baseline costs (development, deployment, operating) as the starting point for return-on-investment analysis, then comparing those costs against measurable benefits.

Without identifying and quantifying full costs (people, model usage, infrastructure, integration, support, retraining, governance), any subsequent metric or benchmark is built on an incomplete denominator. That is why the **first** step is identifying and quantifying all costs.

Why the other options are wrong

A. Establish the AI performance metrics is necessary later in the analysis, but you can't relate performance to ROAI until you know the cost basis you are measuring performance against.

B. Conduct an AI market benchmarking study is useful context but doesn't tell you anything about your own deployment's economics. ROAI accuracy depends on your costs and your benefits, not industry averages.

C. Model the customer experience informs benefits estimation but is not the first step. You need a cost baseline first.

Microsoft Learn references

- Business plan for AI agents (define success metrics, decision gates) - https://learn.microsoft.com/azure/cloud-adoption-framework/ai-agents/business-strategy-plan
- Plan to manage costs for Azure OpenAI - https://learn.microsoft.com/azure/ai-foundry/openai/how-to/manage-costs
- Quantify business value (FinOps) - https://learn.microsoft.com/cloud-computing/finops/framework/quantify/quantify-business-value
- Plan for AI adoption - https://learn.microsoft.com/azure/cloud-adoption-framework/ai/plan

Topic 1 - Question 15

Drag and drop | Domain: Deploy

topic-1/q15

You are designing end-to-end test scenarios for a business solution that uses Microsoft Dynamics 365 Sales and Dynamics 365 Finance.

You need to ensure that the business solution meets the following test requirements:

- Properly exchanges data between the Dynamics 365 apps
- Aligns with defined user workflows and business processes

Which type of testing should you use for each requirement? To answer, drag the appropriate testing types to the correct requirements. Each testing type may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

Drag sources

- Drift
- Exploratory
- Integration
- Performance
- User acceptance

Drop targets

1. Properly exchanges data between the Dynamics 365 apps - drop here
2. Aligns with defined user workflows and business processes - drop here

Answer

Properly exchanges data between the Dynamics 365 apps = Integration; Aligns with defined user workflows and business processes = User acceptance

Correct answer: - Properly exchanges data between the Dynamics 365 apps: **Integration** - Aligns with defined user workflows and business processes: **User acceptance**

Why these answers are correct

Microsoft Learn's "Types of tests" article and the broader Dynamics 365 testing strategy guidance map test types directly to these requirements:

Integration testing is described as "testing how different components or systems work together, such as data flows, interfaces, and APIs." That is precisely the test type for verifying that Dynamics 365 Sales and Dynamics 365 Finance correctly exchange data.

User acceptance testing (UAT) is described as "testing the solution from the user's perspective" and is performed by business users who run real-life process simulations to confirm the solution supports business operations and goals. UAT is the documented test type for validating alignment with user workflows and business processes, and it's a prerequisite for go-live in Success by Design.

Why the other options are wrong (for the slot they don't fit)

Drift typically refers to model drift or configuration drift monitoring, not functional testing of data exchange or workflow alignment.

Exploratory testing finds issues through ad hoc investigation. It is not the structured method for confirming data exchange between two business apps or aligning a solution to defined workflows.

Performance testing measures speed, throughput, and scalability. It does not confirm functional data exchange or workflow alignment.

Microsoft Learn references

- Types of tests - https://learn.microsoft.com/dynamics365/guidance/implementation-guide/testing-strategy-test-types
- Test your Dynamics 365 solution before deployment - https://learn.microsoft.com/dynamics365/guidance/implementation-guide/testing-strategy
- Perform user acceptance testing in finance and operations apps (training) - https://learn.microsoft.com/training/modules/perform-uat-finance-operations/

Topic 1 - Question 16

Multiple choice | Domain: Deploy

topic-1/q16

A company has a Microsoft 365 tenant in Canada and multiple Microsoft Power Platform environments in Canada and the United States.

The company plans to deploy a Microsoft Copilot Studio agent to the Canadian environment that will use:

- Microsoft Dataverse data stored in Canada
- A connector that connects to an Azure OpenAI instance in the United States

You need to ensure that the agent adheres to data residency and data movement policies before being deployed.

What should you do?

A. Ensure that the data processed by Azure OpenAI is stored in the United States.
B. From the Microsoft Purview portal, validate the Data loss prevention settings.
C. Migrate the tenant to the United States.
D. Ensure that cross-region data movement is enabled for the Canadian environment and connector dependencies.

Answer

D

Correct answer: D. Ensure that cross-region data movement is enabled for the Canadian environment and connector dependencies.

Why D is correct

The agent will live in a Canadian environment using Canadian Dataverse data, but it depends on an Azure OpenAI connector hosted in the United States. Microsoft Learn's "Move data across regions for Copilots, AI agents, and generative AI features" article is explicit:

"When the **Move data across regions** checkbox is selected, your inputs (prompts) and outputs (results) might move outside of your region to the location where the generative AI feature is hosted. ... Some regions may have limited capacity, or no capacity at all. To ensure availability of Copilots and generative AI features, we may need to move the data outside of the region for processing."

For the Canada region, the table shows that data may be processed in the United States for Azure OpenAI. The administrator must explicitly enable the "Move data across regions" setting in the Power Platform admin center for the environment, and the related connector dependencies must support that flow. Without it, the agent's cross-region path is blocked or unsupported, and the data residency policy is not addressed.

This is the documented Microsoft administrative step for adhering to data residency and data movement policies in this exact scenario before the agent is deployed.

Why the other options are wrong

A. Ensure that the data processed by Azure OpenAI is stored in the United States is not the residency choice you make at the agent layer. Microsoft Learn states explicitly that "Microsoft doesn't log, store, or retain any input or output data during this process. There is no persistence of the data." You don't configure US storage for processed prompts; you control whether data movement is allowed.

B. From the Microsoft Purview portal, validate the Data loss prevention settings is good practice for sensitive data, but DLP validation alone does not address the cross-region data movement consent requirement. The Power Platform admin center is the location for that consent, not Purview.

C. Migrate the tenant to the United States is disproportionate, breaks Canadian data residency, and is not a Microsoft-recommended response to a US-hosted connector dependency.

Microsoft Learn references

- Move data across regions for Copilots, AI agents, and generative AI features - https://learn.microsoft.com/power-platform/admin/geographical-availability-copilot
- Configure data movement across geographic locations for generative AI - https://learn.microsoft.com/microsoft-copilot-studio/manage-data-movement-outside-us
- Geographic data residency in Copilot Studio - https://learn.microsoft.com/microsoft-copilot-studio/geo-data-residency

Topic 1 - Question 17

Hotspot | Domain: Deploy

topic-1/q17

A company has a Microsoft Copilot Studio agent for customer support.

You are reviewing and validating the following prompts:

- A prompt that has instructions to "help the customer as best you can"
- A prompt that helps retrieve product information from a knowledge base

You need to ensure that the agent delivers consistent and accurate responses.

What should you do for each prompt? To answer, select the appropriate options in the answer area.

Answer area

- A prompt that has instructions to "help the customer as best you can":
  [ ] Add filler words to make the prompt sound more natural and conversational.
  [ ] Keep the prompt vague to enable model flexibility.
  [ ] Rewrite the prompt with clear and task-specific instructions.

- A prompt that helps retrieve product information from a knowledge base:
  [ ] Add several open-ended questions to give the model broader context.
  [ ] Use responses with only reference sources and limit the response scope.
  [ ] Remove the knowledge source so that the model responds freely with general product information.

Answer

Vague prompt = Rewrite the prompt with clear and task-specific instructions; Knowledge-base prompt = Use responses with only reference sources and limit the response scope

Correct answer: - A prompt that has instructions to "help the customer as best you can": **Rewrite the**

prompt with clear and task-specific instructions. - A prompt that helps retrieve product information from a knowledge base: **Use responses with only reference sources and limit the response scope.**

Why these answers are correct

Microsoft Learn's prompt engineering guidance ("Prompt engineering techniques," "Use prompts to make your agent or agent flow perform specific tasks," "Use prompt modification to provide custom instructions to your agent") states the same principle repeatedly: be specific, give task-specific instructions, and start with clear instructions. Vague prompts like "help the customer as best you can" produce inconsistent and unreliable answers because the model has no concrete task definition. The correct fix is to rewrite the prompt with clear, task-specific instructions.

For the knowledge-base retrieval prompt, the goal is grounded, accurate responses based on the source documents. Microsoft Learn's "Best practices for integrating and deploying Copilot Studio" calls out: "Avoid relying solely on model-generated content for critical decisions, always ground outputs in verifiable data." The right pattern for knowledge-base retrieval is to constrain the response to reference sources and limit the response scope. That reduces hallucination and keeps answers traceable through citations.

Why the other options are wrong

Add filler words to make the prompt sound more natural and conversational introduces noise without adding task clarity. Microsoft's guidance is the opposite: keep prompts brief and specific.

Keep the prompt vague to enable model flexibility produces inconsistent responses. Microsoft's prompt engineering guidance explicitly warns against vague or ambiguous language.

Add several open-ended questions to give the model broader context widens the response scope, which is the opposite of what's needed for accurate knowledge-base retrieval.

Remove the knowledge source so that the model responds freely with general product information removes grounding entirely, which is the worst option for accuracy and consistency.

Microsoft Learn references

- Prompt engineering techniques - https://learn.microsoft.com/azure/ai-foundry/openai/concepts/prompt-engineering
- Use prompts to make your agent or agent flow perform specific tasks - https://learn.microsoft.com/microsoft-copilot-studio/nlu-prompt-node
- Use prompt modification to provide custom instructions to your agent - https://learn.microsoft.com/microsoft-copilot-studio/nlu-generative-answers-prompt-modification
- Best practices for integrating and deploying Copilot Studio - https://learn.microsoft.com/microsoft-copilot-studio/system-service-card-copilot-studio

Topic 1 - Question 18

Hotspot | Domain: Deploy

topic-1/q18

You are designing a testing solution for Microsoft Copilot Studio agents.

You need to validate prompt engineering best practices to ensure that the agents generate accurate and contextually relevant responses.

Which prompt validation techniques and metrics should you include in the solution? To answer, select the appropriate options in the answer area.

Answer area

- Prompt validation techniques:

[ ] Exclude domain-specific terminology from the prompts.
[ ] Use prompts that have varied phrasing.

[ ] Use only simple, one-word prompts.

- Metrics:

[ ] The number of words generated per response
[ ] Response relevance and accuracy
[ ] The response generation time

Answer

Prompt validation = Use prompts that have varied phrasing; Metrics = Response relevance and accuracy

Correct answer: - Prompt validation techniques: **Use prompts that have varied phrasing.** - Metrics: **Response relevance and accuracy**

Why these answers are correct

For prompt validation, Microsoft's prompt engineering guidance and Copilot Studio testing guidance both stress that you should test against varied phrasings of the same intent. Real users won't all phrase a question the same way, so validating with varied phrasing is how you confirm the agent generalizes. Microsoft Learn's "Test your agent" and "Best practices for integrating and deploying Copilot Studio" sections call out simulating real-world scenarios and varied user inputs as the way to verify the agent generates accurate, contextually relevant responses.

For the metric, Microsoft Learn's evaluation overview and Copilot Studio analytics guidance describe response relevance and accuracy as the primary qualitative metrics for prompt engineering testing. The Copilot Studio Application Card explicitly lists evaluation as the practice of "assessing agent quality, performance, safety, and reliability," with response quality (relevance and accuracy) at the center. Counting words generated or generation time tells you nothing about whether the response was correct or contextually useful.

Why the other options are wrong

Exclude domain-specific terminology from the prompts would reduce realism and break domain-specific reasoning testing. Real prompts in production use domain terminology, so testing should include it, not exclude it.

Use only simple, one-word prompts is unrealistic and doesn't test how the agent handles real user input.

The number of words generated per response is not a quality signal. A long answer can be wrong; a short answer can be correct.

The response generation time is a performance metric, not a metric for prompt engineering accuracy. You measure latency separately from response correctness.

Microsoft Learn references

- Prompt engineering techniques - https://learn.microsoft.com/azure/ai-foundry/openai/concepts/prompt-engineering
- Test your agent - https://learn.microsoft.com/microsoft-copilot-studio/authoring-test-bot
- Application Card: Microsoft Copilot Studio (Evaluation) - https://learn.microsoft.com/microsoft-copilot-studio/system-service-card-copilot-studio
- Evaluate generative AI applications - https://learn.microsoft.com/azure/ai-foundry/concepts/evaluation-approach-gen-ai

Topic 1 - Question 19

Multiple choice | Domain: Deploy

topic-1/q19

A company has two Microsoft Power Platform environments named Dev1 and Prod1. A Microsoft Copilot Studio agent named Agent1 is built into a solution in the Dev1 environment.

You plan to deploy Agent1 to Prod1.

You need to make Agent1 available to the users in Prod1. The solution must minimize administrative effort.

What should you do?

A. Share Agent1 with the users in Prod1.
B. Export the solution as an unmanaged solution and import the solution into Prod1.
C. Export the solution as a managed solution and import the solution into Prod1.
D. Create a new Copilot Studio agent in Prod1 by replicating the configuration of Agent1.

Answer

C

Correct answer: C. Export the solution as a managed solution and import the solution into Prod1.

Why C is correct

Microsoft's documented application lifecycle management (ALM) pattern for moving Copilot Studio agents between environments is to package the agent in a solution, export the solution as **managed**, and import it into the target environment.

Microsoft Learn's "Export and import agents using solutions" and "Create and manage solutions in Copilot Studio" articles are explicit on this. Managed solutions are the recommended deployment artifact for production environments because: - They prevent direct edits in production (preserving the integrity of the deployment). - They support upgrade and update flows back from Dev1. - They are the artifact type expected by Power Platform pipelines.

This minimizes administrative effort because once the managed solution is imported into Prod1, the agent and all its solution components (topics, tools, environment variables, etc.) are deployed in one operation, and future updates flow as solution upgrades rather than manual configuration.

Why the other options are wrong

A. Share Agent1 with the users in Prod1 does not move the agent to Prod1. Sharing extends access within the source environment; it doesn't deploy the agent to a different environment.

B. Export the solution as an unmanaged solution and import the solution into Prod1 is technically possible but is not the recommended pattern for production. Unmanaged solutions allow direct edits in the target environment, are harder to maintain across environments, and Microsoft's guidance is explicit that production should use managed solutions.

D. Create a new Copilot Studio agent in Prod1 by replicating the configuration of Agent1 is the most administrative effort possible (full manual rebuild) and creates drift between environments. It directly violates the "minimize administrative effort" requirement.

Microsoft Learn references

- Export and import agents using solutions - https://learn.microsoft.com/microsoft-copilot-studio/authoring-solutions-import-export
- Create and manage solutions in Copilot Studio -

https://learn.microsoft.com/microsoft-copilot-studio/authoring-solutions-overview
- Managed and unmanaged solutions - https://learn.microsoft.com/power-platform/alm/solution-concepts-alm
- Use application lifecycle management in Copilot Studio - https://learn.microsoft.com/power-platform/release-plan/2024wave2/data-platform/use-application-lifecycle-management-copilot-studio

Topic 1 - Question 20

Multiple choice | Domain: Deploy

topic-1/q20

A company has a Microsoft Power Platform environment that contains Microsoft Dataverse data.

You create a Microsoft Copilot Studio agent named Agent1 that processes the Dataverse data.

You discover that Agent1 fails to return relevant or accurate results.

You need to improve the quality and reliability of data grounding.

What should you do?

A. Retrain Agent1.
B. Verify and cleanse the Dataverse data.
C. Use an adaptive card in Agent1.
D. Add example user inputs to the training data of Agent1.

Answer

B

Correct answer: B. Verify and cleanse the Dataverse data.

Why B is correct

The agent fails to return relevant or accurate results, and the goal is to improve the quality and reliability of **data grounding**. Grounding quality is a function of the source data quality. Microsoft Learn's Power Platform Well-Architected guidance for Copilot Studio agents states this directly: "The principle of 'garbage in, garbage out' is particularly important for agents and emphasizes the need for high-quality data. Providing accurate information to the agent ensures reliable and correct responses."

Microsoft Learn's "Improve copilot responses from Microsoft Dataverse" article also confirms that the lever for accuracy when an agent grounds in Dataverse is the data itself (cleanliness, descriptions, synonyms, glossary terms). Verifying and cleansing the Dataverse data is the targeted action that improves grounding quality, which directly improves the agent's relevance and accuracy.

Why the other options are wrong

A. Retrain Agent1 is not how Copilot Studio agents work for grounded data. Copilot Studio agents don't get "retrained" on Dataverse data. They retrieve grounded answers from the data sources at runtime, so retraining isn't the lever for grounding quality.

C. Use an adaptive card in Agent1 changes how responses are presented to users. It doesn't change the data the agent grounds in.

D. Add example user inputs to the training data of Agent1 improves intent recognition for classic-orchestration trigger phrases. It doesn't fix grounding quality from Dataverse.

Microsoft Learn references

- Improve copilot responses from Microsoft Dataverse - https://learn.microsoft.com/power-apps/maker/data-platform/data-platform-copilot
- Custom contact center solution with Copilot Studio agent (Reliability: garbage in, garbage out) - https://learn.microsoft.com/power-platform/architecture/solution-ideas/agent-custom-contact-center
- Add Dataverse tables as a knowledge source - https://learn.microsoft.com/microsoft-copilot-studio/knowledge-add-dataverse
- Grounding data design for AI workloads on Azure - https://learn.microsoft.com/azure/well-architected/ai/grounding-data-design

Topic 1 - Question 21

Multiple choice | Domain: Design

topic-1/q21

A company plans to deploy a Microsoft Copilot Studio agent to enhance customer support.

The company stores customer data across ServiceNow, Microsoft Dynamics 365 Finance, Dynamics 365 Supply Chain Management, and Excel files in SharePoint Online.

You need to recommend a solution to ensure that the agent can deliver accurate and timely responses.

What should you recommend?

A. Implement a model router for query handling.
B. Create custom prompts.
C. Implement Microsoft Power Platform connectors.
D. Enable incremental indexing in Azure AI Search.

Answer

C

Correct answer: C. Implement Microsoft Power Platform connectors.

Why C is correct

The agent has to deliver accurate and timely responses while reaching customer data spread across four very different systems: ServiceNow, Dynamics 365 Finance, Dynamics 365 Supply Chain Management, and Excel files in SharePoint Online. The Microsoft-recommended way to bring all four into a single Copilot Studio agent is via Power Platform connectors.

Microsoft Learn's "Add Power Platform connectors as knowledge (preview)" lists supported connectors that include exactly the systems in the question (ServiceNow, Dynamics 365, SharePoint, plus many others). Connectors give the agent live, authenticated access to those systems for both data retrieval (knowledge grounding) and actions. They also respect each source system's permissions and security model.

This is the documented Microsoft approach for an agent that must span heterogeneous enterprise systems and deliver current, accurate responses without copying or staging data into a separate store.

Why the other options are wrong

A. Implement a model router for query handling routes between language models based on cost or capability. It does not address how the agent reaches data spread across ServiceNow, Dynamics 365, and SharePoint.

B. Create custom prompts does not connect the agent to data. Prompts shape model behavior; they don't bridge external systems.

D. Enable incremental indexing in Azure AI Search is a data preparation optimization for Azure AI Search indexes. It is helpful when AI Search is the grounding store, but it doesn't connect the agent to

ServiceNow, Dynamics 365 Finance, Dynamics 365 Supply Chain Management, or SharePoint Excel files. The question is about reaching multiple disparate sources, not about indexing speed in one of them.

Microsoft Learn references

- Add Power Platform connectors as knowledge - https://learn.microsoft.com/microsoft-copilot-studio/knowledge-real-time-connectors
- Use Power Platform connectors in Copilot Studio - https://learn.microsoft.com/microsoft-copilot-studio/advanced-connectors
- Microsoft 365 Copilot connectors overview - https://learn.microsoft.com/microsoft-365/copilot/connectors/overview
- Add unstructured data as a knowledge source (SharePoint) - https://learn.microsoft.com/microsoft-copilot-studio/knowledge-add-unstructured-data

Topic 1 - Question 22

Multiple choice | Domain: Plan

topic-1/q22

A manufacturing company wants to deploy an agent that will automate supplier invoice processing.

You are designing a solution to evaluate the financial implications of the deployment. The company is especially concerned about budget overruns.

You need to ensure that the solution considers the total cost of ownership (TCO), the expected savings from using automation, and whether to extend the existing AI capabilities.

What should you include in the design?

A. a break-even analysis only
B. adopting prebuilt agents to reduce the deployment time
C. training a custom model
D. a return on AI investment (ROAI) analysis

Answer

D

Correct answer: D. a return on AI investment (ROAI) analysis

Why D is correct

The design has to evaluate three things together: total cost of ownership (TCO), expected savings from the automation, and whether to extend existing AI capabilities. That is the exact scope of an ROAI analysis. Microsoft Learn's "Business plan for AI agents" and "Plan for AI adoption" articles describe ROAI as the structured comparison of all costs against measurable benefits, including decisions about whether to expand or modify the AI investment based on the data.

The Copilot Studio "Analyze agent return on investment" feature also lets organizations track time and money savings per agent run, feeding directly into an ROAI view. ROAI captures both the TCO side (development, deployment, operating, governance, model usage) and the benefits side (savings, productivity, cycle-time reductions), and it provides the basis for the extension decision.

Why the other options are wrong

A. a break-even analysis only answers when costs are recouped. It is one input to ROAI but doesn't cover the whole scope (TCO + ongoing savings + extension decision). The word "only" makes it incomplete for the requirements stated.

B. adopting prebuilt agents to reduce the deployment time is a tactical choice, not a financial evaluation framework. It doesn't analyze TCO or savings.

C. training a custom model is an implementation choice and increases costs. It doesn't evaluate the financial implications of the deployment.

Microsoft Learn references

- Business plan for AI agents - https://learn.microsoft.com/azure/cloud-adoption-framework/ai-agents/business-strategy-plan
- Plan for AI adoption - https://learn.microsoft.com/azure/cloud-adoption-framework/ai/plan
- Analyze agent return on investment - https://learn.microsoft.com/power-platform/release-plan/2025wave1/microsoft-copilot-studio/analyze-agent-return-investment
- Quantify business value (FinOps) - https://learn.microsoft.com/cloud-computing/finops/framework/quantify/quantify-business-value

Topic 1 - Question 23

Multiple choice | Domain: Plan

topic-1/q23

A company has a Microsoft Power Platform solution that contains the following components:

- Microsoft Dataverse tables
- A Microsoft Power BI workspace named WS1
- A canvas app named App1 that uses Dataverse
- A Power BI semantic model that connects to Dataverse by using DirectQuery

You plan to use generative AI to provide answers to queries based on a subset of corporate data.

You need to ensure that the data is available as a grounding data source for AI systems.

What should you do?

A. Populate a Dataverse table.
B. Share WS1.
C. Endorse the semantic model.
D. Export the semantic model.

Answer

C

Correct answer: C. Endorse the semantic model.

Why C is correct

The objective is to make the data available as a grounding source for AI systems that answer queries against a subset of corporate data. The data already lives in Dataverse, the canvas app already uses it, and a Power BI semantic model already connects to Dataverse via DirectQuery.

Microsoft Learn's "Endorse your content" and "Add a Power BI semantic model as a data source to Fabric data agent" articles explain that semantic models endorsed (promoted or certified) in Power BI are clearly identified, prioritized in searches, and discoverable. Endorsement is the documented way to mark a semantic model as authoritative and ready for reuse, including for AI grounding scenarios such as Power BI Copilot, Fabric Data Agents, and Microsoft 365 Copilot. The semantic model best practices article for data agents calls out that endorsement combined with Prep for AI is what makes a semantic model a reliable grounding source for AI systems.

Endorsing the semantic model surfaces it in the OneLake catalog and in AI tools as a trusted grounding source, which directly satisfies the requirement to "make the data available as a grounding data source for AI systems."

Why the other options are wrong

A. Populate a Dataverse table doesn't add anything new for grounding. The Dataverse data is already there and is already queried by the existing semantic model.

B. Share WS1 controls workspace access but does not endorse or expose the model as a discoverable, trusted AI grounding source.

D. Export the semantic model moves it out of the Power BI service, which removes the integrated AI grounding pathway. Exported model files are not a grounding source for live AI queries.

Microsoft Learn references

- Endorse your content (Power BI promotion and certification) - https://learn.microsoft.com/power-bi/collaborate-share/service-endorse-content
- Promote and certify Power BI content with endorsement - https://learn.microsoft.com/power-bi/collaborate-share/service-endorsement-overview
- Add a Power BI semantic model as a data source to Fabric data agent - https://learn.microsoft.com/fabric/data-science/data-agent-semantic-model
- Semantic model best practices for data agent - https://learn.microsoft.com/fabric/data-science/semantic-model-best-practices

Topic 1 - Question 24

Drag and drop | Domain: Design

topic-1/q24

A company plans to implement an AI business solution for a consumer goods company.

You need to create agents that meet the following requirements:

- Orchestrate the sales order fulfillment and shipping of goods to customers.
- Analyze historical data and trends to replenish stock.

Which type of agent should you use for each requirement? To answer, drag the appropriate agent types to the correct requirements. Each agent type may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

Drag sources

- Autonomous
- Prompt-and-response
- Task

Drop targets

1. Orchestrate the sales order fulfillment and shipping: - drop here
2. Analyze historical data and trends: - drop here

Answer

Orchestrate the sales order fulfillment and shipping = Task; Analyze historical data and trends = Autonomous

Correct answer: - Orchestrate the sales order fulfillment and shipping: **Task** - Analyze historical data and trends to replenish stock: **Autonomous**

Why these answers are correct

Microsoft Learn's Cloud Adoption Framework for AI agents categorizes agents as productivity (retrieval), task (action), and automation (autonomous) agents. The Microsoft Copilot Studio guidance and Dynamics 365 Business Central agent documentation also describe these roles in the same way:

Task agent for sales order fulfillment and shipping. A task agent "performs specific tasks within defined workflows, such as updating records or triggering processes." That maps cleanly to orchestrating sales order fulfillment and shipping, which is a workflow with defined steps and structured handoffs (verify items, create the order, generate shipping documents, notify the customer).

Autonomous agent for analyzing historical data and trends to replenish stock. Microsoft Learn's "Design autonomous agent capabilities" article describes autonomous agents as agents that "perceive events, make decisions, and execute tasks independently by using triggers, instructions, and guardrails," and that "operate continuously in the background, monitoring data, reacting to conditions, and running workflows at scale." Replenishment driven by historical-trend analysis is exactly that pattern: the agent monitors stock levels and historical sales, decides when and how much to reorder, and triggers the replenishment process without waiting for a user prompt.

Why the other option is wrong

Prompt-and-response agents respond to direct user questions with retrieved or generated answers. Neither orchestration of multi-step fulfillment nor autonomous trend monitoring fits a simple prompt-and-response model.

Microsoft Learn references

- AI agent adoption (productivity, task, automation agents) - https://learn.microsoft.com/azure/cloud-adoption-framework/ai-agents/
- Design autonomous agent capabilities - https://learn.microsoft.com/microsoft-copilot-studio/guidance/autonomous-agents
- Sales Order Agent overview (task automation example) - https://learn.microsoft.com/dynamics365/business-central/sales-order-agent
- Commerce agent templates (Inventory Fulfillment Agent) - https://learn.microsoft.com/dynamics365/guidance/agent-templates/commerce-agent-templates

Topic 1 - Question 25

Multiple choice | Domain: Deploy

topic-1/q25

A company has an AI agent that automates the review of customer feedback stored in a cloud database.

You plan to generate monthly reports from the agent's output to provide insights into customer sentiment and guide product development and marketing.

You need to ensure that the data ingested by the agent is clean and suitable for the intended use.

What should you do to prepare the data?

A. Create a workflow in Microsoft Power Automate.
B. Identify and address biased data.
C. Create an agent flow in Microsoft Copilot Studio.
D. Sort the database by customer last name.

Answer

B

Correct answer: B. Identify and address biased data.

Why B is correct

The objective is to ensure the data ingested by the agent is **clean and suitable for the intended use** before generating monthly insights into customer sentiment. Microsoft Learn's responsible AI guidance for Copilot Studio and the Customer Insights sentiment analysis documentation both call out the same risk: biased input data leads to biased outputs and decisions, especially in sentiment analysis (where digital-only feedback or skewed sources misrepresent the customer base).

The "Apply responsible AI principles" article for Copilot Studio recommends using "diverse and representative data" and using "bias detection and mitigation tools" to "detect biases in the AI models" and apply corrective actions. The "Analyze sentiment in customer feedback" Customer Insights documentation explicitly warns about potential bias in feedback data and the need to address it before analysis. Identifying and addressing biased data is the data-preparation step that makes the ingested data suitable for the intended use.

Why the other options are wrong

A. Create a workflow in Microsoft Power Automate is an automation step, not a data quality / bias step. It moves data; it doesn't make it clean or unbiased.

C. Create an agent flow in Microsoft Copilot Studio is also automation orchestration, not data preparation.

D. Sort the database by customer last name has no impact on data quality, suitability, or bias.

Microsoft Learn references

- Apply responsible AI principles (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/responsible-ai
- Analyze sentiment in customer feedback (potential bias) - https://learn.microsoft.com/dynamics365/customer-insights/data/sentiment-analysis
- Understanding Data Curation and Management for AI Projects - https://learn.microsoft.com/ai/playbook/capabilities/data-curation/
- Design training data for AI workloads on Azure - https://learn.microsoft.com/azure/well-architected/ai/training-data-design

Topic 1 - Question 26

Multiple choice | Domain: Plan

topic-1/q26

A company is designing a Microsoft Power Platform solution to reduce the manual steps of a business process by deploying an existing AI model.

You need to calculate the return on AI investment (ROAI) by identifying the metadata and telemetry of the solution.

What should you use?

A. Microsoft Power Platform admin center
B. Success by Design
C. the Business value toolkit
D. Microsoft Cloud Adoption Framework for Azure

Answer

C

Correct answer: C. the Business value toolkit

Why C is correct

The requirement is to calculate ROAI for a Power Platform solution by identifying the **metadata and telemetry** of the solution. That is the exact purpose of the Business value toolkit.

Microsoft Learn's "Choose the best methods and tools to measure business value" article maps use cases to recommended methods and explicitly states: for "Everybody" (large numbers of users, smallest number of solutions), the identification method is "Metadata and telemetry" and the recommended methodology is the **Business value toolkit**. The "Capture and communicate value with the Business value toolkit" article describes the toolkit as part of the Center of Excellence starter kit, designed to capture and communicate the value of Power Platform solutions, including AI-driven solutions, using a structured storytelling process aligned to organizational objectives.

The Business value toolkit is also the documented Power Platform tool for ROI/ROAI measurement when the input is solution metadata and telemetry (which is what the question states).

Why the other options are wrong

A. Microsoft Power Platform admin center provides administrative monitoring and capacity views. It is not the ROAI calculation tool.

B. Success by Design is the prescriptive implementation methodology for Dynamics 365. It is not the ROAI / business-value measurement tool for Power Platform solutions.

D. Microsoft Cloud Adoption Framework for Azure is the framework for adopting Azure infrastructure (used for the broader migration in the Fabrikam case study). It is not the Power Platform business-value measurement tool.

Microsoft Learn references

- Choose the best methods and tools to measure business value - https://learn.microsoft.com/power-platform/guidance/adoption/business-value-methods
- Capture and communicate value with the Business value toolkit - https://learn.microsoft.com/power-platform/guidance/coe/business-value-toolkit
- Measure value and realize return on investment - https://learn.microsoft.com/power-platform/guidance/adoption/common-vision/realize-value
- Measure and communicate the business value of Power Platform solutions - https://learn.microsoft.com/power-platform/guidance/adoption/business-value

Topic 1 - Question 27

Hotspot | Domain: Design

topic-1/q27

A company has a Microsoft Power Platform environment.

You need to build two agents named Agent1 and Agent2. The solution must meet the following requirements:

- Agent1 must be extendable by using the Semantic Kernel and must connect to multiple business apps and APIs.
- Agent2 must connect directly to data stored in Microsoft Dataverse and must be embeddable in a Microsoft Power Apps canvas app.

What should you use to build each agent? To answer, select the appropriate options in the answer area.

Answer area

- Agent1:

[ ] Microsoft Foundry
[ ] Azure Logic Apps
[ ] Copilot in Power Apps
[ ] Microsoft Copilot Studio

- **Agent2:**
  [ ] Microsoft Foundry
  [ ] Azure Logic Apps
  [ ] Copilot in Power Apps
  [ ] Microsoft Copilot Studio

Answer

Agent1 = Microsoft Foundry; Agent2 = Copilot in Power Apps

Correct answer: - Agent1: **Microsoft Foundry** - Agent2: **Copilot in Power Apps**

Why these answers are correct

Agent1 must be extendable using the **Semantic Kernel** and connect to multiple business apps and APIs. Microsoft Foundry (Azure AI Foundry / Microsoft Foundry Agent Service) is the pro-code agent platform that integrates natively with Semantic Kernel and the Microsoft Agent Framework. Microsoft Learn's Semantic Kernel documentation, the "Connect to a Microsoft Foundry agent" article, and the Foundry Agent transparency note describe Foundry as the place to build agents that can be orchestrated via Semantic Kernel and that integrate with broad sets of APIs, tools, and connectors. That is the natural fit for Agent1.

Agent2 must connect directly to Dataverse data and be **embeddable in a Microsoft Power Apps canvas app**. Copilot in Power Apps (Agent builder in Power Apps) is the documented option for that. Microsoft Learn's "Agent builder in Power Apps" and "Copilot for Power Apps makers and users" articles describe the Power Apps Agent builder as a fast way to build agents directly within Power Apps Studio, using existing Dataverse data, logic, and actions, embedded in the canvas app's user experience. That maps directly to Agent2's requirements.

Why the other options are wrong (for the slot they don't fit)

Azure Logic Apps is a workflow orchestration service. It is not an agent-building platform with Semantic Kernel extensibility, and it is not embeddable as an agent inside a Power Apps canvas app.

Microsoft Copilot Studio is a low-code agent platform and is a strong general-purpose option for both agents, but for Agent1's specific Semantic Kernel extensibility requirement, Foundry is the documented Microsoft platform. For Agent2's embed-in-canvas-app requirement, Copilot in Power Apps (Agent builder in Power Apps) is the documented in-product experience for building and embedding the agent directly in the canvas app, which is closer to the question's wording than a Copilot Studio agent surfaced separately.

Microsoft Learn references

- Exploring the Semantic Kernel CopilotStudioAgent / Foundry agents - https://learn.microsoft.com/semantic-kernel/frameworks/agent/agent-types/copilot-studio-agent
- Connect to a Microsoft Foundry agent - https://learn.microsoft.com/microsoft-copilot-studio/add-agent-foundry-agent
- Agent builder in Power Apps - https://learn.microsoft.com/power-platform/release-plan/2025wave1/power-apps/ga-agent-builder-power-apps
- Copilot for Power Apps makers and users - https://learn.microsoft.com/power-platform/release-plan/2025wave2/power-apps/copilot-power-apps-makers-users
- Bring intelligence into model driven apps and custom components using Agent Xrm and PCF APIs - https://learn.microsoft.com/power-platform/release-plan/2025wave1/power-apps/bring-intelligence-into-model-driven-apps-custom-components-using-agent-xrm-pcf-apis

Topic 1 - Question 28

Multiple choice | Domain: Plan

topic-1/q28

A company has an Azure environment that supports multiple business units.

The company plans to implement an AI solution that will perform sentiment analysis on customer product reviews.

You need to evaluate the potential cost of the solution to support return on AI investment (ROAI) analysis.

What should you use?

A. Cost Management + Billing
B. Microsoft Fabric SKU Estimator
C. Total Cost of Ownership (TCO) Calculator
D. Azure Reservations

Answer

C

Correct answer: C. Total Cost of Ownership (TCO) Calculator

Why C is correct

The requirement is to **evaluate the potential cost** of an Azure AI solution to support an ROAI analysis. That is a forward-looking estimate of cost before deployment, and it is exactly the use case for the Microsoft TCO Calculator.

Microsoft Learn's "Make an inventory and collect data" article describes the TCO Calculator as the tool to "compare the costs of running your workloads on-premises versus on Azure" and to gain "insights into potential cost savings and efficiencies." The "Describe cost management in Azure" learning module teaches the TCO Calculator alongside the Pricing Calculator as the planning tools you use **before** deploying workloads, which is the input ROAI analysis needs.

For an organization evaluating the cost of a new sentiment-analysis solution before deployment, the TCO Calculator (or Pricing Calculator) is the documented planning tool, and TCO is the option offered.

Why the other options are wrong

A. Cost Management + Billing analyzes **actual** Azure spending after deployment. It is not a planning/estimation tool for a solution that hasn't been built yet.

B. Microsoft Fabric SKU Estimator estimates Microsoft Fabric capacity costs specifically. It is too narrow for evaluating a broader AI sentiment-analysis solution that may use Azure AI services beyond Fabric.

D. Azure Reservations is a purchasing option that gives discounts for committed usage. It is not an estimation tool for evaluating potential cost.

Microsoft Learn references

- Make an inventory and collect data (TCO Calculator) - https://learn.microsoft.com/azure/app-modernization-guidance/assess/make-an-inventory-and-collect-data
- Compare the Pricing and Total Cost of Ownership calculators (training) - https://learn.microsoft.com/training/modules/describe-cost-management-azure/3-compare-pricing-total-cost-of-ownership-calculators

- Plan to manage costs for Azure OpenAI - https://learn.microsoft.com/azure/ai-foundry/openai/how-to/manage-costs
- Estimate costs with the Azure pricing calculator - https://learn.microsoft.com/azure/cost-management-billing/costs/pricing-calculator

Topic 1 - Question 29

Hotspot | Domain: Design

topic-1/q29

You need to recommend a Microsoft Power Platform solution for customer support. The solution must include AI capabilities in Microsoft Power Automate and must meet the following requirements:

- Use a centralized workspace for AI models.
- Generate short overviews from large amounts of unstructured text, such as case notes or transcripts, without requiring additional training or coding.

What should you include in the recommendation for each requirement? To answer, select the appropriate options in the answer area.

Answer area

- Use a centralized workspace:

[ ] An Microsoft Foundry hub
[ ] Azure OpenAI Foundry
[ ] Microsoft Copilot Studio
[ ] Microsoft Dataverse

- Generate short overviews:

[ ] An AI Builder prebuilt model
[ ] An AI Builder prebuilt prompt
[ ] Azure OpenAI
[ ] GitHub Copilot
[ ] Microsoft Copilot in Power Automate

Answer

Centralized workspace = A Microsoft Foundry hub; Generate short overviews = An AI Builder prebuilt prompt

Correct answer: - Use a centralized workspace: **A Microsoft Foundry hub** - Generate short overviews: **An AI Builder prebuilt prompt**

Why these answers are correct

For the centralized AI workspace requirement, **a Microsoft Foundry hub** is the documented Microsoft pattern. Microsoft Foundry hubs are the central workspace for managing AI models, projects, and resources in Azure. They give teams a shared location for model deployment, governance, and collaboration. Microsoft Learn's Foundry control-plane and architecture documentation describe hubs as the centralized organizational container for AI work.

For generating short overviews of unstructured text **without requiring additional training or coding**, **an AI Builder prebuilt prompt** is the right answer. Microsoft Learn's "Get started with prebuilt prompts" article documents the AISummarize prebuilt prompt, which "summarizes the text that you provide" (input: text, output: summarized text). It is available out of the box in Power Automate, Power Apps, and Dataverse low-code plug-ins. No training is required and no code has to be written, which exactly matches the requirement.

Why the other options are wrong (for the slot they don't fit)

Azure OpenAI Foundry is a model service, not a centralized AI workspace abstraction in the way a Foundry hub is.

Microsoft Copilot Studio is the platform for building agents, not the centralized workspace for AI models in the Foundry sense.

Microsoft Dataverse is the data backbone, not an AI model workspace.

An AI Builder prebuilt model refers to AI Builder's prebuilt models (form processing, receipt processing, etc.), which are separate from prebuilt prompts. The prebuilt prompts (AISummarize, AISentiment, AIReply, AIClassify, AIExtract, AITranslate) are the no-training, no-coding text-generation features.

Azure OpenAI requires more setup than the requirement allows ("without requiring additional training or coding") and is not an AI Builder out-of-the-box prompt.

GitHub Copilot is a developer code-completion tool, not a Power Automate text-summarization feature.

Microsoft Copilot in Power Automate assists makers in building flows; it is not the no-code summarization element you place inside a flow.

Microsoft Learn references

- Get started with prebuilt prompts (AISummarize) - https://learn.microsoft.com/ai-builder/prebuilt-prompts
- Microsoft Foundry hubs (architecture overview) - https://learn.microsoft.com/azure/foundry/concepts/foundry-models-overview
- AI Builder in Power Automate overview - https://learn.microsoft.com/ai-builder/use-in-flow-overview
- Prompt builder - https://learn.microsoft.com/power-platform/release-plan/2025wave1/ai-builder/prompt-builder

Topic 1 - Question 30

Hotspot | Domain: Design

topic-1/q30

A company uses Microsoft Dynamics 365 Finance for accounts payable and customer debt recovery.

You are designing an AI finance process that meets the following requirements:

- Provides AI-driven details to help staff identify overdue vendor invoices and outstanding balances
- Helps staff reduce how long it takes to review overdue invoices and payment history

You need to recommend which Microsoft Copilot features to include in the design.

What should you recommend for each requirement? To answer, select the appropriate options in the answer area.

Answer area

- Help identify vendor overdue invoices and outstanding balances:
 [ ] Agent management
 [ ] AI Summaries with Copilot
 [ ] The Account Reconciliation agent
 [ ] The Supplier Communication agent

- Reduce how long it takes to review overdue invoices and payment history:
 [ ] Analyze demand plans with Copilot
 [ ] Collections coordinator summary
 [ ] The Account Reconciliation agent
 [ ] The Supplier Communication agent

Answer

Identify overdue invoices and outstanding balances = AI Summaries with Copilot; Reduce review time = Collections coordinator summary

Correct answer: - Help identify vendor overdue invoices and outstanding balances: **AI Summaries with Copilot** - Reduce how long it takes to review overdue invoices and payment history: **Collections coordinator summary**

Why these answers are correct

The Dynamics 365 Finance Copilot feature set in Microsoft Learn maps each capability to a specific job:

AI Summaries with Copilot in Dynamics 365 Finance generate AI-driven detail summaries (vendor and customer page summaries, workflow history summary, customer page summary). For accounts payable identification of overdue vendor invoices and outstanding balances, AI Summaries with Copilot give finance staff the AI-driven detail view across the relevant pages. Microsoft Learn's "Agents, Copilot, and AI capabilities in Dynamics 365 apps" article and the TechTalk on Copilot capabilities in Dynamics 365 Finance describe these context-aware summarization features (vendor summary, customer summary, workflow history summary) as the way to surface AI-driven details to staff.

Collections coordinator summary is purpose-built to "reduce the time reviewing collections details for customers." Microsoft Learn's "Collections coordinator summary" article states the feature generates an AI summary of overdue invoices, payment history, and outstanding balances on the Collections coordinator workspace, and that it is "designed to reduce the time that you must spend reviewing collections details for your customers." That is a direct match to the second requirement: reducing how long it takes to review overdue invoices and payment history.

Why the other options are wrong (for the slot they don't fit)

Agent management is the preview framework for managing autonomous AI agents in finance and operations apps; it is not a feature that identifies vendor overdue invoices.

The Account Reconciliation agent automates subledger-to-general-ledger reconciliation, not overdue-invoice identification or collections review reduction.

The Supplier Communication agent handles supplier communications, not overdue-balance identification or AI-driven collections review.

Analyze demand plans with Copilot is a Supply Chain Management capability for demand planning, not an accounts payable / accounts receivable review feature.

Microsoft Learn references

- Agents, Copilot, and AI capabilities in Dynamics 365 apps (Dynamics 365 Finance) - https://learn.microsoft.com/dynamics365/copilot/ai-get-started#dynamics-365-finance
- Collections coordinator summary - https://learn.microsoft.com/dynamics365/finance/accounts-receivable/collectionscoordinatorsummary
- Collections coordinator summary: FAQ - https://learn.microsoft.com/dynamics365/finance/accounts-receivable/collections-coordinator-summary-faq
- TechTalk for Copilot capabilities in Dynamics 365 Finance and Supply Chain Management - https://learn.microsoft.com/dynamics365/guidance/techtalks/dynamics-365-finance-supply-chain-management-copilot-capabilities

Topic 1 - Question 31

Multi-select | Domain: Plan | Case study: Contoso

topic-1/q31

Which two components in the custom AI agent design should the CFO evaluate in the quarterly agent analysis? Each correct answer presents part of the solution.

A. the GPT models used for the agent
B. the average characters in a chat message
C. the agent orchestration method
D. the average session time per agent

Select all that apply.

Answer

A, C

Correct answers: A. the GPT models used for the agent, and C. the agent orchestration method

Why these answers are correct

The Contoso case study states: "The CFO will analyze all the AI solutions quarterly to compare the estimated ROI against actual measured efficiencies and adoption. The CFO will use the Copilot Studio agent usage estimator to perform this analysis."

Microsoft Learn's Copilot Studio billing/licensing pages describe the inputs to the Copilot Studio agent usage estimator directly: "Forecast your agent's Copilot Credits volume using the Microsoft Copilot Studio agent usage estimator. Create estimates and potential consumption impacts by selecting from agent type, traffic, **orchestration**, knowledge, and tools."

The estimator (and Copilot Credits cost in general) is sensitive to two design choices that the CFO would evaluate:

A. The GPT models used for the agent. Copilot Studio's models page and the Application Card describe how different language models (GPT-4o, GPT-4.1, GPT-5, third-party models) drive different consumption and capability profiles. Model choice directly affects Copilot Credit consumption per interaction, which is a primary cost driver in the estimator and therefore central to the CFO's quarterly ROI analysis.

C. The agent orchestration method. Orchestration (generative vs. classic) is an explicit input to the agent usage estimator. Generative orchestration consumes more Copilot Credits per interaction because the model is used to plan and select tools, topics, and knowledge dynamically. Classic orchestration is more deterministic and typically consumes fewer credits. The CFO must consider this in the quarterly analysis.

Why the other options are wrong

B. The average characters in a chat message is not an input to the agent usage estimator and is not a primary Copilot Credit cost driver. Copilot Credits are billed by activity type and orchestration model usage, not character count.

D. The average session time per agent is a usage observation metric, not an input to the agent usage estimator. The CFO's tool of choice is the estimator, which forecasts based on agent type, traffic, orchestration, knowledge, and tools. Session time is closer to engagement analytics than ROAI cost forecasting.

Microsoft Learn references

- Copilot Studio licensing (agent usage estimator) - https://learn.microsoft.com/microsoft-copilot-studio/billing-licensing

- Microsoft Copilot Studio agent usage estimator - https://learn.microsoft.com/microsoft-copilot-studio/agent-usage-estimator
- Billing rates and management (Copilot Credits) - https://learn.microsoft.com/microsoft-copilot-studio/requirements-messages-management
- Select a primary AI model for your agent - https://learn.microsoft.com/microsoft-copilot-studio/authoring-select-agent-model
- Orchestrate agent behavior with generative AI - https://learn.microsoft.com/microsoft-copilot-studio/advanced-generative-actions

Topic 1 - Question 32

Multiple choice | Domain: Design | Case study: Contoso

topic-1/q32

What should you configure for the custom AI agent?

A. AI-assisted evaluators
B. classic orchestration
C. generative orchestration
D. Azure OpenAI reasoning models

Answer

C

Correct answer: C. generative orchestration

Why C is correct

The Contoso case study calls out two requirements that map directly to generative orchestration:

1. "The custom AI agent must be designed to eventually connect to other agents that can be selected based on their description."
2. "The topics used in the custom AI agent will be selected based NOT on a trigger phrase, but on a description of the purpose of the query, to make the interactions more conversational."

Microsoft Learn's "Orchestrate agent behavior with generative AI" article maps these requirements to generative orchestration explicitly:

- "With generative orchestration, an agent answers user queries or responds to event triggers by selecting the most appropriate combination of topics, tools, and knowledge. Each topic has a description that informs the agent of its purpose."
- "The agent selects topics based on the description of their purpose."
- "The agent selects child and connected agents based on their description."

By contrast, classic orchestration uses trigger phrases to match topics, which is exactly what the case study says Contoso doesn't want. Classic orchestration also doesn't support "Child and connected agents" selection by description.

Generative orchestration is the only option that supports both requirements (description-based topic selection and description-based agent-to-agent routing).

Why the other options are wrong

A. AI-assisted evaluators are agent quality evaluators used in test sets to grade response quality. They don't determine how the agent selects topics or connects to other agents.

B. Classic orchestration uses trigger phrases for topic selection, which directly contradicts Contoso's requirement that topics be selected by description. Classic orchestration also does not support agent-to-agent connection selection by description.

D. Azure OpenAI reasoning models are a class of language models. The orchestration mode is a separate setting from the model choice. The case study's requirements describe orchestration behavior (description-based selection, multi-agent connections), not a model class.

Microsoft Learn references

- Orchestrate agent behavior with generative AI - https://learn.microsoft.com/microsoft-copilot-studio/advanced-generative-actions
- FAQ for using generative orchestration - https://learn.microsoft.com/microsoft-copilot-studio/faqs-generative-orchestration
- Set topic triggers (The agent chooses) - https://learn.microsoft.com/microsoft-copilot-studio/authoring-triggers
- Add other agents overview (connect to other agents) - https://learn.microsoft.com/microsoft-copilot-studio/authoring-add-other-agents
- Apply generative orchestration capabilities - https://learn.microsoft.com/microsoft-copilot-studio/guidance/generative-orchestration

Topic 2

29 questions

Topic 2 - Question 1

Multiple choice | Domain: Design | Case study: Fabrikam

topic-2/q01

Which template should you use for the AI agent to meet the requirements for the sales executives?

A. IT Helpdesk in Microsoft Copilot Studio
B. AI agents in Microsoft Foundry
C. Voice in Microsoft Copilot Studio
D. AI chat in Microsoft Foundry

Answer

C

Correct answer: C. Voice in Microsoft Copilot Studio

Why C is correct

The case study tells you the sales executives "use handsfree headsets to interact with an AI agent when they have questions about internal policies or customer data." The Voice agent template in Microsoft Copilot Studio is built specifically for that scenario.

From Microsoft Learn:

"By using the Voice agent template, you can build a voice-enabled agent that provides an effective self-service, hands-free solution from a phone. Customers can interact with the agent by using natural language and choosing options from a touch-tone menu."

The Voice template enables the Telephony channel by default, includes Speech and DTMF modality, and ships with voice-related system topics (Silence detection, Speech unrecognized, Unknown dial pad press). It is the Microsoft-recommended template when the conversational user experience must be hands-free and speech-driven, which is exactly what Fabrikam's stated requirement specifies.

Why the other options are wrong

A. IT Helpdesk in Microsoft Copilot Studio is a custom agent template aimed at IT support scenarios (password resets, ticket creation, common helpdesk topics). It has no inherent voice or hands-free capability and does not match Fabrikam's hands-free headset requirement.

B. AI agents in Microsoft Foundry is a code-first or workflow-first agent platform for prompt agents, workflow agents, and hosted agents. It is more developer-oriented and does not provide a low-code, hands-free voice template aligned to Fabrikam's "single AI agent that has Dataverse as its core component" and conversational sales-executive scenario.

D. AI chat in Microsoft Foundry is a text-based chat capability in Foundry. It does not deliver the speech-driven, telephony-style interaction the case study calls for.

Microsoft Learn references

- Voice agent template (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/voice-build-from-template
- Configure basic voice agents (SSML, DTMF, transfer, termination) - https://learn.microsoft.com/microsoft-copilot-studio/voice-configuration
- Create a custom agent from an agent template - https://learn.microsoft.com/microsoft-copilot-studio/template-fundamentals
- Microsoft Foundry Agent Service overview (for contrast) - https://learn.microsoft.com/azure/foundry/agents/overview

Topic 2 - Question 2

Multiple choice | Domain: Design | Case study: Fabrikam

topic-2/q02

Which tool should you use for the prospect communication requirements in Dynamics 365 Sales?

A. Azure AI Search
B. Copilot email assist
C. the Voice template Microsoft Copilot Studio
D. Deep Research in Microsoft Foundry Agent Service

Answer

B

Correct answer: B. Copilot email assist

Why B is correct

The Fabrikam case study requires that "the sales team must use Dynamics 365 Sales to correspond with prospects more quickly and efficiently than currently" and that sales executives "send follow-up communications to prospects." Copilot email assist in Dynamics 365 Sales is the Microsoft-recommended capability for that requirement.

From Microsoft Learn:

"The Copilot for email assist feature helps sellers to generate content for emails that is specific to customer's needs with clarity, conciseness, and compelling content."

"Copilot helps you spend less time composing emails. With this feature, you can: select a predefined category or enter your text, and Copilot suggests content; adjust the tone and length of the email; customize the suggested content before you send it."

Copilot email assist is supported on lead and opportunity records through Dynamics 365 email and is enabled in the model-driven app under Settings, Features, "Contextual email drafting with AI." It directly addresses faster, more efficient prospect correspondence inside Dynamics 365 Sales without requiring a separate Microsoft 365 Copilot license.

Why the other options are wrong

A. Azure AI Search is a vector and keyword retrieval service used to ground generative AI responses in indexed content. It does not draft prospect emails inside Dynamics 365 Sales.

C. The Voice template in Microsoft Copilot Studio builds hands-free, telephony-based voice agents. The case study uses voice for the sales executives' question-and-answer agent, not for email correspondence with prospects.

D. Deep Research in Microsoft Foundry Agent Service performs deep, multi-step research over public and grounded sources. It is not the tool for drafting outbound prospect emails inside Dynamics 365 Sales.

Microsoft Learn references

- Enable Copilot email assist (Dynamics 365 Sales) - https://learn.microsoft.com/dynamics365/sales/enable-copilot-email-assist
- Elevate your sales pitch using Copilot email assistance - https://learn.microsoft.com/dynamics365/release-plan/2024wave2/sales/dynamics365-sales/elevate-sales-pitch-using-copilot-email-assistance
- Copilot in Dynamics 365 Sales overview - https://learn.microsoft.com/dynamics365/sales/copilot-overview
- FAQ about Copilot for emails - https://learn.microsoft.com/dynamics365/sales/faqs-sales-copilot-for-email

Topic 2 - Question 3

Hotspot | Domain: Design | Case study: Fabrikam

topic-2/q03

HOTSPOT -

Which components should you use to meet the sales cycle enablement requirements? To answer, select the appropriate options in the answer area.

Answer area

- For AI agent creation:

[ ] Microsoft Foundry
[ ] Dynamics 365 Sales
[ ] Microsoft Copilot Studio
[ ] the Power Platform admin center

- For unexpected AI agent actions:

[ ] a custom connector
[ ] an event trigger
[ ] a Fallback topic
[ ] a REST API

Answer

AI agent creation = Microsoft Copilot Studio; Unexpected AI agent actions = a Fallback topic

Correct answers: - For AI agent creation: **Microsoft Copilot Studio** - For unexpected AI agent actions: **a Fallback topic**

Why these answers are correct

AI agent creation: Microsoft Copilot Studio

The case study calls for "low-code development to create a single AI agent that has Dataverse as its core component" and a hands-free voice experience for sales executives. Microsoft Copilot Studio is Microsoft's low-code authoring tool for building agents, and it includes the Voice template, Dataverse knowledge integration, and connectors required by Fabrikam.

From Microsoft Learn:

"Copilot Studio is a graphical, low-code tool for building agents and agent flows."

Microsoft Foundry is code-first or developer-oriented. Dynamics 365 Sales doesn't author agents directly. The Power Platform admin center governs environments and capacity but does not author agents.

Unexpected AI agent actions: a Fallback topic

The case study states: "Unexpected AI agent actions must end in an escalation to a live representative. For example, a sales executive must be rerouted to a representative if the agent cannot answer a question after two failed attempts." That is exactly the documented behavior of the system Fallback topic in Copilot Studio.

From Microsoft Learn:

"If the agent can't determine the user's intent, it prompts the user again. After two prompts, the custom or Copilot agent escalates to a live representative through a system topic called Escalate."

"Customize the fallback topic and behavior in the default system Fallback topic. A fallback topic triggers On Unknown Intent to capture the unrecognized input."

The Fallback topic is the built-in, configurable mechanism that handles unrecognized intent, prompts the user up to two times, and then redirects to the Escalate topic for live handoff. A custom connector, an event trigger, or a REST API are not the right primitives for this requirement.

Microsoft Learn references

- Configure the system fallback topic - https://learn.microsoft.com/microsoft-copilot-studio/authoring-system-fallback-topic
- Topics in Copilot Studio (default system topics, including Fallback and Escalate) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/topics-overview
- Copilot Studio overview - https://learn.microsoft.com/microsoft-copilot-studio/fundamentals-what-is-copilot-studio
- Hand off to a live agent (Escalate topic) - https://learn.microsoft.com/microsoft-copilot-studio/advanced-hand-off

Topic 2 - Question 4

Multiple choice | Domain: Design

topic-2/q04

You need to design a Microsoft 365 Copilot solution to optimize employee productivity. The solution must meet the following requirements:

Ensure that the employees can query content stored in a subset of Microsoft SharePoint Online sites and in Teams by using natural language-based prompt actions.

Ensure that employees receive contextually relevant responses in Microsoft 365 Copilot.

What should you include in the design?

A. Build a Microsoft Power Automate desktop flow to read the SharePoint content and post the responses to Teams.
B. Modify SharePoint settings.
C. Create a custom REST API that crawls the SharePoint content.
D. Configure Microsoft Graph access.

Answer

D

Correct answer: D. Configure Microsoft Graph access.

Why D is correct

Microsoft 365 Copilot uses Microsoft Graph to bring tenant data (SharePoint sites, OneDrive, Teams chats and meetings, email, calendar) into prompts and responses. To let employees query a subset of SharePoint Online sites and Teams content with natural language and get contextually relevant responses, the design depends on Microsoft Graph access.

From Microsoft Learn:

"Microsoft Graph includes information on users, their activities, and the organization data they can access. The Microsoft Graph API brings a personalized context into the prompt, like information from a user's emails, chats, documents, and meetings."

"Microsoft 365 Copilot enhances search relevance and accuracy by using advanced lexical and semantic understanding of Microsoft Graph data, resulting in more contextually precise information retrieval. Copilot preserves security, compliance, and privacy, ensuring organizational boundaries are respected."

Configuring Microsoft Graph access (including semantic indexing for Microsoft 365 Copilot and the user's permissioned SharePoint and Teams content) is the foundational design step for natural-language queries that return contextually relevant answers in Microsoft 365 Copilot.

Why the other options are wrong

A. Build a Microsoft Power Automate desktop flow to read the SharePoint content and post the responses to Teams. Desktop flows automate UI tasks on a local machine. They don't enable natural-language prompts in Microsoft 365 Copilot or contextual grounding in Graph.

B. Modify SharePoint settings. Adjusting site-level SharePoint settings alone does not deliver Microsoft 365 Copilot natural-language querying. Copilot's grounding still requires Microsoft Graph access and the appropriate Copilot service plan.

C. Create a custom REST API that crawls the SharePoint content. Building a custom crawler duplicates work that Microsoft Graph and the semantic index already perform, bypasses native security trimming, and isn't the design Microsoft recommends.

Microsoft Learn references

- Microsoft 365 Copilot overview (Graph and semantic index) - https://learn.microsoft.com/microsoft-365/copilot/microsoft-365-copilot-overview
- Microsoft 365 Copilot architecture - https://learn.microsoft.com/microsoft-365/copilot/microsoft-365-copilot-architecture
- Overview of Microsoft Graph - https://learn.microsoft.com/graph/overview
- Microsoft 365 Copilot service plans (SharePoint and Teams plans) - https://learn.microsoft.com/microsoft-365/copilot/microsoft-365-copilot-service-plans

Topic 2 - Question 5

Multiple choice | Domain: Design

topic-2/q05

A company uses Microsoft Dynamics 365 Finance to manage accounts payable.

You are designing an AI invoice processing solution.

You need to recommend the prerequisites to configure a prebuilt copilot for accounts payable.

What should you recommend?

A. From Microsoft Copilot Studio, create an accounts payable agent.
B. Extend Microsoft 365 Copilot for Sales to an accounts payable agent.
C. Build an AI tool in Microsoft Foundry.
D. From the Power Platform admin center, assign the Finance and Operations AI security role to users.

Answer

D

Correct answer: D. From the Power Platform admin center, assign the Finance and Operations AI security role to users.

Why D is correct

The question asks about prerequisites to configure a prebuilt Copilot for accounts payable in Dynamics 365 Finance. Microsoft Learn documents this exact prerequisite.

From Microsoft Learn:

"Users who should have access to the functionality must be assigned the Finance and Operations AI security role in Dataverse."

"In the detail view of the environment, in the Access field, select Users or Teams. Select the users or teams that should have access, and assign the Finance and Operations AI security role."

This assignment is performed in the Dataverse environment that's associated with finance and operations apps, which is reached through the Power Platform admin center. Without the role, users cannot see Copilot summaries or get answers from the prebuilt Copilot in Dynamics 365 Finance.

Why the other options are wrong

A. From Microsoft Copilot Studio, create an accounts payable agent. The question is about a prebuilt Copilot in Dynamics 365 Finance, not building a custom agent in Copilot Studio. Creating one bypasses the prebuilt experience and is not the prerequisite to configure it.

B. Extend Microsoft 365 Copilot for Sales to an accounts payable agent. Microsoft 365 Copilot for Sales is a sales-focused Outlook and Teams add-on tied to CRM. It is not the foundation for an accounts payable Copilot in Dynamics 365 Finance.

C. Build an AI tool in Microsoft Foundry. Foundry tools are for building custom AI agents and prompts. The question is about prerequisites to enable the prebuilt Copilot, which is a Dataverse role-assignment task, not a Foundry build.

Microsoft Learn references

- Enable Copilot in Dynamics 365 Finance (Finance and Operations AI security role) - https://learn.microsoft.com/dynamics365/finance/accounts-receivable/enable-copilot-in-finance
- Overview of Copilot capabilities in finance and operations apps - https://learn.microsoft.com/dynamics365/fin-ops-core/fin-ops/copilot/copilot-for-finance-operations
- Power Platform admin center (finance and operations apps overview) - https://learn.microsoft.com/power-platform/admin/unified-experience/finance-operations-apps-overview

Topic 2 - Question 6

Multi-select | Domain: Design

topic-2/q06

A company plans to deploy a Microsoft Dynamics 365 Contact Center agent.

You need to ensure that the agent can transfer the conversation to a live customer service representative.

Which two components should you include in the solution? Each correct answer presents part of the solution.

A. Microsoft Foundry
B. Microsoft Copilot Studio
C. Microsoft 365 Agents Toolkit
D. an Azure AI Bot Service skill
E. Customer engagement hub

Select all that apply.

Answer

B, E

Correct answers: B. Microsoft Copilot Studio and E. Customer engagement hub

Why B and E are correct

Microsoft's documented pattern for handing a Dynamics 365 Contact Center conversation off to a live customer service representative requires two pieces:

1. **The agent built in Microsoft Copilot Studio**, with the Escalate system topic (or a Transfer conversation node) configured. From Microsoft Learn:

"When your customers need to speak with a representative, your agent can seamlessly hand off the conversation. When your agent hands off a conversation, it can share the full history of the conversation, and all relevant variables. Make sure you have an escalation article configured in your agent to hand off the conversation to a representative."

1. **A customer engagement hub** (such as Omnichannel for Customer Service in Dynamics 365 Contact Center) that receives the handoff and routes the conversation to a live representative. From Microsoft Learn:

"An engagement hub that is being used by live agents, such as Omnichannel for Customer Service... you need to configure the connection, as described in Configure handoff to Omnichannel for Customer Service."

The agent emits the handoff event from Copilot Studio; the engagement hub catches it and connects the customer to a live representative.

Why the other options are wrong

A. Microsoft Foundry is a developer platform for building, deploying, and scaling AI agents (prompt agents, workflow agents, hosted agents). It is not the Copilot Studio Contact Center handoff mechanism.

C. Microsoft 365 Agents Toolkit is a developer toolkit (formerly Teams Toolkit) used to build and ship custom agents for Microsoft 365 Copilot and Teams. It is not used to enable live-representative handoff in Dynamics 365 Contact Center.

D. Azure AI Bot Service skill is a Bot Framework skill pattern. While Copilot Studio does support skills, the standard Microsoft-recommended Contact Center handoff pattern is Copilot Studio plus an engagement hub, not an Azure AI Bot Service skill.

Microsoft Learn references

- Hand off to a live agent (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/advanced-hand-off
- Configure handoff to Omnichannel for Customer Service - https://learn.microsoft.com/microsoft-copilot-studio/configuration-hand-off-omnichannel
- Integrate a Copilot agent (Dynamics 365 Contact Center) - https://learn.microsoft.com/dynamics365/customer-service/administer/configure-bot-virtual-agent
- Channels and choosing an approach to enable handoff to a live agent - https://learn.microsoft.com/microsoft-copilot-studio/guidance/channels

Topic 2 - Question 7

Hotspot | Domain: Design

topic-2/q07

HOTSPOT -

A company uses Microsoft Dynamics 365 Supply Chain Management.

You are designing an AI supply chain process that meets the following requirements:

Provides managers with AI-driven insights that surface key information from customer orders

Helps planners use AI to anticipate future product needs more accurately

You need to recommend which Microsoft Copilot features to include in the design.

What should you recommend for each requirement? To answer, select the appropriate options in the answer area.

Answer area

- Provide AI-driven insights from customer orders:
- [ ] AI Summaries with Copilot
- [ ] Generative insights for Demand planning
- [ ] The Customer credit and collections workspace
- [ ] Workload insights with Copilot

- Anticipate future product needs:
- [ ] Generative insights for Demand planning
- [ ] Microsoft Power BI
- [ ] Product information management
- [ ] The Supplier Communications Agent

Answer

AI-driven insights from customer orders = AI Summaries with Copilot; Anticipate future product needs = Generative insights for Demand planning

Correct answers: - Provide AI-driven insights from customer orders: **AI Summaries with Copilot** - Anticipate future product needs: **Generative insights for Demand planning**

Why these answers are correct

AI summaries with Copilot for customer order insights

From Microsoft Learn:

"AI summaries with Copilot... provide a quick overview of the most important information that's related to the page, personalized for the current user. Summaries can include information such as the number of lines on a purchase order, the number of items in a warehouse, or the number of overdue invoices for a vendor."

That description maps directly to the requirement to "surface key information from customer orders" for managers in Dynamics 365 Supply Chain Management. AI Summaries with Copilot is the Microsoft-recommended feature for that scenario.

The Customer credit and collections workspace targets credit and AR workflows, not generalized order insights. Workload insights with Copilot are specific to the Warehouse Management mobile app.

Generative insights for Demand planning to anticipate future product needs

Demand planning in Dynamics 365 Supply Chain Management uses Copilot and generative insights to forecast demand based on multiple signals and patterns, with cell-level explainability and best-fit forecast models.

From Microsoft Learn:

"Demand Planning: Copilot, generative insights, and cell-level explainability features improve forecasting accuracy."

"Forecast with signals: Forecast based on multiple input signals by using calculations based on the Prophet forecast algorithm, XGBoost forecast algorithm, or best-fit forecast model."

That capability is the documented way to help planners "anticipate future product needs more accurately." Power BI is generic analytics. Product information management is master data, not forecasting. The Supplier Communications Agent automates procure-to-pay communications with vendors.

Microsoft Learn references

- Agents, Copilot, and AI capabilities in Dynamics 365 apps (Supply Chain Management) - https://learn.microsoft.com/dynamics365/copilot/ai-get-started#dynamics-365-supply-chain-management
- Analyze demand plans with Copilot - https://learn.microsoft.com/dynamics365/supply-chain/demand-planning/demand-planning-copilot
- Enhance your demand forecasting and planning (generative insights) - https://learn.microsoft.com/dynamics365/release-plan/2025wave1/finance-supply-chain/dynamics365-supply-chain-management/enhance-demand-forecasting-planning
- AI summaries with Copilot - https://learn.microsoft.com/dynamics365/supply-chain/get-started/copilot-summaries-overview

Topic 2 - Question 8

Hotspot | Domain: Design

topic-2/q08

HOTSPOT -

A company has a Microsoft 365 E5 subscription and uses Microsoft Copilot Studio.

The company has a Microsoft SharePoint Online library that contains 10,000 policy PDFs from various departments. The library contains a populated column named Department for each PDF.

You need to design a Copilot Studio agent that will use the SharePoint library as a knowledge source. The solution must meet the following requirements:

Enable the agent to answer user questions about company policies.

Ensure that the agent can identify which departments and policies are connected.

What should you include in the design for each requirement? To answer, select the appropriate options in the answer area.

Answer area

- Enable the agent to answer questions about company policies:

[ ] Build a custom model in Microsoft Foundry.
[ ] From Copilot Studio, add SharePoint as a knowledge source.
[ ] Import the PDFs into Microsoft Dataverse.
[ ] Use AI Builder to process and feed SharePoint content.

- Identify which departments and policies are connected:

[ ] Apply Microsoft Purview sensitivity labels.
[ ] Create a Microsoft Dataverse table for the departments.
[ ] From Copilot Studio, configure the SharePoint tool.
[ ] Upgrade to SharePoint Premium.

Answer

Answer questions about company policies = From Copilot Studio, add SharePoint as a knowledge source; Identify which departments and policies are connected = From Copilot Studio, configure the SharePoint tool

Correct answers: - Enable the agent to answer questions about company policies: **From Copilot Studio, add SharePoint as a knowledge source.** - Identify which departments and policies are connected: **From Copilot Studio, configure the SharePoint tool.**

Why these answers are correct

Add SharePoint as a knowledge source

From Microsoft Learn:

"SharePoint as a knowledge source works by pairing your agent with a SharePoint URL or SharePoint lists."

"When a user asks a question and the agent doesn't have a topic to use for an answer, the agent searches the URL and all subpaths... Generative answers summarize this content into a targeted response."

Adding the SharePoint document library as a knowledge source is the Microsoft-documented way to let a Copilot Studio agent answer policy questions grounded in 10,000 PDFs. Building a custom Foundry model, importing PDFs into Dataverse, or using AI Builder to feed SharePoint content are not the Microsoft-recommended low-code path here.

Configure the SharePoint tool to use the Department column

The agent needs to understand that each policy PDF is associated with a Department column value. The SharePoint tool (along with SharePoint metadata filters and advanced settings on a SharePoint knowledge source) lets the agent reason over SharePoint metadata such as Title, Author, Modified by, Modified on, and other refinable managed properties.

From Microsoft Learn:

"Improve response accuracy with SharePoint metadata filters. Use metadata like filename, owner, and modified date to refine knowledge retrieval and ensure responses come from the most relevant, up-to-date documents."

"Makers can aid the performance of their agent's SharePoint knowledge source by specifying search query parameters... Build your filters to include or exclude information from your SharePoint knowledge source."

Configuring the SharePoint tool exposes that metadata to the agent so it can identify which department a policy belongs to. Purview sensitivity labels classify content for protection, not for connecting policies to departments. A Dataverse table for departments duplicates the existing SharePoint column. SharePoint Premium is not a prerequisite for this scenario.

Microsoft Learn references

- Add SharePoint as a knowledge source (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/knowledge-add-sharepoint
- Filter your SharePoint source (advanced settings, metadata filters) - https://learn.microsoft.com/microsoft-copilot-studio/knowledge-add-sharepoint#filter-your-sharepoint-source
- Add knowledge to an agent - https://learn.microsoft.com/microsoft-copilot-studio/knowledge-add-existing-copilot
- What's new in Copilot Studio (SharePoint metadata filters) - https://learn.microsoft.com/microsoft-copilot-studio/whats-new

Topic 2 - Question 9

Drag and drop | Domain: Design

topic-2/q09

DRAG DROP -

You need to design a Microsoft Copilot Studio agent that meets the following requirements:

Supports interactive speech responses

Optimizes decision-making and the accuracy of responses

What should you include in the design for each requirement? To answer, drag the appropriate options to the correct requirements. Each option may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

Drag sources

- A deep reasoning model
- Azure Language in Foundry Tools
- Azure AI Speech
- Copilot Studio voice features
- Speech Synthesis Markup Language (SSML)

Drop targets

1. Supports interactive speech responses - drop here
2. Optimizes decision-making and response accuracy - drop here

Answer

Interactive speech responses = Copilot Studio voice features; Optimize decision-making and accuracy = A deep reasoning model

Correct answers: - Supports interactive speech responses: **Copilot Studio voice features** - Optimizes decision-making and the accuracy of responses: **A deep reasoning model**

Why these answers are correct

Copilot Studio voice features for interactive speech

Copilot Studio's built-in voice features support speech and DTMF input, text-to-speech output, SSML formatting, barge-in, silence detection, real-time voice agents, and the Voice template. They are the Microsoft-recommended capability for interactive speech responses inside a Copilot Studio agent.

From Microsoft Learn:

"Voice interaction is the most natural way for humans to communicate and remains a critical medium in copilot engagements. With Copilot Studio, businesses can easily create and customize agents that interact naturally with customers through voice by accepting speech and dual-tone multi-frequency (DTMF) input from telephony. It provides natural text-to-speech capabilities so end users can hear the agent's response."

Azure AI Speech and SSML can be used inside Copilot Studio voice agents, but they are component services. The end-to-end agent capability that supports interactive speech responses is the Copilot Studio voice feature set. Azure Language in Foundry Tools is for natural language understanding and entity extraction, not speech.

A deep reasoning model for better decisions and accuracy

From Microsoft Learn:

"Deep reasoning models are advanced large language models designed to solve complex problems. They carefully consider each question, generating a detailed internal chain of thought before providing a response back to the user."

"Models like Azure OpenAI o3 use deep reasoning to enhance agent decision making and return more accurate responses."

Deep reasoning models are explicitly positioned to optimize agent decision-making and response accuracy for complex, multi-step tasks.

Microsoft Learn references

- Speech and IVR overview (Copilot Studio voice features) - https://learn.microsoft.com/power-platform/release-plan/2025wave2/microsoft-copilot-studio/speech-ivr
- Basic voice agents overview - https://learn.microsoft.com/microsoft-copilot-studio/voice-basic-overview
- Use deep reasoning models for complex tasks (preview) - https://learn.microsoft.com/microsoft-copilot-studio/authoring-reasoning-models
- FAQ for deep reasoning - https://learn.microsoft.com/microsoft-copilot-studio/faqs-reasoning

Topic 2 - Question 10

Multiple choice | Domain: Design

topic-2/q10

You are designing a low-code AI business solution by using Microsoft Copilot Studio.

The solution must include an agent that automates tasks by simulating user interactions across third-party apps and websites, such as clicking buttons, entering text, and extracting information from screens.

You need to recommend what to include in the agent.

What should you recommend?

A. Model Context Protocol (MCP)
B. a natural language understanding + (NLU+) model in Copilot Studio
C. Computer Use in Copilot Studio
D. Copilot skills

Answer

C

Correct answer: C. Computer Use in Copilot Studio

Why C is correct

The scenario describes an agent that simulates user interactions on third-party apps and websites: clicking buttons, entering text, and extracting information from screens. Computer Use is the Microsoft-documented Copilot Studio tool built specifically for that pattern.

From Microsoft Learn:

"Computer use is a feature in Copilot Studio that enables your agent to interact with any system with a graphical user interface (GUI), such as websites or desktop applications. It performs actions like pressing buttons, selecting menu options, and typing into fields-just like a human user. You simply describe the task in natural language, and the agent executes it on a configured computer with a virtual mouse and keyboard."

"Computer use runs on Computer-Using Agents (CUA), an AI model that combines computer vision and reasoning to navigate and interact with GUIs. It adapts to interface changes... It's useful for automated data entry, invoice processing, and data extraction."

That maps line-for-line to the question's requirements.

Why the other options are wrong

A. Model Context Protocol (MCP) is a standardized protocol for connecting agents to external tools, data, and services through a server. It does not click buttons or interact with arbitrary GUIs.

B. A natural language understanding + (NLU+) model in Copilot Studio improves intent recognition and entity extraction from user utterances. It does not perform GUI automation.

D. Copilot skills are reusable Bot Framework conversational components. They are not designed to drive a virtual mouse and keyboard across third-party UIs.

Microsoft Learn references

- Automate web and desktop apps with computer use (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/computer-use
- Computer use release plan (Copilot Studio) - https://learn.microsoft.com/power-platform/release-plan/2025wave1/microsoft-copilot-studio/automate-web-desktop-apps-computer-use
- FAQ for the computer use tool - https://learn.microsoft.com/microsoft-copilot-studio/faqs-computer-use

Topic 2 - Question 11

Multiple choice | Domain: Design

topic-2/q11

You need to recommend a solution to integrate a Microsoft Copilot agent with a Microsoft Dynamics 365 Contact Center chat channel.

The agent must respond to customer questions and hand off the conversation to a live customer service representative when the customer requests an escalation.

What should you recommend?

A. Build an agent flow.
B. Configure the Conversation Start topic.
C. Configure a skill.
D. Call a Microsoft Power Automate connector.
E. Configure the Escalate topic.

Answer

E

Correct answer: E. Configure the Escalate topic.

Why E is correct

When integrating a Copilot agent with a Dynamics 365 Contact Center chat channel, the documented Microsoft pattern for handing the conversation off to a live representative on customer request is to configure the Escalate system topic. The Escalate topic is what triggers the handoff to the connected engagement hub (Dynamics 365 Contact Center, Customer Service, or Omnichannel for Customer Service).

From Microsoft Learn:

"The Escalate topic is used to hand off the conversation to an external system, typically to a live service representative (when configured-for example to Omnichannel for Customer Service). When this topic is reached, the session outcome is escalated."

"When you create an agent from Dynamics 365 Customer Service, the Escalate system topic already includes a Transfer conversation node. However, agents created in Copilot Studio aren't configured with this node by default. To add a Transfer conversation node to the Escalate system topic... go to Topics > System tab > Escalate, then add Topic Management > Transfer conversation."

That is the exact requirement: respond to questions and escalate to a live representative on request.

Why the other options are wrong

A. Build an agent flow. Agent flows automate business processes through deterministic steps. They aren't the mechanism that hands the chat off to a live representative.

B. Configure the Conversation Start topic. Conversation Start fires at the beginning of the conversation. It greets the user, but it doesn't initiate live-representative handoff on escalation request.

C. Configure a skill. Bot Framework skills are reusable conversational components. They aren't the supported handoff mechanism for Dynamics 365 Contact Center; the documented mechanism is Escalate plus a Transfer conversation node.

D. Call a Microsoft Power Automate connector. Calling a connector can perform automation, but it isn't how Copilot Studio routes a conversation to a live customer service representative in Dynamics 365 Contact Center.

Microsoft Learn references

- Hand off to a live agent (Escalate system topic, Transfer conversation node) - https://learn.microsoft.com/microsoft-copilot-studio/advanced-hand-off
- Topics in Copilot Studio (Escalate system topic) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/topics-overview
- Configure handoff to Omnichannel for Customer Service - https://learn.microsoft.com/microsoft-copilot-studio/configuration-hand-off-omnichannel
- Integrate a Copilot agent (Dynamics 365 Contact Center) - https://learn.microsoft.com/dynamics365/customer-service/administer/configure-bot-virtual-agent

Topic 2 - Question 12

Multiple choice | Domain: Design

topic-2/q12

A company has a customer order system that creates sales orders manually.

You need to design an AI solution to automate the following tasks as part of the system:

Save the order details to a database.

Update the order status in the database.

Extract the order details from an order file.

Prepare and send a confirmation email to customers.

The solution must minimize development effort and support intelligent automation and solution integration.

What should you include in the design?

A. a workflow in Azure Logic Apps
B. a multi-agent solution that uses the Semantic Kernel SDK
C. a multi-agent solution that uses Microsoft Foundry Agent Service
D. a Microsoft Copilot Studio agent that uses Microsoft Power Automate workflows

Answer

D

Correct answer: D. a Microsoft Copilot Studio agent that uses Microsoft Power Automate workflows

Why D is correct

The scenario calls for a low-development-effort solution that supports intelligent automation, integrates with multiple systems, and handles a connected sequence of tasks: extract order details from files, save and update records in a database, and send confirmation emails. A Copilot Studio agent that uses Power Automate workflows (cloud flows or agent flows) is the Microsoft-recommended low-code pattern for this combination of conversational, intelligent, and deterministic automation.

From Microsoft Learn:

"Agents built in Copilot Studio gain new capabilities through integration with other online services... Power Platform offers a rich ecosystem of built-in connectors that are available to Copilot Studio."

"Agent flows execute a series of actions in a predefined sequence. They use the low-code actions found in Power Platform connectors. Agents can pass values as input to an agent flow and receive their outputs."

"Streamline document processing with AI Builder... Power Automate orchestrates workflows for document processing."

This pattern is also reflected in Microsoft's reference architectures for vendor invoice processing and document processing, where Power Automate cloud flows orchestrate the work, AI Builder or AI prompts handle extraction, and Dataverse or back-end systems store the records.

Why the other options are wrong

A. A workflow in Azure Logic Apps can orchestrate the work, but it is a pro-developer service and does not minimize development effort or provide a built-in conversational front door for review and exception handling.

B. A multi-agent solution that uses the Semantic Kernel SDK is a code-first multi-agent framework. It is heavier than the requirement and increases development effort.

C. A multi-agent solution that uses Microsoft Foundry Agent Service is a code- and configuration-driven platform for hosted, prompt, or workflow agents. It does more than required and is not the low-code, low-effort solution Microsoft recommends for this scenario.

Microsoft Learn references

- Plan and design integration strategies (Copilot Studio integrations) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/integrations
- Use agent flows with your agent - https://learn.microsoft.com/microsoft-copilot-studio/advanced-flow
- Streamline document processing with AI Builder - https://learn.microsoft.com/power-platform/architecture/reference-architectures/ai-document-processing
- Automate vendor invoice processing with Power Automate and AI Builder - https://learn.microsoft.com/power-platform/architecture/reference-architectures/vendor-invoice-integration

Topic 2 - Question 13

Hotspot | Domain: Design
topic-2/q13

HOTSPOT -

You are designing an AI strategy for Microsoft Dynamics 365 finance and operations apps. You are evaluating the use of Microsoft Copilot Studio to provide in-app help and guidance based on generative AI general knowledge.

You need to recommend which knowledge sources to include in the generative help and guidance agent. The solution must minimize the risk of generating inaccurate responses.

What should you recommend? To answer, select the appropriate options in the answer area.

Answer area
- **Custom knowledge sources:**
  - [ ] Must be uploaded to the agent
  - [ ] Must be excluded from the agent
  - [ ] Are not supported
- **AI general knowledge:**
  - [ ] Must be enabled for the agent
  - [ ] Must be disabled for the agent
  - [ ] Is not supported

Answer

Custom knowledge sources = Must be uploaded to the agent; AI general knowledge = Must be disabled for the agent

Correct answers: - Custom knowledge sources: **Must be uploaded to the agent** - AI general knowledge: **Must be disabled for the agent**

Why these answers are correct

The agent in question provides in-app help and guidance for finance and operations apps. The requirement is to minimize the risk of inaccurate responses.

Custom knowledge sources must be uploaded to the agent

From Microsoft Learn:

"You can add knowledge to the Copilot help and guidance feature by uploading files to Copilot Studio (in file formats such as PDF, RTF, or Word). If you want to add other types of knowledge (such as from Sharepoint or other data sources), then you must add your own topic in the Copilot for Finance and Operations apps agent in Copilot Studio."

Uploading approved, organization-specific content into the Copilot for Finance and Operations apps agent grounds answers in trusted sources, which reduces the risk of inaccurate responses.

AI general knowledge must be disabled

From Microsoft Learn:

"You can add the capability for Copilot to use AI general knowledge to answer general questions. This content includes information that is part of the language model that is used, web content that is identified through Bing Search, and information from other sources. Copilot uses this information after it exhausts the information from your custom knowledge sources."

"Warning: When the use of general knowledge is enabled, it can increase the richness of Copilot responses. However, it also increases the risk that answers might include incorrect content. Before you publish AI general knowledge to your production system, be sure to carefully consider and test it."

When the requirement is to minimize the risk of inaccurate responses, the Microsoft-documented guidance is to keep AI general knowledge disabled so the agent answers only from grounded, trusted sources.

Microsoft Learn references

- Add knowledge to generative help and guidance with Copilot - https://learn.microsoft.com/dynamics365/fin-ops-core/dev-itpro/copilot/extend-copilot-generative-help
- Generative help and guidance with Copilot (overview) - https://learn.microsoft.com/dynamics365/fin-ops-core/fin-ops/copilot/copilot-generative-help
- Enable generative help and guidance with Copilot - https://learn.microsoft.com/dynamics365/fin-ops-core/dev-itpro/copilot/enable-copilot-generative-help

Topic 2 - Question 14

Hotspot | Domain: Design

topic-2/q14

HOTSPOT -

You need to design a multi-agent solution that will include a custom agent. The solution must meet the following requirements:

Define the rules and constraints that the agent must follow.

Automate a backend process that involves data movement between services and runs independently of the agent's reasoning steps.

What should you include in the design for each requirement? To answer, select the appropriate options in the answer area.

Answer area

- Define rules and constraints:

[ ] Agent flows

[ ] Conversation topics

[ ] Microsoft Power Automate cloud flow

- Automate a backend process:

[ ] Conversation topics

[ ] Microsoft Power Automate cloud flow

[ ] Microsoft Power Pages

Answer

Define rules and constraints = Conversation topics; Automate a backend process = Microsoft Power Automate cloud flow

Correct answers: - Define rules and constraints: **Conversation topics** - Automate a backend process: **Microsoft Power Automate cloud flow**

Why these answers are correct

Conversation topics define rules and constraints

Topics in Copilot Studio are deterministic conversational pathways. They let the maker define exactly what the agent must say, ask, and do, including conditions, branches, variables, and Power Fx expressions. They are the place where the rules and constraints the agent must follow are encoded.

From Microsoft Learn:

"Topics let you design and control how Copilot responds to user interactions. You can use topics to define specific business logic that Copilot must follow when a topic is triggered."

"With topics, you can: guide conversations by defining the sequence of dialogs, messages, tools, agents, and knowledge sources; implement complex workflows by controlling how Copilot responds at each step; use a low-code canvas to define logic and manage variables."

Agent flows are for nonconversational business-process automation. Power Automate cloud flows are general-purpose automation. Neither is the right primitive for defining the agent's conversational rules and constraints.

Power Automate cloud flow for the backend data-movement process

The second requirement is a backend process that moves data between services and runs independently of the agent's reasoning. That description matches a Power Automate cloud flow.

From Microsoft Learn:

"Cloud flows in Power Automate provide a flexible, low-code approach to automating workflows across various applications and services. They offer a wide range of connectors and actions to integrate with different systems and services. Cloud flows are designed for general automation scenarios and can be used independently or with agent flows to create end-to-end automation solutions."

A cloud flow can run on its own schedule or trigger, independent of the agent's reasoning steps. Conversation topics are conversational. Power Pages is a portal builder for external-facing websites and is not the right tool here.

Microsoft Learn references

- Topics overview (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/topics-overview
- Agent flows in Microsoft Copilot Studio FAQ - https://learn.microsoft.com/microsoft-copilot-studio/flows-faqs
- Plan and design integration strategies (cloud flows vs agent flows vs topics) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/integrations
- Power Automate cloud flows overview - https://learn.microsoft.com/power-automate/overview-cloud

Topic 2 - Question 15

Multiple choice | Domain: Plan

topic-2/q15

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.

After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.

A company uses Microsoft 365 and Dynamics 365.

You need to recommend a solution to automatically summarize email threads, generate suggested replies in Microsoft Outlook, and provide meeting preparation summaries that include relevant customer relationship management (CRM) data.

Solution: You recommend Microsoft 365 Copilot for Sales.

Does this meet the goal?

A. Yes
B. No

Answer

A

Correct answer: A. Yes

Why Yes is correct

The requirement is to automatically summarize email threads, generate suggested replies in Microsoft Outlook, and provide meeting preparation summaries that include CRM data. Microsoft 365 Copilot for Sales is built specifically for that combination of capabilities.

From Microsoft Learn:

"Microsoft 365 Copilot for Sales is an AI assistant designed for sellers. The Copilot for Sales app in Outlook gives you recommendations and information to help you stay connected to your customers, reduce data entry, and personalize your engagements."

"With Copilot, you can get related information from your CRM system, such as contact details, account history, and opportunities... You get an overview of recent interactions with your customers, including email summaries, meeting notes, and action items. Use AI capabilities to draft emails, summarize conversations, and generate follow-ups."

"Sales agent integrates Microsoft 365 Copilot in Outlook with your CRM system to help you build customer relationships and close deals faster. Microsoft 365 Copilot summarizes your emails, helping you quickly grasp key details. Sales agent enriches those summaries with sales information from your Dynamics 365 or Salesforce CRM."

That covers all three goals: email-thread summaries, suggested replies in Outlook, and CRM-grounded meeting preparation summaries.

Microsoft Learn references

- Microsoft Outlook experiences for Microsoft 365 Copilot for Sales - https://learn.microsoft.com/copilot/release-plan/2026wave1/copilot-sales/outlook-experiences
- Boost sales efficiency with CRM-enriched email summaries in Outlook - https://learn.microsoft.com/microsoft-sales-copilot/email-summary-premium
- Agents, Copilot, and AI capabilities in Dynamics 365 apps (Sales in Microsoft 365 Copilot) - https://learn.microsoft.com/dynamics365/copilot/ai-get-started
- Copilot in Dynamics 365 Sales overview - https://learn.microsoft.com/dynamics365/sales/copilot-overview

Topic 2 - Question 16

Multiple choice | Domain: Plan

topic-2/q16

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.

After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.

A company uses Microsoft 365 and Dynamics 365.

You need to recommend a solution to automatically summarize email threads, generate suggested replies in Microsoft Outlook, and provide meeting preparation summaries that include relevant customer relationship management (CRM) data.

Solution: You recommend a classic Microsoft Dataverse workflow.

Does this meet the goal?

A. Yes
B. No

Answer

B

Correct answer: B. No

Why No is correct

A classic Microsoft Dataverse workflow is a deterministic, rule-based process designed to automate record updates and business logic in Dataverse tables. It is not a generative AI capability.

The stated requirements are to summarize email threads, generate suggested replies in Outlook, and produce meeting preparation summaries that include CRM data. Those are all generative AI features delivered by Microsoft 365 Copilot and Microsoft 365 Copilot for Sales (with the Sales agent), not by classic Dataverse workflows.

From Microsoft Learn:

"Microsoft 365 Copilot for Sales is an AI assistant designed for sellers... With Copilot, you can get related information from your CRM system... You get an overview of recent interactions with your customers, including email summaries, meeting notes, and action items. Use AI capabilities to draft emails, summarize conversations, and generate follow-ups."

A classic Dataverse workflow cannot summarize emails, generate suggested replies in Outlook, or produce CRM-enriched meeting summaries. The proposed solution does not meet any of the stated goals.

Microsoft Learn references

- Microsoft Outlook experiences for Microsoft 365 Copilot for Sales - https://learn.microsoft.com/copilot/release-plan/2026wave1/copilot-sales/outlook-experiences
- Boost sales efficiency with CRM-enriched email summaries in Outlook - https://learn.microsoft.com/microsoft-sales-copilot/email-summary-premium

- Copilot in Dynamics 365 Sales overview - https://learn.microsoft.com/dynamics365/sales/copilot-overview
- Dataverse classic workflows (background processes, not generative AI) - https://learn.microsoft.com/power-automate/workflow-processes

Topic 2 - Question 17

Multiple choice | Domain: Plan

topic-2/q17

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.

After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.

A company uses Microsoft 365 and Dynamics 365.

You need to recommend a solution to automatically summarize email threads, generate suggested replies in Microsoft Outlook, and provide meeting preparation summaries that include relevant customer relationship management (CRM) data.

Solution: You recommend a Microsoft 365 Copilot agent template.

Does this meet the goal?

A. Yes
B. No

Answer

B

Correct answer: B. No

Why No is correct

A Microsoft 365 Copilot agent template is a starting point for building a custom declarative or custom-engine agent for Microsoft 365 Copilot. It is not the productized Outlook and CRM integration that delivers email summaries, suggested replies in Outlook, and CRM-grounded meeting preparation summaries out of the box.

The Microsoft-recommended product for the stated goals is Microsoft 365 Copilot for Sales, which integrates Copilot in Outlook with Dynamics 365 or Salesforce and provides those exact capabilities natively.

From Microsoft Learn:

"Microsoft 365 Copilot for Sales is an AI assistant that empowers sellers with insights, recommendations, actions, and up-to-date CRM data. It works directly within the Microsoft 365 applications sellers already use daily. Copilot for Sales is a Copilot-first application that works with both Dynamics 365 Sales and Salesforce."

"Sales agent integrates Microsoft 365 Copilot in Outlook with your CRM system to help you build customer relationships and close deals faster. Microsoft 365 Copilot summarizes your emails, helping you quickly grasp key details. Sales agent enriches those summaries with sales information from your Dynamics 365 or Salesforce CRM."

A generic Microsoft 365 Copilot agent template would still require substantial custom build work for Outlook email summaries, suggested replies, and CRM-enriched meeting prep. It is not the right fit and does not meet the goal as designed.

Microsoft Learn references

- Microsoft Outlook experiences for Microsoft 365 Copilot for Sales - https://learn.microsoft.com/copilot/release-plan/2026wave1/copilot-sales/outlook-experiences
- Boost sales efficiency with CRM-enriched email summaries in Outlook - https://learn.microsoft.com/microsoft-sales-copilot/email-summary-premium
- Copilot in Dynamics 365 Sales overview - https://learn.microsoft.com/dynamics365/sales/copilot-overview
- Microsoft 365 Copilot agent templates (declarative agents) - https://learn.microsoft.com/microsoft-365/copilot/extensibility/overview-declarative-agent

Topic 2 - Question 18

Multiple choice | Domain: Deploy

topic-2/q18

You need to design an application lifecycle management (ALM) process for a Microsoft Power Platform environment that contains a solution named Solution1.

Solution1 must include a custom connector for Copilot in Microsoft Dynamics 365 Customer Service. Solution1 must meet the following requirements:

* Ensure that the custom connector can be deployed consistently across environments as part of the ALM process.
* Allow the custom connector to be edited only in the development environment.

What should you include in the design?

A. Add the custom connector to GitHub.
B. Share the custom connector.
C. Create the custom connector in the default solution.
D. Add the custom connector to Solution1.

Answer

D

Correct answer: D. Add the custom connector to Solution1.

Why D is correct

Power Platform application lifecycle management (ALM) treats solutions as the unit of transport between environments. To deploy a custom connector consistently across development, test, and production environments, the connector must be inside the solution that is exported from development and imported into the downstream environments.

From Microsoft Learn:

"Use solutions as containers to transport artifacts and customizations across environments."

"When you create an agent in Copilot Studio, the agent is created within a Power Platform solution. You can create custom solutions to manage your agents across multiple environments, or for pipeline deployments and other application lifecycle management (ALM) scenarios."

"Don't customize outside of a development environment. Always work in the context of solutions... Export and deploy solutions as managed, unless setting up a development environment."

By adding the connector to Solution1 and exporting Solution1 as managed, the connector is deployed consistently across environments and is editable only in the development environment, which matches both stated requirements.

Why the other options are wrong

A. Add the custom connector to GitHub. Source-controlling the connector definition in GitHub doesn't deploy it through Power Platform ALM. Solution-based deployment is still required to package and import the connector into target environments.

B. Share the custom connector. Sharing grants user permissions. It doesn't move the component across environments or restrict edits to development only.

C. Create the custom connector in the default solution. Components in the default solution can't be exported. Microsoft's ALM golden rule is to always work in custom (unmanaged) solutions during development and never customize the default solution for ALM-bound components.

Microsoft Learn references

- Establish an application lifecycle management strategy (Copilot Studio ALM) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/alm
- Create and manage solutions in Copilot Studio - https://learn.microsoft.com/microsoft-copilot-studio/authoring-solutions-overview
- Export and import agents using solutions - https://learn.microsoft.com/microsoft-copilot-studio/authoring-solutions-import-export
- Power Platform solution concepts - https://learn.microsoft.com/power-platform/alm/solution-concepts-alm

Topic 2 - Question 19

Multiple choice | Domain: Design

topic-2/q19

A company uses a Microsoft Copilot Studio agent to automate tasks in a web app.

During testing, you discover that the automation sometimes fails because of frequent changes to the app's user interface.

You need to recommend a solution to ensure that the agent successfully automates the tasks. The solution must minimize changes to the agent.

What should you include in the recommendation?

A. Computer Use in Copilot Studio
B. custom models in Azure AI Studio
C. conversation topics in Copilot Studio
D. an agent flow in Copilot Studio

Answer

A

Correct answer: A. Computer Use in Copilot Studio

Why A is correct

The web app's UI changes frequently, which breaks brittle automations that target specific selectors or coordinates. Computer Use in Copilot Studio is built on a Computer-Using Agent model that combines computer vision and reasoning to interpret the screen at runtime, so it adapts when buttons or layouts shift, with no maker code changes.

From Microsoft Learn:

"Computer use is a feature in Copilot Studio that enables your agent to interact with any system with a graphical user interface (GUI), such as websites or desktop applications."

"Computer use runs on Computer-Using Agents (CUA), an AI model that combines computer vision and reasoning to navigate and interact with GUIs. It adapts to interface changes, so it continues working even if buttons or layouts shift."

That is exactly the requirement: keep the agent working through frequent UI changes, while minimizing changes to the agent.

Why the other options are wrong

B. Custom models in Azure AI Studio introduce model-training effort and don't, on their own, perform GUI automation that adapts to UI changes.

C. Conversation topics in Copilot Studio are conversational dialog steps. They aren't a GUI automation engine and don't help with UI drift in a target web app.

D. An agent flow in Copilot Studio runs deterministic actions through Power Platform connectors. If the target web app has no API and the UI keeps changing, an agent flow doesn't adapt to UI changes the way Computer Use does.

Microsoft Learn references

- Automate web and desktop apps with computer use (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/computer-use
- Computer use release plan - https://learn.microsoft.com/power-platform/release-plan/2025wave1/microsoft-copilot-studio/automate-web-desktop-apps-computer-use
- FAQ for the computer use tool - https://learn.microsoft.com/microsoft-copilot-studio/faqs-computer-use

Topic 2 - Question 20

Multiple choice | Domain: Design

topic-2/q20

A company processes invoices stored across multiple systems in multiple formats.

You need to implement an AI solution to automate the invoice processing. The solution must meet the following requirements:

* Automate multi-step invoice processing tasks, including document analysis, data validation, and approval routing.
* Enable users to interact directly via Microsoft Teams to review and approve invoices.
* Minimize development efforts to define and customize approval workflows.

What should you include in the solution?

A. Azure Document Intelligence in Foundry Tools and Azure Logic Apps
B. a SharePoint agent
C. Microsoft Copilot Studio and AI Builder
D. Azure OpenAI and Azure Functions

Answer

C

Correct answer: C. Microsoft Copilot Studio and AI Builder

Why C is correct

The requirements line up with the Microsoft-documented invoice processing pattern that uses AI Builder for document understanding and Copilot Studio for the agent and approval experience inside Microsoft Teams.

From Microsoft Learn (Streamline document processing with AI Builder):

"AI Builder: Extracts key data from documents using prebuilt or custom models. Power Automate: Orchestrates workflows for document processing. Microsoft Dataverse: Serves as the central data store for extracted document data and tracks document progress through the business process."

From Microsoft Learn (Multistage and AI approvals in agent flows):

"AI approvals are intelligent, automated decision steps in approval workflows. AI approvals use AI to evaluate approval requests against your business rules and return an 'Approved' or 'Rejected' decision with a rationale."

"Responding to approvals... Users assigned to approvals can respond through: Microsoft Teams approvals app, Outlook, Power Automate portal."

AI Builder handles document analysis and data validation across multiple invoice formats. Copilot Studio (with agent flows and approvals) automates multi-step processing, exposes the agent and approval interaction inside Microsoft Teams, and minimizes development effort through low-code configuration of approval workflows.

Why the other options are wrong

A. Azure Document Intelligence in Foundry Tools and Azure Logic Apps is a pro-developer pattern that requires more development effort than the low-code Copilot Studio plus AI Builder pattern. It also does not provide the native Teams agent and approval experience the question requires.

B. A SharePoint agent is a content-grounded Q&A agent over SharePoint sites. It is not designed to extract invoice fields from many formats, validate data, route approvals, or drive multi-step processing.

D. Azure OpenAI and Azure Functions is a pro-code pattern that requires significant custom development. It does not minimize development effort or provide built-in approval workflows.

Microsoft Learn references

- Streamline document processing with AI Builder - https://learn.microsoft.com/power-platform/architecture/reference-architectures/ai-document-processing
- Automate vendor invoice processing with Power Automate and AI Builder - https://learn.microsoft.com/power-platform/architecture/reference-architectures/vendor-invoice-integration
- Multistage and AI approvals in agent flows - https://learn.microsoft.com/microsoft-copilot-studio/flows-advanced-approvals
- FAQ for AI Approvals - https://learn.microsoft.com/microsoft-copilot-studio/faqs-ai-approvals

Topic 2 - Question 21

Multiple choice | Domain: Design

topic-2/q21

You need to design a Microsoft Copilot Studio agent for customer support.

The agent must securely retrieve product warranty data from a REST API. The solution must minimize development effort.

What should you include in the design?

A. Export the agent as a managed solution and customize the agent in Power Apps.
B. Create a custom connector in Copilot Studio and use the connector to call the API.
C. Use a Microsoft Power Automate desktop flow to screen scrape the warranty data.
D. Add the warranty data to the Fallback topic.

Answer

B

Correct answer: B. Create a custom connector in Copilot Studio and use the connector to call the API.

Why B is correct

A custom connector in Copilot Studio is a low-code wrapper around a REST API. It supports standard authentication (Microsoft Entra ID, OAuth 2.0, basic auth, API key), can be reused across agents, and is the Microsoft-recommended way to securely call a REST API from a Copilot Studio agent with minimal development effort.

From Microsoft Learn:

"A custom connector is a wrapper around a REST API that allows Logic Apps, Power Automate, Power Apps, or Copilot Studio to communicate with that REST or SOAP API."

"Custom connectors: When there's no prebuilt connector available, you can build your own connector for a service. They're a no-code or low-code wrapper for REST APIs."

"Use one of these standard authentication methods for your APIs and connectors (Microsoft Entra ID is recommended): Generic OAuth 2.0; OAuth 2.0 for specific services... Basic authentication; API Key."

That covers both the secure REST API call and the minimal-development-effort requirement.

Why the other options are wrong

A. Export the agent as a managed solution and customize the agent in Power Apps. Exporting and re-customizing in Power Apps doesn't securely call a REST API and doesn't minimize development effort.

C. Use a Microsoft Power Automate desktop flow to screen scrape the warranty data. Screen scraping is brittle, doesn't use the REST API, and isn't a secure or low-effort way to retrieve warranty data.

D. Add the warranty data to the Fallback topic. The Fallback topic captures unrecognized intent. It doesn't retrieve live data from a REST API and isn't the design pattern for secure API integration.

Microsoft Learn references

- Use Power Platform connectors as tools (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/advanced-connectors
- Custom connectors overview - https://learn.microsoft.com/connectors/custom-connectors/

- Plan and design integration strategies (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/integrations
- Extend your agent with tools from a REST API (preview) - https://learn.microsoft.com/microsoft-copilot-studio/agent-extend-action-rest-api

Topic 2 - Question 22

Drag and drop | Domain: Design

topic-2/q22

DRAG DROP -

A company has an ecommerce support portal that uses Microsoft Dataverse.

You are designing a Microsoft Copilot Studio agent for the portal. The agent must meet the following requirements:

* Respond with a default help message when the user input is unclear.
* Initiate external processes, such as retrieving the order status, when users make specific requests.

Generative orchestration will be enabled for the solution.

You need to recommend a feature for each requirement.

What should you recommend? To answer, drag the appropriate features to the correct requirements. Each feature may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

Drag sources

- A tool (connector)
- A trigger phrase
- The Fallback topic
- A skill
- The Escalate topic

Drop targets

1. Respond with a default help message when the user input is unclear - drop here
2. Initiate external processes when requested - drop here

Answer

Default help when input is unclear = The Fallback topic; Initiate external processes = A tool (connector)

Correct answers: - Respond with a default help message when the user input is unclear: **The Fallback topic** - Initiate external processes when requested: **A tool (connector)**

Why these answers are correct

Fallback topic for unclear input

From Microsoft Learn:

"The Fallback system topic triggers when a Copilot Studio agent doesn't understand a user utterance and doesn't have enough confidence to trigger any of the existing topics."

"If the agent can't determine the user's intent, it prompts the user again. After two prompts, the custom or

Copilot agent escalates to a live representative through a system topic called Escalate."

The Fallback topic is the documented place to send a default help message when the user's input is unclear, even with generative orchestration enabled.

A tool (connector) to initiate external processes

In Copilot Studio with generative orchestration enabled, tools are how the agent reaches out to external systems to perform actions, like retrieving an order status. Connectors register as tools the orchestrator can choose at runtime.

From Microsoft Learn:

"Tools and connectors: External actions, APIs, or automation flows that the agent can call as part of a plan. Each tool has a defined interface: input parameters (with expected types), output variables, and possibly error conditions. They're essentially the agent's 'skills' for performing operations, such as looking up an order, sending an email, or running a script."

A trigger phrase is for classic orchestration topic routing, not for invoking external processes. The Escalate topic is for live-representative handoff. A skill is a Bot Framework pattern, not the Microsoft-recommended choice for new generative-orchestration designs that need to call connectors.

Microsoft Learn references

- Use the Fallback topic - https://learn.microsoft.com/microsoft-copilot-studio/guidance/fallback-topic
- Configure the system fallback topic - https://learn.microsoft.com/microsoft-copilot-studio/authoring-system-fallback-topic
- Apply generative orchestration capabilities (tools and connectors) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/generative-orchestration
- Use Power Platform connectors as tools - https://learn.microsoft.com/microsoft-copilot-studio/advanced-connectors

Topic 2 - Question 23

Multi-select | Domain: Design

topic-2/q23

A company uses Microsoft Dynamics 365 to manage service operations. Dispatchers coordinate service requests, and technicians perform scheduled on-site work.

You need to design a solution that will use Microsoft Copilot to improve the efficiency of the service operations. The solution must meet the following requirements:

* Provide AI-driven assistance to help staff organize and resolve work orders.
* Deliver contextual AI support to frontline workers as they prepare for and complete customer appointments.

Which two components should you include in the design? Each correct answer presents part of the solution.

A. Copilot Service workspace
B. Copilot in Outlook
C. Dynamics 365 Customer Service
D. Copilot in Customer Service
E. Copilot in Field Service
F. the Dynamics 365 Field Service mobile app

Select all that apply.

Answer

A, E

Correct answers: A. Copilot Service workspace and E. Copilot in Field Service

Why A and E are correct

The two requirements split cleanly between dispatcher work-order organization and frontline-worker support during appointments. Microsoft documents two distinct Copilot experiences for each.

Copilot Service workspace for dispatcher productivity

From Microsoft Learn:

"Copilot Service workspace helps customer service representatives (service representatives or representatives) increase productivity with a browser-like, tabbed experience. Service representatives can use the app to work on multiple cases and conversations."

"The application uses artificial intelligence in productivity tools like Smart Assist to identify similar cases and relevant articles, thereby boosting representative productivity. Features such as scripts and macros provide representatives with guidance and resources to automate repetitive tasks."

Copilot Service workspace is the AI-powered workspace where dispatchers and representatives organize and act on cases and work orders.

Copilot in Field Service for frontline workers

From Microsoft Learn:

"Copilot in Dynamics 365 Field Service helps boost productivity and minimize manual work by generating work order summaries, creating inspection templates, and enabling frontline workers using natural language."

"Technicians can use Copilot to summarize their work orders within the Field Service mobile application. The summary gives them meaningful context for the work they're about to perform. It can include notes, diagnostic information, key events in the work order lifecycle, and recommendations."

Copilot in Field Service is the documented capability that gives frontline workers contextual AI support before and during customer appointments, including in the Field Service mobile app and the web application.

Why the other options are wrong

B. Copilot in Outlook drafts and summarizes email. It is not the Microsoft-recommended capability for organizing work orders or supporting frontline appointment work.

C. Dynamics 365 Customer Service is the underlying customer service application; the question asks for the Copilot capability layered on it, which is Copilot Service workspace.

D. Copilot in Customer Service focuses on case-handling assistance for customer service representatives. It overlaps with A but does not deliver field-worker support; the documented service-operations workspace experience is Copilot Service workspace.

F. The Dynamics 365 Field Service mobile app is the mobile client. The Copilot capability inside it that delivers contextual AI is Copilot in Field Service. The question asks for the Copilot component, not the host app.

Microsoft Learn references

- Get started with Copilot Service workspace - https://learn.microsoft.com/dynamics365/customer-service/implement/csw-overview
- Manage Copilot features in Customer Service - https://learn.microsoft.com/dynamics365/customer-service/administer/configure-copilot-features
- Copilot features in Dynamics 365 Field Service -

https://learn.microsoft.com/dynamics365/field-service/copilot-overview
- Summarize work orders with Copilot in the Field Service mobile app - https://learn.microsoft.com/dynamics365/release-plan/2025wave1/service/dynamics365-field-service/summarize-work-orders-copilot-new-mobile-experience

Topic 2 - Question 24

Multiple choice | Domain: Design

topic-2/q24

A company uses Microsoft Foundry agents.

You need to ensure that an agent can dynamically use external tools at runtime without updating the agent.

What should you include in the solution?

A. a Microsoft Foundry hub
B. a Model Context Protocol (MCP) server
C. Azure AI Search
D. Microsoft Copilot Studio

Answer

B

Correct answer: B. a Model Context Protocol (MCP) server

Why B is correct

Microsoft Foundry agents (and Copilot Studio agents) use MCP to discover and call external tools at runtime through a standardized protocol. When tools are added or changed on an MCP server, the agent dynamically picks them up without an agent update.

From Microsoft Learn:

"Each tool or resource that a connected MCP server publishes is automatically available for use in Copilot Studio. The server provides the name, description, inputs, and outputs. When you update or remove tools and resources on the MCP server, Copilot Studio dynamically reflects these changes. This process ensures users always have the latest versions and removes obsolete tools and resources."

"MCP is especially valuable when upstream APIs change frequently. Instead of updating every agent that consumes the API, you modify the definition once on the MCP server, and all agents automatically use the updated version without republishing."

That is precisely the requirement: dynamically use external tools at runtime without updating the agent.

Why the other options are wrong

A. A Microsoft Foundry hub is a top-level Foundry workspace organization construct. It does not, by itself, expose external tools to agents at runtime.

C. Azure AI Search is a retrieval service for grounded knowledge. It returns search results, not callable external tools that change dynamically.

D. Microsoft Copilot Studio is an authoring product. It is not the dynamic-tool mechanism the question is asking for; MCP is.

Microsoft Learn references

- Extend your agent with Model Context Protocol - https://learn.microsoft.com/microsoft-copilot-studio/agent-extend-action-mcp
- Use agent tools to extend, automate, and enhance your agents (when to use MCP) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/agent-tools
- Tool best practices for Microsoft Foundry Agent Service (MCP support) - https://learn.microsoft.com/azure/foundry/agents/concepts/tool-best-practice
- What is Microsoft Foundry Agent Service? - https://learn.microsoft.com/azure/foundry/agents/overview

Topic 2 - Question 25

Multiple choice | Domain: Design

topic-2/q25

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.

After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.

A company has a Microsoft Dynamics 365 Sales environment that has Microsoft Copilot enabled.

You need to customize Copilot by tailoring how opportunity summaries are generated or how they are presented to users.

Solution: You add fields to the opportunity summary.

Does this meet the goal?

A. Yes
B. No

Answer

A

Correct answer: A. Yes

Why Yes is correct

The requirement is to customize Copilot in Dynamics 365 Sales by tailoring how opportunity summaries are generated or presented. Microsoft documents adding fields as a supported, intended customization for opportunity summaries.

From Microsoft Learn:

"By default, Copilot uses a set of predefined fields to generate summaries, a list of recent changes for accounts, leads, and opportunities, and prepare for meetings. You can add other fields from lead, opportunity, account, and related tables to make the summaries and recent changes list more relevant for your business."

"Select Add fields. Select at least four fields, up to a maximum of 15 for summaries and 10 for recent changes. You can add fields from the current table or related tables."

"Copilot generates its summary from a set of predefined fields. However, other fields might be more important to you... Work with your Dynamics 365 Sales administrator to add those fields to the summary."

Adding fields to the opportunity summary is the documented, supported way to tailor how Copilot generates the summary content. The proposed solution meets the goal.

Microsoft Learn references

- Configure fields for generating summaries and recent changes - https://learn.microsoft.com/dynamics365/sales/copilot-configure-summary-fields
- Summarize records with Copilot - https://learn.microsoft.com/dynamics365/sales/copilot-summarize-records
- Customize Copilot in Dynamics 365 Sales - https://learn.microsoft.com/dynamics365/sales/extend-copilot-chat
- FAQ about Copilot in Dynamics 365 Sales - https://learn.microsoft.com/dynamics365/sales/sales-copilot-faq

Topic 2 - Question 26

Multiple choice | Domain: Design

topic-2/q26

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.

After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.

A company has a Microsoft Dynamics 365 Sales environment that has Microsoft Copilot enabled.

You need to customize Copilot by tailoring how opportunity summaries are generated or how they are presented to users.

Solution: You build Microsoft Power Automate flows to trigger customized Copilot summaries.

Does this meet the goal?

A. Yes
B. No

Answer

B

Correct answer: B. No

Why No is correct

Power Automate flows automate business processes and can update records, send notifications, and orchestrate connector actions. They do not customize how Copilot in Dynamics 365 Sales generates or presents opportunity summaries.

The Microsoft-documented way to tailor opportunity summary generation and presentation is to configure summary fields (and related-info sections) in the Sales Hub Copilot settings, or extend Copilot through Copilot Studio with glossary, synonyms, and topic-level customization.

From Microsoft Learn:

"By default, Copilot uses a set of predefined fields to generate summaries... You can add other fields from lead, opportunity, account, and related tables to make the summaries and recent changes list more relevant for your business."

A Power Automate flow doesn't change which fields Copilot uses for summary generation, doesn't reshape the summary widget, and isn't a documented mechanism for customizing the opportunity summary. The proposed solution does not meet the goal.

Microsoft Learn references

- Configure fields for generating summaries and recent changes - https://learn.microsoft.com/dynamics365/sales/copilot-configure-summary-fields
- Summarize records with Copilot - https://learn.microsoft.com/dynamics365/sales/copilot-summarize-records
- Customize Copilot in Dynamics 365 Sales - https://learn.microsoft.com/dynamics365/sales/extend-copilot-chat
- FAQ about Copilot in Dynamics 365 Sales - https://learn.microsoft.com/dynamics365/sales/sales-copilot-faq

Topic 2 - Question 27

Multiple choice | Domain: Design

topic-2/q27

Note: This section contains one or more sets of questions with the same scenario and problem. Each question presents a unique solution to the problem. You must determine whether the solution meets the stated goals. More than one solution in the set might solve the problem. It is also possible that none of the solutions in the set solve the problem.

After you answer a question in this section, you will NOT be able to return. As a result, these questions do not appear on the Review Screen.

A company has a Microsoft Dynamics 365 Sales environment that has Microsoft Copilot enabled.

You need to customize Copilot by tailoring how opportunity summaries are generated or how they are presented to users.

Solution: You configure AI Builder lead scoring models to influence opportunity summaries.

Does this meet the goal?

A. Yes
B. No

Answer

B

Correct answer: B. No

Why No is correct

AI Builder lead scoring models predict lead-to-opportunity conversion likelihood. They produce a score and reasons for that score. They are not used to customize how Copilot in Dynamics 365 Sales generates or presents opportunity summaries.

The Microsoft-documented mechanisms for customizing opportunity summaries are:

- Configure summary fields and related-info sections in the Sales Hub Copilot settings, including which fields, sections (enriched key info, competitor insights, product insights, quote insights), and the opportunity summary widget.
- Extend Copilot in Dynamics 365 Sales through Copilot Studio (glossary, synonyms, topics).

From Microsoft Learn:

"By default, Copilot uses a set of predefined fields to generate summaries, a list of recent changes for accounts, leads, and opportunities, and prepare for meetings. You can add other fields from lead, opportunity, account, and related tables to make the summaries and recent changes list more relevant for your business."

Lead scoring models do not influence opportunity summary generation. The proposed solution does not meet the goal.

Microsoft Learn references

- Configure fields for generating summaries and recent changes - https://learn.microsoft.com/dynamics365/sales/copilot-configure-summary-fields
- Customize Copilot in Dynamics 365 Sales - https://learn.microsoft.com/dynamics365/sales/extend-copilot-chat
- AI Builder predictive lead scoring overview - https://learn.microsoft.com/dynamics365/sales/predictive-lead-scoring
- Summarize records with Copilot - https://learn.microsoft.com/dynamics365/sales/copilot-summarize-records

Topic 2 - Question 28

Multiple choice | Domain: Plan | Case study: Contoso

topic-2/q28

What should you recommend to assist the CTO with the prebuilt agent selection process?

A. Agent management
B. Copilot Studio
C. Lifecycle Services (LCS)
D. Immersive Home

Answer

A

Correct answer: A. Agent management

Why A is correct

The Contoso case study (questions 28 and 29) says the CTO "wants to view available agent templates to identify which agent will add the most business value" for Dynamics 365 Supply Chain Management. The Microsoft-documented experience for discovering, previewing, configuring, and activating prebuilt Dynamics 365 finance and operations agents is the agent management feature, accessed through Immersive Home in finance and operations apps.

From Microsoft Learn:

"By using the agent management feature in finance and operations apps, you can use autonomous, AI-powered agents to perform predefined tasks in your business ecosystem. Users can discover, configure, and manage agents that automate routine operational tasks."

"Use the following tabs on the Agent page to discover and activate new agents: Library - Browse available agent templates. Preview agent capabilities. Select and configure new agents for deployment. Manage - View agents that are currently active. Edit existing agent configurations. Activate or deactivate individual agent tasks."

The Library tab in Agent management is exactly where the CTO previews available prebuilt agent templates and selects the one that adds the most business value.

Why the other options are wrong

B. Copilot Studio is for building and customizing custom AI agents. The CTO is selecting a prebuilt agent from Dynamics 365 Supply Chain Management; the discovery and activation experience is Agent management, not Copilot Studio.

C. Lifecycle Services (LCS) is the legacy environment management portal for finance and operations apps. Microsoft is moving environment management to the Power Platform admin center, and LCS is not where the CTO browses prebuilt agent templates. (From Microsoft Learn: "Starting February 2026, new customers can't create projects in Microsoft Dynamics Lifecycle Services for Microsoft Dynamics 365 Finance, Microsoft Dynamics 365 Human Resources, Microsoft Dynamics 365 Supply Chain Management, and Microsoft Dynamics 365 Project Operations.")

D. Immersive Home is the navigation hub that gives access to agent-related pages, but the documented feature for selecting prebuilt agent templates is Agent management. Immersive Home is the entry point, not the selection feature itself.

Microsoft Learn references

- Enable agent management (production ready preview) - https://learn.microsoft.com/dynamics365/fin-ops-core/fin-ops/copilot/agent-mgmt
- Agent management release notes - https://learn.microsoft.com/dynamics365/release-plan/2024wave2/finance-supply-chain/dynamics365-finance/agent-management
- Agents, Copilot, and AI capabilities in Dynamics 365 apps - https://learn.microsoft.com/dynamics365/copilot/ai-get-started
- Immersive Home (production ready preview) - https://learn.microsoft.com/dynamics365/fin-ops-core/fin-ops/copilot/immersive-home

Topic 2 - Question 29

Hotspot | Domain: Design | Case study: Contoso

topic-2/q29

HOTSPOT -

What should you include in the custom AI agent design to meet the R&D product specifications and the compliance information requirements? To answer, select the appropriate options in the answer area.

Answer area

- To expose the data to the agent, create:

[ ] an Azure AI Bot Service channel
[ ] a custom connector
[ ] a custom OData entity
[ ] the Semantic Kernel

- Add to the agent:

[ ] an event trigger
[ ] the MCP server
[ ] a REST API
[ ] a tool

Answer

Expose data to the agent = a custom connector; Add to the agent = the MCP server

Correct answers: - To expose the data to the agent, create: **a custom connector** - Add to the agent: **the MCP server**

Why these answers are correct

The Contoso case study states that the R&D department already has a custom Model Context Protocol (MCP) server containing comprehensive product specifications and compliance data. The custom AI agent must answer questions about product specifications using existing technologies, and it must be designed to eventually connect to other agents that can be selected based on description (a generative orchestration design pattern).

Expose the data to the agent: a custom connector

In Copilot Studio, MCP servers are connected to agents through a custom connector. The connector wraps the MCP server's HTTPS endpoint and includes the `x-ms-agentic-protocol: mcp-streamable-1.0` property required for MCP communication.

From Microsoft Learn:

"The Copilot Studio agent connects to MCP servers by using a custom connector."

"Create a new custom connector using the Create from blank option... Set Scheme to HTTPS. Set Host to the CONTAINER_APP_URL value... Toggle Swagger editor to enter the editor view. Expose a POST method at the root path with a custom x-ms-agentic-protocol: mcp-streamable-1.0 property. This property is required for the custom connector to interact with the API by using the MCP protocol."

A custom connector (not a Bot Service channel, custom OData entity, or Semantic Kernel) is the documented way to expose the existing MCP server's data to the Copilot Studio agent.

Add to the agent: the MCP server

After the custom connector is in place, the maker adds the MCP server itself as a tool on the agent. Each MCP tool and resource published by the server then becomes available, by name and description, for the agent to choose at runtime.

From Microsoft Learn:

"Each tool or resource that a connected MCP server publishes is automatically available for use in Copilot Studio. The server provides the name, description, inputs, and outputs. When you update or remove tools and resources on the MCP server, Copilot Studio dynamically reflects these changes."

"Add MCP server tools and resources to your agent so that your Copilot Studio agent can use them."

That fits the case study's requirements to use the existing R&D MCP server, integrate with existing technologies, and select capabilities based on description for conversational interactions.

An event trigger is for autonomous agent responses to events, a REST API tool would require building a separate tool against an OpenAPI spec rather than using the existing MCP server, and a generic tool isn't the specific named option that maps to the MCP server scenario.

Microsoft Learn references

- Extend your agent with Model Context Protocol (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/agent-extend-action-mcp
- Connect your agent to an existing MCP server - https://learn.microsoft.com/microsoft-copilot-studio/mcp-add-existing-server-to-agent
- Add tools and resources from a Model Context Protocol (MCP) server to your agent - https://learn.microsoft.com/microsoft-copilot-studio/mcp-add-components-to-agent
- Deploy a self-hosted remote Azure MCP Server and connect to it using Copilot Studio (custom connector pattern) - https://learn.microsoft.com/azure/developer/azure-mcp-server/how-to/deploy-remote-mcp-server-copilot-studio

Topic 3

42 questions

Topic 3 - Question 1

Multiple choice | Domain: Deploy | Case study: Fabrikam

topic-3/q01

Which tool should you recommend to address the sensitive information concerns in the sales process?

A. the Analytics tab in Microsoft Copilot Studio
B. Model Context Protocol (MCP)
C. Application Insights
D. Microsoft Foundry Tracing UI
E. Monitoring in Microsoft Foundry

Answer

C

Correct answer: C. Application Insights

Why C is correct

The case study describes a single low-code Copilot Studio agent grounded in Dataverse. Two requirements steer the answer:

1. "Detailed telemetry must be logged for the first created AI agent to help troubleshoot and optimize the agent during the initial AI agent adoption process."
2. "Any sensitive information, such as user IDs and names, shared via the AI agent must be tracked for future auditing."

Application Insights is the integrated telemetry option for Copilot Studio agents. Microsoft Learn shows that Copilot Studio captures detailed telemetry (conversations, exceptions, tool usage, latency, custom events) into Application Insights, and the Copilot Studio dashboard workbook in Application Insights surfaces those signals for troubleshooting and optimization. That telemetry is queryable in Log Analytics, which gives auditors the durable, indexable trail needed to track who shared which sensitive values with the agent.

Why the other options are wrong

A. The Analytics tab in Microsoft Copilot Studio shows aggregate KPIs (engagement, resolution, escalation, CSAT, generated answer quality). It does not capture per-message telemetry suitable for auditing sensitive information at the user level.

B. Model Context Protocol (MCP) is an open protocol for connecting agents to external tools and data sources. It is not an audit or telemetry tool.

D. Microsoft Foundry Tracing UI traces Microsoft Foundry agents (prompt, hosted, workflow agents) using OpenTelemetry. The Fabrikam agent is built in Copilot Studio with Dataverse, so Foundry Tracing UI is not the integrated telemetry surface for it.

E. Monitoring in Microsoft Foundry also targets Foundry agents and Foundry-hosted models. It is not the path for a Copilot Studio agent.

Microsoft Learn references

- Capture telemetry with Application Insights (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/advanced-bot-framework-composer-capture-telemetry
- Audit Copilot Studio activities in Microsoft Purview - https://learn.microsoft.com/microsoft-copilot-studio/admin-logging-copilot-studio
- Analytics overview for Copilot Studio - https://learn.microsoft.com/microsoft-copilot-studio/analytics-overview

Topic 3 - Question 2

Hotspot | Domain: Plan | Case study: Fabrikam

topic-3/q02

HOTSPOT -

Which existing tool and data should you use to gather the required metrics for stakeholder signoff for the AI agents? To answer, select the appropriate options in the answer area.

Answer area

- Tool:

[ ] Microsoft Foundry
[ ] Azure Resource Monitor (ARM)
[ ] Dynamics 365 Sales
[ ] Microsoft Copilot Studio

- Data required for the tool:

[ ] the cumulative time spent on the task over the past year
[ ] the current cost to complete the tasks per instance
[ ] the current time to complete the task today per instance
[ ] the current cost of the Dynamics 365 Sales licenses

Answer

Microsoft Copilot Studio; the current cost to complete the tasks per instance

Correct answer: - **Tool: Microsoft Copilot Studio** - **Data required for the tool: the current cost to complete the tasks per instance**

Why these selections are correct

The case study requires "the return on investment (ROI) of switching from the current process to the future process" for stakeholder sign-off, and the agent in question is the low-code Copilot Studio agent grounded in Dataverse. Copilot Studio includes a built-in **Savings calculator** on the Analytics page that estimates ROI for the agent. The calculator compares the agent's runs against the existing manual method on a per-run or per-tool basis and reports time and money saved over the selected period.

Microsoft Learn says: "Use the Savings calculator to estimate your savings per run or to calculate the savings of one or more of the tools your agent uses." To compute that comparison, you enter "the estimated cost required for another method to complete the task of the agent." That maps directly to the **current cost to complete the tasks per instance** baseline.

Microsoft Foundry, Azure Resource Monitor, and Dynamics 365 Sales do not provide a comparable built-in ROI calculator for the Copilot Studio agent in the case study. Cumulative time over the past year, current time per instance, and Dynamics 365 Sales license cost do not match the calculator's expected input, which is the per-instance cost of the existing method to complete the task.

Microsoft Learn references

- Analyze time and cost savings for agents (Savings calculator) - https://learn.microsoft.com/microsoft-copilot-studio/analytics-cost-savings
- Estimate savings per run for an agent - https://learn.microsoft.com/microsoft-copilot-studio/analytics-cost-savings#estimate-savings-per-run-for-an-agent
- Analyze agent return on investment (release plan) - https://learn.microsoft.com/power-platform/release-plan/2025wave1/microsoft-copilot-studio/analyze-agent-return-investment

Topic 3 - Question 3

Multiple choice | Domain: Deploy | Case study: Fabrikam

topic-3/q03

Which tool should you recommend to help secure funding for future AI agent development?

A. Evaluations in Microsoft Foundry
B. the Azure Cost Optimization workbook
C. Azure Operator Insights
D. the Analytics tab in Microsoft Copilot Studio
E. Direct Preference Optimization (DPO)

Answer

D

Correct answer: D. the Analytics tab in Microsoft Copilot Studio

Why D is correct

The case study states: "Sales managers must report on the adoption of the AI agent to key Fabrikam stakeholders on a monthly basis," and "funding for future AI initiatives will depend on demonstrating an increase in AI adoption month over month."

The Analytics tab in Copilot Studio is the built-in surface for measuring agent adoption and usage. Microsoft Learn describes it as the place to view "summaries of agent usage, conversation volume, and engagement," along with KPIs (engagement rate, resolution rate, escalation rate, CSAT, savings, generated answer quality, themes). Sales managers can review monthly trend lines on the Analytics page to demonstrate adoption growth and justify continued AI investment.

Why the other options are wrong

A. Evaluations in Microsoft Foundry measures AI quality, safety, and groundedness for Foundry-hosted models and agents. It is not designed to surface adoption trends for a Copilot Studio agent.

B. The Azure Cost Optimization workbook tracks Azure resource spend. It does not measure agent adoption or usage trends.

C. Azure Operator Insights is for telecommunications operators analyzing network data. It is not relevant to Copilot Studio adoption tracking.

E. Direct Preference Optimization (DPO) is a fine-tuning technique for language models. It does not measure or report adoption.

Microsoft Learn references

- Analytics overview (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/analytics-overview
- Analyze conversational agent effectiveness - https://learn.microsoft.com/microsoft-copilot-studio/analytics-improve-agent-effectiveness
- Measure and improve agent performance with KPIs and analytics - https://learn.microsoft.com/microsoft-copilot-studio/guidance/analytics

Topic 3 - Question 4

Multiple choice | Domain: Design

topic-3/q04

You are designing a Microsoft Copilot Studio agent that uses a custom Microsoft Foundry model to generate responses.

You need to ensure that the agent can securely connect to and invoke the custom model during user interactions.

What should you include in the design?

A. Configure the agent to use classic orchestration.
B. Create a connection to Microsoft Foundry in the agent.
C. Add the Microsoft Foundry model as a Copilot Studio skill.
D. Create a custom engine agent.

Answer

B

Correct answer: B. Create a connection to Microsoft Foundry in the agent.

Why B is correct

Copilot Studio supports invoking custom Microsoft Foundry models through a native connection to Azure AI Foundry. Microsoft Learn's "Bring your own model for your prompts" article shows the pattern: in the prompt configuration, select the Model field, choose the plus sign to "connect to a model from Azure AI Foundry," and enter the deployment and base model names. The connection authenticates and routes invocations securely to the Foundry-hosted model. The same path is used to attach Foundry agents through the "Add an agent" dialog with a Microsoft Foundry connection.

Why the other options are wrong

A. Configure the agent to use classic orchestration. Classic orchestration controls how topics are matched (deterministic vs. generative). It does not establish secure connectivity to a Foundry-hosted model.

C. Add the Microsoft Foundry model as a Copilot Studio skill. Skills are bot-to-bot integration patterns from Bot Framework. The current, supported integration for Foundry models is the Azure AI Foundry connection in prompt builder, not a skill registration.

D. Create a custom engine agent. Custom engine agents are pro-code agents (Teams AI, Microsoft 365 Agents SDK, Foundry SDK) that bring your own model and orchestrator. The question is about a Copilot Studio agent calling a Foundry model, not replacing the Copilot Studio engine.

Microsoft Learn references

- Bring your own model for your prompts (Copilot Studio + Azure AI Foundry) - https://learn.microsoft.com/microsoft-copilot-studio/bring-your-own-model-prompts
- Connect to a Microsoft Foundry agent (preview) - https://learn.microsoft.com/microsoft-copilot-studio/add-agent-foundry-agent
- Use your own generative AI model from Azure AI Foundry in prompt builder - https://learn.microsoft.com/power-platform/release-plan/2025wave1/ai-builder/use-own-generative-ai-model-azure-ai-foundry-prompt-builder

Topic 3 - Question 5

Multiple choice | Domain: Deploy

topic-3/q05

You are designing an AI business solution that contains the following components:

A Microsoft Power Automate workflow

A Microsoft Copilot Studio agent

A Microsoft Dataverse database -

A Microsoft Power Apps app -

As part of the application lifecycle management (ALM) process, you plan to package the components, so that they can be deployed to other environments as a group.

You need to recommend a solution that supports versioning, dependencies, and deployments.

What should you include in the recommendation?

A. GitHub Actions
B. Azure DevOps
C. Microsoft Power Platform solutions

Answer

C

Correct answer: C. Microsoft Power Platform solutions

Why C is correct

Power Platform solutions are the packaging mechanism for ALM across Power Apps, Power Automate, Microsoft Dataverse, and Microsoft Copilot Studio. Microsoft Learn defines solutions as "the mechanism for implementing ALM; you use them to distribute components across environments through export and import. A component represents an artifact used in your application... Anything that can be included in a solution is a component, such as tables, columns, canvas and model-driven apps, Power Automate flows, agents, charts, and plug-ins."

Solutions track component dependencies, support versioning, and can be exported as managed or unmanaged for deployment between environments. That matches all three requirements (versioning, dependencies, deployments) for the four components listed in the question.

Why the other options are wrong

A. GitHub Actions automates CI/CD and can deploy Power Platform artifacts, but it does not itself package the components or track dependencies. It typically operates on top of solutions.

B. Azure DevOps provides CI/CD pipelines and source control. Like GitHub Actions, it automates the deployment of solutions but is not the packaging mechanism itself.

Microsoft Learn references

- Overview of ALM with Microsoft Power Platform - https://learn.microsoft.com/power-platform/alm/overview-alm
- Solution concepts (ALM) - https://learn.microsoft.com/power-platform/alm/solution-concepts-alm
- Create and manage solutions in Copilot Studio -

- https://learn.microsoft.com/microsoft-copilot-studio/authoring-solutions-overview
- Dependency tracking for solution components - https://learn.microsoft.com/power-platform/alm/dependency-tracking-solution-components

Topic 3 - Question 6

Multi-select | Domain: Deploy

topic-3/q06

A company has Microsoft 365 Copilot agents.

You need to design a security solution for the agents. The solution must meet the following requirements:

Identify and mitigate potential risks that relate to AI use.

Protect AI apps and the sensitive data processed or generated by the agents.

Support responsible AI governance by retaining and logging interactions, detecting policy violations, and investigating incidents.

Which two components should you include in the design? Each correct answer presents part of the solution.

A. Microsoft Purview
B. Azure AI Content Safety
C. role-based access control (RBAC) in Microsoft Foundry
D. Microsoft Defender

Select all that apply.

Answer

A, D

Correct answer: A. Microsoft Purview and D. Microsoft Defender

Why A and D are correct

The three requirements break cleanly across the two services Microsoft positions for AI security and governance.

Microsoft Defender (for Cloud / Defender XDR for AI) identifies and mitigates risks that relate to AI use. Microsoft Learn describes Defender for Cloud's AI security posture management and AI threat protection: discovering generative AI workloads and agents (including Copilot Studio agents through Defender for Cloud Apps), surfacing vulnerabilities and attack paths, and detecting threats such as prompt injection and abnormal execution against Microsoft Foundry workloads.

Microsoft Purview protects AI apps and the data they process and supports responsible AI governance. Microsoft Learn calls out DSPM for AI, DLP, sensitivity labels, audit, communication compliance, and insider risk management as the controls for monitoring AI interactions, retaining and logging prompts and responses, detecting policy violations, and investigating incidents.

Together they cover all three requirements: risk identification (Defender), data protection (Purview), and responsible AI governance through retention, logging, policy detection, and investigation (Purview).

Why the other options are wrong

B. Azure AI Content Safety filters harmful content in prompts and responses (violence, hate, sexual, self-harm, jailbreak detection). It is one piece of model safety, but it does not log interactions, detect policy violations across the broader environment, or investigate incidents.

C. Role-based access control (RBAC) in Microsoft Foundry restricts who can configure or fine-tune Foundry resources. It is an identity control, not a risk-mitigation, data-protection, or incident-investigation platform for Microsoft 365 Copilot agents.

Microsoft Learn references

- Microsoft Purview data security and compliance protections for generative AI apps - https://learn.microsoft.com/purview/ai-microsoft-purview
- DSPM for AI in Microsoft Purview - https://learn.microsoft.com/purview/data-security-posture-management-learn-about
- Microsoft Defender for Cloud AI security and threat protection - https://learn.microsoft.com/azure/defender-for-cloud/defender-for-cloud-introduction#ai-security-and-threat-protection
- Discover AI agents and assess security posture using Microsoft Defender - https://learn.microsoft.com/defender-xdr/security-for-ai/ai-agent-inventory

Topic 3 - Question 7

Multiple choice | Domain: Deploy

topic-3/q07

You are creating validation criteria for a custom generative AI model that produces business reports based on internal enterprise data.

You need to assess whether the model's outputs are appropriate and meaningful for the business reports.

Which metric should you use?

A. the number of active users interacting with the model
B. alignment of the output to domain-specific tasks
C. the average system resource usage during inference
D. the model training duration

Answer

B

Correct answer: B. alignment of the output to domain-specific tasks

Why B is correct

The question asks whether the model's outputs are appropriate and meaningful for business reports based on internal enterprise data. That is a quality-of-output question, and Microsoft's evaluation guidance maps it to AI quality metrics that measure how well the response reflects the intended task and domain. Microsoft Foundry evaluators include task adherence, intent resolution, relevance, groundedness, and similar AI-assisted quality measures designed to test whether the output matches the domain task. Domain-specific alignment is the dimension that determines whether the output is appropriate and meaningful for the use case.

Why the other options are wrong

A. The number of active users interacting with the model is an adoption metric. It says nothing about whether the outputs themselves are appropriate or meaningful.

C. The average system resource usage during inference is an operational metric (compute, latency, cost). It does not measure output quality.

D. The model training duration is a process metric for fine-tuning. It does not measure whether the resulting outputs are appropriate for business reports.

Microsoft Learn references

- Built-in evaluators reference (Microsoft Foundry) - https://learn.microsoft.com/azure/foundry/concepts/built-in-evaluators
- Observability in generative AI (evaluation, monitoring, tracing) - https://learn.microsoft.com/azure/foundry/concepts/observability
- Evaluation metrics built-in (Azure AI Studio) - https://learn.microsoft.com/azure/ai-studio/concepts/evaluation-metrics-built-in

Topic 3 - Question 8

Multiple choice | Domain: Deploy

topic-3/q08

A company has Microsoft Foundry agents that generate responses by using Azure OpenAI resources. The agents are deployed to both the United States and Europe.

A company mandate states that the agents and their grounding data must adhere to data residency and movement regulations.

You need to recommend a governance solution for the agents.

What should you include in the recommendation?

A. Microsoft Defender for Cloud
B. Azure Policy
C. Azure Monitor
D. Microsoft Purview

Answer

B

Correct answer: B. Azure Policy

Why B is correct

The requirement is that agents and their grounding data must adhere to data residency and movement regulations. Azure Policy is the Azure service for enforcing where resources can be deployed. Microsoft Learn calls out the pattern explicitly: "If you have data residency requirements that require you to keep all your data in a single Azure region, you can enforce... by using an Azure Policy. You can also enforce a policy that the [resources] are not geo-replicated to other regions." Policy can restrict allowed locations for Azure OpenAI accounts, Foundry resources, storage, and other services that hold grounding data, and it continuously evaluates resources for compliance.

Why the other options are wrong

A. Microsoft Defender for Cloud assesses security posture and detects threats. It does not enforce regional deployment constraints.

C. Azure Monitor collects logs and metrics for diagnostics and alerts. It does not enforce data residency.

D. Microsoft Purview governs and protects data (DSPM for AI, DLP, sensitivity labels, audit). It is not the tool for enforcing the Azure region in which resources are deployed.

Microsoft Learn references

- Use Azure Policy to enforce residency requirements - https://learn.microsoft.com/azure/cosmos-db/data-residency#use-azure-policy-to-enforce-the-residency-requirements
- Move data across regions for Copilots and generative AI features - https://learn.microsoft.com/power-platform/admin/geographical-availability-copilot
- Azure Policy overview - https://learn.microsoft.com/azure/governance/policy/overview

Topic 3 - Question 9

Multiple choice | Domain: Deploy

topic-3/q09

A company has a Microsoft Copilot Studio agent that uses custom connectors to interact with enterprise APIs.

You need to recommend an application lifecycle management (ALM) process to ensure that the connectors are deployed consistently across development, test, and production environments and meet governance and traceability requirements.

What should you recommend?

A. Deploy the APIs as Azure Functions.
B. Manage the connectors as solution components and deploy the components by using ALM pipelines.
C. Maintain connector definitions in environment variables.
D. Export and import the connectors between the environments as unmanaged solutions.

Answer

B

Correct answer: B. Manage the connectors as solution components and deploy the components by using ALM pipelines.

Why B is correct

Custom connectors are solution-aware components in Power Platform. Microsoft's ALM guidance is to package the connector (along with the agent and any related artifacts) inside a Dataverse solution and promote it through an automated ALM pipeline (Pipelines in Power Platform, Azure DevOps with the Power Platform Build Tools, or GitHub Actions for Power Platform). This approach gives consistent deployment across development, test, and production, and it produces traceable deployment records that satisfy governance requirements.

Why the other options are wrong

A. Deploy the APIs as Azure Functions. Hosting the API differently does not address the ALM problem. The connector and the agent still need a packaging and promotion strategy.

C. Maintain connector definitions in environment variables. Environment variables hold values that change between environments (URLs, IDs, secrets). They do not move connector definitions between environments.

D. Export and import the connectors between the environments as unmanaged solutions. Production should receive managed solutions, not unmanaged ones. Unmanaged in production prevents clean upgrade and rollback paths and undermines traceability.

Microsoft Learn references

- Establish an application lifecycle management strategy (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/alm
- Pipelines in Power Platform - https://learn.microsoft.com/power-platform/alm/pipelines
- Power Platform Build Tools for Azure DevOps - https://learn.microsoft.com/power-platform/alm/devops-build-tools
- GitHub Actions for Power Platform - https://learn.microsoft.com/power-platform/alm/devops-github-actions

Topic 3 - Question 10

Hotspot | Domain: Deploy

topic-3/q10

HOTSPOT -

A company plans to implement an AI solution that will contain a Microsoft Copilot Studio agent and a Microsoft Foundry agent. The solution will be stored in a source code repository.

You need to recommend a deployment method for each agent. The solution must meet the following requirements:

A test environment must be used before a deployment to production.

Production must be isolated from development and testing.

The deployment must be repeatable and fully automated.

The solution must NOT require manual intervention.

Which deployment method should you recommend for each agent? To answer, select the appropriate options in the answer area.

Answer area

- Copilot Studio:
  [ ] Export from the source code repository and import to the target environment.
  [ ] Use a Bicep file.
  [ ] Use a Microsoft Power Platform deployment pipeline.

- Microsoft Foundry:
  [ ] Use a Bicep file.
  [ ] Use a Microsoft Power Platform deployment pipeline.
  [ ] Use an Azure DevOps pipeline.

Answer

Copilot Studio = Microsoft Power Platform deployment pipeline; Microsoft Foundry = Azure DevOps pipeline

Correct answer: - **Copilot Studio: Use a Microsoft Power Platform deployment pipeline.** - **Microsoft Foundry: Use an Azure DevOps pipeline.**

Why these selections are correct

Copilot Studio agents are Power Platform solution components. Microsoft's documented ALM path for Copilot Studio is to package the agent inside a solution and promote it across dev, test, and production with Pipelines in Power Platform. Pipelines are repeatable, automated, and require no manual intervention, and they natively understand Dataverse and Copilot Studio dependencies. Bicep does not deploy Copilot Studio agents.

Microsoft Foundry agents are Azure resources. Microsoft Learn provides Bicep templates and Azure CLI/PowerShell deployment patterns for Foundry. To make the deployment fully repeatable across dev, test, and production with isolation and no manual intervention, the template should run inside an Azure DevOps pipeline (or GitHub Actions equivalent) using the AzureResourceManagerTemplateDeployment task. Pipelines in Power Platform and Bicep on its own do not satisfy "repeatable and fully automated" end-to-end for a Foundry agent stored in source control. An Azure DevOps pipeline does, and it can invoke the Bicep template as part of its job.

Why the other options are wrong

Copilot Studio with Bicep file is not supported. Bicep cannot deploy Copilot Studio agents.

Copilot Studio export from source repo, import to target is a manual operation. It violates the "no manual intervention" rule.

Microsoft Foundry with Power Platform deployment pipeline is not supported. Power Platform pipelines target Dataverse solution components, not Azure resources.

Microsoft Foundry with Bicep file alone describes the artifact, not the automation. A Bicep file by itself does not deploy automatically; you need a pipeline orchestrating the deployment.

Microsoft Learn references

- Pipelines in Power Platform - https://learn.microsoft.com/power-platform/alm/pipelines
- Solution management in Microsoft Copilot Studio - https://learn.microsoft.com/microsoft-copilot-studio/authoring-solutions-overview
- Quickstart: Deploy a Microsoft Foundry resource by using a Bicep file - https://learn.microsoft.com/azure/foundry/how-to/create-resource-template
- Integrate Bicep with Azure Pipelines - https://learn.microsoft.com/azure/azure-resource-manager/bicep/add-template-to-azure-pipelines
- AzureResourceManagerTemplateDeployment@3 task - https://learn.microsoft.com/azure/devops/pipelines/tasks/reference/azure-resource-manager-template-deployment-v3

Topic 3 - Question 11

Hotspot | Domain: Deploy

topic-3/q11

HOTSPOT -

A company has a Microsoft Copilot Studio prompt-and-response agent.

You need to ensure that the agent meets the following requirements:

Provides effective and relevant responses

Provides conversational outcomes

Which metric should you use for each requirement? To answer, select the appropriate options in the answer area.

Answer area

- Provides effective and relevant responses:

[ ] Generated answer rate and quality

[ ] Reactions

[ ] Tool use

- Provides conversational outcomes:

[ ] Satisfaction

[ ] Tool use
[ ] Topics by outcome

Answer

Effective and relevant responses = Generated answer rate and quality; Conversational outcomes = Topics by outcome

Correct answer: - **Provides effective and relevant responses: Generated answer rate and quality** - **Provides conversational outcomes: Topics by outcome**

Why these selections are correct

Generated answer rate and quality measures whether the agent answered the user's question and whether the answer is good. Microsoft Learn states that Copilot Studio "looks at a sample set of answered questions and analyzes different quality, including completeness, relevance, and level of groundedness of a response." That maps directly to "effective and relevant responses." Reactions show user thumbs up/down feedback only, and Tool use measures tool invocations rather than answer quality, so Generated answer rate and quality is the correct match.

Topics by outcome shows the relationship between topics and the outcomes of the conversations they led to (resolved, escalated, abandoned). Microsoft Learn describes the topic analytics pane as showing "a chart showing the aggregated breakdown by outcome for the selected period." Mapping conversational outcomes to the topics that produced them is precisely what this metric is for. Satisfaction tracks CSAT and sentiment but not outcomes, and Tool use measures tool calls.

Microsoft Learn references

- Analyze conversational agent effectiveness (Generated answer rate and quality, Conversation outcomes) - https://learn.microsoft.com/microsoft-copilot-studio/analytics-improve-agent-effectiveness
- Analyze topic usage in Copilot Studio - https://learn.microsoft.com/microsoft-copilot-studio/analytics-topic-details
- Measure agent outcomes - https://learn.microsoft.com/microsoft-copilot-studio/guidance/measuring-outcomes

Topic 3 - Question 12

Multi-select | Domain: Deploy
topic-3/q12

A company extends Copilot in Microsoft Dynamics 365 Customer Service.

You need to recommend an automated application lifecycle management (ALM) process so that the Copilot components can be safely developed, tested, and promoted to production.

Which two actions should you include in the ALM process? Each correct answer presents part of the solution.

A. Use an unmanaged solution in production.
B. Rebuild the agents in each environment.
C. Use Microsoft Power Platform pipelines.
D. Include the components in a solution.
E. Store the agent transcripts in source control.

Select all that apply.

Answer

C, D

Correct answer: C. Use Microsoft Power Platform pipelines and D. Include the components in a solution.

Why C and D are correct

Copilot extensions in Dynamics 365 Customer Service are Copilot Studio components, and Power Platform's prescribed ALM pattern packages them in a solution and promotes them through pipelines.

Microsoft Learn states: "When you create an agent in Copilot Studio, the agent is created within a Power Platform solution. You can create custom solutions to manage your agents across multiple environments, or for pipeline deployments and other application lifecycle management (ALM) scenarios." Pipelines in Power Platform "democratize ALM... by bringing ALM automation and continuous integration and continuous delivery (CI/CD) capabilities into the service."

Together, packaging in a solution (D) and promoting through pipelines (C) provide safe, repeatable development, test, and production movement.

Why the other options are wrong

A. Use an unmanaged solution in production. Microsoft's ALM golden rule is to deploy managed solutions to production. Unmanaged in production prevents clean uninstall and upgrade.

B. Rebuild the agents in each environment. Manual rebuilds break consistency and traceability and bypass ALM entirely.

E. Store the agent transcripts in source control. Transcripts are runtime conversation data, not ALM artifacts. They do not belong in a deployment pipeline.

Microsoft Learn references

- Establish an application lifecycle management strategy (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/alm
- Pipelines in Power Platform - https://learn.microsoft.com/power-platform/alm/pipelines
- Create and manage solutions in Copilot Studio - https://learn.microsoft.com/microsoft-copilot-studio/authoring-solutions-overview
- ALM basics with Microsoft Power Platform - https://learn.microsoft.com/power-platform/alm/basics-alm

Topic 3 - Question 13

Hotspot | Domain: Deploy

topic-3/q13

HOTSPOT -

You are designing a testing solution for a Microsoft Copilot Studio agent that integrates with Microsoft Dynamics 365 Customer Service and Dynamics 365 Sales.

You need to design end-to-end scenarios to test the agent's ability to perform the following actions:

Coordinate tasks and data interactions across both Dynamics 365 apps.

Interpret user input and provide contextually relevant outputs.

Which test scenario and metric should you include in the design? To answer, select the appropriate options in the answer area.

Answer area

- Test scenario:

[ ] In each app, test isolated tasks without using workflows.

[ ] Run task-based scenarios that involve both apps.
[ ] Test visual consistency across both apps.

- Metric:

[ ] Measure the initial prompt response time for each app.
[ ] Track the average click rate across both apps.
[ ] Track the successful completion of cross-app tasks.

Answer

Test scenario = Run task-based scenarios that involve both apps; Metric = Track the successful completion of cross-app tasks

Correct answer: - **Test scenario: Run task-based scenarios that involve both apps.** - **Metric: Track the successful completion of cross-app tasks.**

Why these selections are correct

The agent must coordinate tasks across Dynamics 365 Customer Service and Dynamics 365 Sales and interpret user input to produce contextually relevant outputs. End-to-end testing for that requires task-based scenarios that exercise both applications together. Isolated tests in each app would not validate cross-app coordination, and visual consistency tests do not validate functional behavior.

Microsoft's evaluation guidance for Copilot Studio agents focuses on real task execution. The Copilot Studio Kit and Power Platform pipelines run automated test cases that validate agent behavior end-to-end before promotion. Tracking successful completion of those cross-app tasks is the metric that reflects whether the agent actually performed the intended work. Initial prompt response time measures latency only, and click rate is not a meaningful agent quality metric.

Microsoft Learn references

- Automate testing and deployment of agents using pipelines in Power Platform - https://learn.microsoft.com/microsoft-copilot-studio/guidance/kit-automate-test-deploy
- Copilot Studio Kit overview - https://learn.microsoft.com/microsoft-copilot-studio/guidance/kit-overview
- Improve your Copilot Studio projects (analytics and outcomes) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/improve-overview
- Measure agent outcomes - https://learn.microsoft.com/microsoft-copilot-studio/guidance/measuring-outcomes

Topic 3 - Question 14

Multiple choice | Domain: Deploy

topic-3/q14

A company has multiple AI models that support generation of sales transactions.

Each release of the models must be reviewed by a security and compliance team before being deployed to the production environment. The security and compliance team must have access to prior versions to properly determine potential exposures introduced.

You need to recommend a solution to evaluate the impact of each deployment to production. The solution must enhance business continuity.

What should you recommend?

A. Create a central model registry that uses version history.
B. Establish a promotion process by using a quality gate.
C. Implement version control for all the AI system components.

D. Track model retirement schedules to prevent service disruptions.

Answer

A

Correct answer: A. Create a central model registry that uses version history.

Why A is correct

The requirement combines three things: a security and compliance review of each release, access to prior versions for impact analysis, and business continuity. A central model registry with version history is the canonical pattern for that. It tracks every published model version with metadata, lineage, and lifecycle stage; security and compliance teams can review the new version and compare it against prior versions, and the registry preserves rollback paths if a release introduces regressions or exposures, supporting business continuity.

Microsoft's GenAIOps guidance and Foundry Control Plane both center on registered model assets with versioning, evaluation results, and observability tied to each version. That is the structure that lets reviewers see what changed release over release and enables rapid rollback.

Why the other options are wrong

B. Establish a promotion process by using a quality gate. A quality gate enforces tests before promotion, but it does not by itself give the security team access to prior versions or guarantee business continuity if a deployed release fails.

C. Implement version control for all the AI system components. Source-level version control covers prompts, code, and config, but not the model artifacts and their evaluation history. A model registry is the purpose-built complement.

D. Track model retirement schedules to prevent service disruptions. Retirement tracking helps with vendor model deprecation, not the impact assessment of new releases.

Microsoft Learn references

- Generative AI operations for organizations with MLOps investments - https://learn.microsoft.com/azure/architecture/ai-ml/guide/genaiops-for-mlops
- Microsoft Foundry Control Plane (assets, observability, compliance) - https://learn.microsoft.com/azure/foundry/control-plane/overview
- Manage agents at scale in Microsoft Foundry Control Plane - https://learn.microsoft.com/azure/foundry/control-plane/how-to-manage-agents

Topic 3 - Question 15

Drag and drop | Domain: Deploy

topic-3/q15

DRAG DROP -

A company has an AI solution that uses a Microsoft Copilot Studio agent.

You need to monitor the agent's performance. The solution must meet the following requirements:

Monitor the agent's telemetry in near-real-time (NRT).

Download transcripts of full conversations.

Monitor the agent's usage and performance.

What should you use for each requirement? To answer, drag the appropriate options to the correct requirements. Each option may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

Drag sources

- Application Insights
- Copilot Studio
- Log Analytics
- Microsoft Power Apps

Drop targets

1. Monitor the agent's telemetry in NRT: - drop here
2. Download transcripts of full conversations: - drop here
3. Monitor the agent's usage and performance: - drop here

Answer

NRT telemetry = Application Insights; Download transcripts = Copilot Studio; Usage and performance = Copilot Studio

Correct answer: - **Monitor the agent's telemetry in NRT: Application Insights** - **Download transcripts of full conversations: Copilot Studio** - **Monitor the agent's usage and performance: Copilot Studio**

Why these selections are correct

Telemetry in near-real-time: Microsoft Learn shows that Copilot Studio sends detailed telemetry (conversations, tool usage, exceptions, latency, custom events) to Application Insights, which is designed for near-real-time analysis through Live Metrics, Log Analytics, and the Copilot Studio Dashboard workbook. Application Insights is the integrated NRT signal source.

Download transcripts of full conversations: Conversation transcripts are stored in Dataverse and accessible from Copilot Studio. Microsoft Learn states: "You can download conversation transcripts a few minutes after the conversation times out... You can download them in Dataverse using the Power Apps portal and as session chat transcripts using the Copilot Studio app." Copilot Studio is the surface for downloading transcripts.

Usage and performance: The Analytics page in Copilot Studio aggregates usage metrics (sessions, engagement, resolution, escalation, savings, generated answer quality, themes). It is the maker-facing view for usage and performance.

Log Analytics could host the data behind Application Insights, but the question lists Application Insights as a separate option, and Log Analytics is not the labeled NRT entry point. Microsoft Power Apps is unrelated to agent monitoring.

Microsoft Learn references

- Capture telemetry with Application Insights (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/advanced-bot-framework-composer-capture-telemetry
- Download agent session transcripts - https://learn.microsoft.com/microsoft-copilot-studio/analytics-sessions
- Analytics overview (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/analytics-overview
- Analyze conversational agent effectiveness - https://learn.microsoft.com/microsoft-copilot-studio/analytics-improve-agent-effectiveness

Topic 3 - Question 16

Hotspot | Domain: Deploy

topic-3/q16

HOTSPOT -

A company deploys a Microsoft Copilot Studio agent that integrates with a Microsoft Power Automate desktop flow.

You need to recommend a testing solution that meets the following requirements:

Test cases must validate the most recent changes to the agent before the agent is released.

The flow must be validated as part of the agent's orchestration.

What should you recommend for each requirement? To answer, select the appropriate options in the answer area.

Answer area

- Validate the most recent changes to the agent before release:

[ ] Publish the agent to a channel and test the agent on live users.

[ ] Run tests against the latest unpublished version of the agent.

[ ] Run tests against the production version of the agent.

- Validate the flow as part of the agent's orchestration:

[ ] Add the flow to the agent as a tool.

[ ] Add the flow to a canvas app.

[ ] Use the Power Automate for desktop console.

Answer

Recent changes = Run tests against the latest unpublished version of the agent; Validate the flow = Add the flow to the agent as a tool

Correct answer: - **Validate the most recent changes to the agent before release: Run tests against the latest unpublished version of the agent.** - **Validate the flow as part of the agent's orchestration: Add the flow to the agent as a tool.**

Why these selections are correct

Latest unpublished version: Microsoft's testing guidance for Copilot Studio is to test agents in the maker authoring environment before publishing. The Test panel and the Copilot Studio Kit run tests against the unpublished, in-progress version so changes are validated before they reach users. Publishing the agent to a channel and testing on live users skips validation, and the production version is what you are trying to gate against.

Flow as a tool: Agent flows are added to a Copilot Studio agent as tools so the orchestrator can invoke them as part of conversation logic. Microsoft Learn says that to add a flow to an agent it "must be a solution flow and contain the When an agent calls the flow trigger and the Respond to the agent action," and the agent invokes it through the orchestrator. Validating the flow as part of the agent's orchestration therefore requires adding the flow as a tool. A canvas app is unrelated, and the Power Automate for desktop console is for unattended desktop flow runs, not for orchestrated agent calls.

Microsoft Learn references

- Use agent flows with your agent - https://learn.microsoft.com/microsoft-copilot-studio/advanced-flow
- Call an agent flow - https://learn.microsoft.com/microsoft-copilot-studio/advanced-use-flow
- Agent flows in Microsoft Copilot Studio FAQ -

https://learn.microsoft.com/microsoft-copilot-studio/flows-faqs
- Automate testing and deployment of agents using pipelines in Power Platform - https://learn.microsoft.com/microsoft-copilot-studio/guidance/kit-automate-test-deploy

Topic 3 - Question 17

Multi-select | *Domain: Deploy*

topic-3/q17

A company has a Microsoft Copilot Studio agent that provides answers based on a knowledge base for customer support.

Users report that, occasionally, the agent provides inaccurate answers.

You need to use metrics from the Analytics tab in Copilot Studio to identify the cause of the inaccuracies.

Which two options should you use? Each correct answer presents part of the solution.

A. survey results
B. session information and session outcomes
C. topic usage and topics with low resolution
D. engagement, resolution, and escalation rates
E. quality of generated answers

Select all that apply.

Answer

C, E

Correct answer: C. topic usage and topics with low resolution and E. quality of generated answers

Why C and E are correct

The agent draws on a knowledge base, and users report inaccurate answers. Two Analytics signals point at the cause.

Topic usage and topics with low resolution identifies topics that route a high share of sessions into unresolved or escalated outcomes. Microsoft Learn lists "Topics with low resolution" among the optimization metrics to focus on monthly: "Resolution rate for an individual topic." A low resolution rate on a specific topic typically means the topic's logic, knowledge mapping, or grounding has gaps, which directly correlates with inaccurate answers.

Quality of generated answers measures whether the AI-generated responses are good. Microsoft Learn describes the Generated answer rate and quality section: "Copilot Studio looks at a sample set of answered questions and analyzes different quality, including completeness, relevance, and level of groundedness of a response. If the answer meets a set standard, Copilot Studio labels the answer as Good quality." The "Poor" reasons (irrelevant, incomplete, ungrounded) explain why answers are inaccurate and identify which knowledge sources are responsible.

Why the other options are wrong

A. Survey results capture user satisfaction (CSAT) but do not tell you why an answer was wrong.

B. Session information and session outcomes show whether sessions resolved, escalated, or were abandoned. They do not isolate inaccuracy in a specific knowledge-grounded answer.

D. Engagement, resolution, and escalation rates are top-line KPIs. They show that something is wrong but not which topic or knowledge source produced the inaccurate answer.

Microsoft Learn references

- Analyze conversational agent effectiveness (Generated answer rate and quality) - https://learn.microsoft.com/microsoft-copilot-studio/analytics-improve-agent-effectiveness
- Measure and improve agent performance with KPIs and analytics - https://learn.microsoft.com/microsoft-copilot-studio/guidance/analytics
- Analyze topic usage in Copilot Studio - https://learn.microsoft.com/microsoft-copilot-studio/analytics-topic-details
- Topic escalation analysis - https://learn.microsoft.com/microsoft-copilot-studio/guidance/deflection-topic-escalation-analysis

Topic 3 - Question 18

Multi-select | Domain: Deploy

topic-3/q18

A company uses a fine-tuned Microsoft Foundry model that requires frequent updates as new customer feedback becomes available.

You need to design an application lifecycle management (ALM) process that meets the following requirements:

Data changes must be tracked and versioned.

The model must be retrained consistently by using approved training data.

Which two actions should you include in the design? Each correct answer presents part of the solution.

A. Associate the storage location to the fine-tuning job.
B. Create a content filter.
C. Store the training data in Azure Files.
D. Upload the training data to Microsoft Foundry data files
E. Store the training data in Azure Blob Storage that has version control enabled.

Select all that apply.

Answer

A, E

Correct answer: A. Associate the storage location to the fine-tuning job and E. Store the training data in Azure Blob Storage that has version control enabled.

Why A and E are correct

The two requirements are tracking and versioning training data, and consistent retraining with approved data.

Azure Blob Storage with versioning enabled is the durable, versioned data store Microsoft recommends for fine-tuning datasets. Microsoft Learn states: "For large data files, we recommend that you import from an Azure Blob store. Large files can become unstable when uploaded through multipart forms because the requests are atomic and can't be retried or resumed." Blob Storage's blob versioning automatically maintains previous versions of objects, so you can audit changes and roll back to an approved dataset version.

Associating the storage location to the fine-tuning job ensures every retraining run pulls from the same governed source. The fine-tuning job references the dataset in Blob Storage, the version is recorded with the job, and the team retrains consistently against approved data.

Why the other options are wrong

B. Create a content filter. Content filters control output safety. They do not version training data.

C. Store the training data in Azure Files. Azure Files is an SMB/NFS file share. It is not the recommended pattern for fine-tuning datasets and lacks the same versioning model. Microsoft Learn specifically recommends Blob Storage for large fine-tuning files.

D. Upload the training data to Microsoft Foundry data files. Direct uploads through Foundry are appropriate for small files, but they do not give you versioning or change tracking. The recommendation is Blob Storage with versioning.

Microsoft Learn references

- Customize a model with fine-tuning (upload training data) - https://learn.microsoft.com/azure/ai-foundry/openai/how-to/fine-tuning
- What is Azure Blob Storage? - https://learn.microsoft.com/azure/storage/blobs/storage-blobs-overview
- Blob versioning - https://learn.microsoft.com/azure/storage/blobs/versioning-overview
- Design training data for AI workloads on Azure - https://learn.microsoft.com/azure/well-architected/ai/training-data-design

Topic 3 - Question 19

Hotspot | Domain: Deploy

topic-3/q19

HOTSPOT -

A company deploys agents that generate responses by using Azure OpenAI resources. The agents are deployed to both the United States and Europe.

You need to recommend a governance solution that meets the following requirements:

Enforces the deployment of the resources to only approved Azure regions

Provides continuous compliance verification of the resources

What should you include in the recommendation for each requirement? To answer, select the appropriate options in the answer area.

Answer area

- Enforces the deployment of the resources to only approved regions:

[ ] Azure Monitor
[ ] Azure Policy
[ ] Microsoft Defender for Cloud
[ ] Microsoft Purview
[ ] Microsoft Sentinel

- Provides continuous compliance verification of the resources:

[ ] Azure Monitor
[ ] Azure Policy
[ ] Microsoft Defender for Cloud
[ ] Microsoft Purview
[ ] Microsoft Sentinel

Answer

Enforce regions = Azure Policy; Continuous compliance verification = Microsoft Defender for Cloud

Correct answer: - **Enforces the deployment of the resources to only approved regions: Azure Policy** - **Provides continuous compliance verification of the resources: Microsoft Defender for Cloud**

Why these selections are correct

Azure Policy is the Azure service for enforcing what can be deployed and where. The built-in `allowedLocations` policy and similar definitions block deployments to non-approved regions and audit existing resources. That is exactly the "enforce deployment to only approved regions" requirement.

Microsoft Defender for Cloud provides continuous regulatory compliance assessment through the Compliance dashboard. Defender for Cloud continuously evaluates resources against built-in or custom standards (NIST, ISO 27001, PCI DSS, custom standards), produces compliance scores, and surfaces non-compliant resources for remediation. Microsoft Learn describes this as cloud security posture management with continuous monitoring.

Azure Monitor handles logs and metrics, Microsoft Purview governs data, and Microsoft Sentinel is the SIEM for threat investigation. None of those is the continuous-compliance evaluator that Defender for Cloud is.

Microsoft Learn references

- Use Azure Policy to enforce residency - https://learn.microsoft.com/azure/cosmos-db/data-residency#use-azure-policy-to-enforce-the-residency-requirements
- Azure Policy overview - https://learn.microsoft.com/azure/governance/policy/overview
- Microsoft Defender for Cloud overview (CSPM and compliance) - https://learn.microsoft.com/azure/defender-for-cloud/defender-for-cloud-introduction
- Regulatory compliance in Defender for Cloud - https://learn.microsoft.com/azure/defender-for-cloud/concept-regulatory-compliance

Topic 3 - Question 20

Multiple choice | Domain: Deploy

topic-3/q20

A company has an AI solution that uses Azure OpenAI models.

You need to recommend a governance solution that monitors and audits changes to model configurations and data usage. The solution must minimize administrative effort.

What should you include in the recommendation?

A. Azure Monitor
B. Azure Stream Analytics
C. Azure API Management
D. Azure Policy
E. Microsoft Purview

Answer

E

Correct answer: E. Microsoft Purview

Why E is correct

The requirement is governance that monitors and audits changes to model configurations and data usage with minimal administrative effort. Microsoft Purview is the unified governance and compliance platform Microsoft positions for AI workloads. Microsoft Learn describes Purview's native integration with Azure AI Foundry: "Azure Admins can turn on the setting for any given Azure subscription. There's no action your developers need to take. This setting enables data from all Azure AI based applications running in that subscription to be sent to Microsoft Purview to support governance and compliance outcomes." Purview Audit captures activity logs, DSPM for AI provides visibility into data interactions and policy violations, and the integration is configured at the subscription level, which minimizes administrative effort.

Why the other options are wrong

A. Azure Monitor collects logs and metrics for diagnostics. It can ingest activity but does not provide AI-specific governance, audit, and data usage tracking out of the box.

B. Azure Stream Analytics processes streaming data. It is not a governance tool.

C. Azure API Management front-doors APIs and enforces gateway policies. It does not audit AI model configuration changes.

D. Azure Policy enforces rules on resource configurations and detects non-compliance. It does not capture detailed audit trails for data usage or AI app interactions.

Microsoft Learn references

- Use Microsoft Purview to manage data security and compliance for Microsoft Foundry - https://learn.microsoft.com/purview/ai-azure-foundry
- Microsoft Purview data security and compliance protections for generative AI apps - https://learn.microsoft.com/purview/ai-microsoft-purview
- Audit Copilot Studio activities in Microsoft Purview - https://learn.microsoft.com/microsoft-copilot-studio/admin-logging-copilot-studio
- Enable Data Security for Azure AI with Microsoft Purview - https://learn.microsoft.com/azure/defender-for-cloud/ai-onboarding#enable-data-security-for-azure-ai-with-microsoft-purview

Topic 3 - Question 21

Hotspot | Domain: Deploy

topic-3/q21

HOTSPOT -

A company uses Azure OpenAI models that use grounding data from Microsoft Fabric for agents. The models are fine-tuned by using proprietary datasets.

You need to design a governance solution that meets the following requirements:

Restricts access to the grounding data to only assigned roles

Restricts model fine-tuning to only the AI engineering team

What should you include in the design? To answer, select the appropriate options in the answer area.

Answer area

- Restricts access to the grounding data:

[ ] Azure AI Content Safety
[ ] Azure Monitor alerts
[ ] Azure Policy compliance rules
[ ] Microsoft Purview access policies

- Restricts model fine-tuning:

[ ] Azure Policy compliance rules
[ ] Azure Resource Manager (ARM) resource locks
[ ] Microsoft Entra Conditional Access
[ ] Role-based access control (RBAC) in Microsoft Foundry

Answer

Grounding data = Microsoft Purview access policies; Fine-tuning = Role-based access control (RBAC) in Microsoft Foundry

Correct answer: - **Restricts access to the grounding data: Microsoft Purview access policies** - **Restricts model fine-tuning: Role-based access control (RBAC) in Microsoft Foundry**

Why these selections are correct

Grounding data in Microsoft Fabric is governed by Microsoft Purview. Purview access policies (and the Unified Catalog) control who can read or use the data assets that ground the AI models. Microsoft Learn describes Purview as the platform for "data security solutions" and "unified data governance solutions" and details data access governance for Fabric assets through Purview. Azure Policy enforces resource-level configuration, Azure Monitor alerts surface signals, and Content Safety filters output, none of which is the data access control plane.

Fine-tuning is controlled by RBAC in Microsoft Foundry. Foundry uses Azure RBAC to scope what each principal can do at the Foundry resource and project level. Roles such as Microsoft Foundry Developer, Microsoft Foundry Project Manager, and similar control whether a user can create or run fine-tuning jobs. Microsoft's RBAC for Foundry article enumerates these built-in roles. Conditional Access governs sign-in conditions for users, ARM resource locks prevent resource deletion or modification, and Azure Policy governs resource configuration. None of those scope a user's ability to launch a fine-tuning job the way RBAC does.

Microsoft Learn references

- Microsoft Purview overview - https://learn.microsoft.com/purview/purview
- Data governance with Microsoft Purview (Unified Catalog) - https://learn.microsoft.com/purview/data-governance-overview
- Role-based access control for Microsoft Foundry - https://learn.microsoft.com/azure/foundry/concepts/rbac-foundry
- Microsoft Fabric and Purview integration - https://learn.microsoft.com/fabric/governance/microsoft-purview-fabric

Topic 3 - Question 22

Drag and drop | Domain: Deploy

topic-3/q22

DRAG DROP -

A company has a Microsoft Copilot Studio agent that has been in production for three months.

The agent has received positive feedback from users.

You need to identify the number of questions unanswered by the agent and the number of abandoned sessions between the users and the agent.

Which Copilot Studio insights should you use? To answer, drag the appropriate insights to the correct requirements. Each insight may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

Drag sources

- Conversation outcomes
- Generated answer rate and quality
- Reactions
- Survey results

Drop targets

1. The number of unanswered questions: - drop here
2. The number of abandoned sessions: - drop here

Answer

Unanswered questions = Generated answer rate and quality; Abandoned sessions = Conversation outcomes

Correct answer: - **The number of unanswered questions: Generated answer rate and quality** - **The number of abandoned sessions: Conversation outcomes**

Why these selections are correct

Generated answer rate and quality explicitly tracks unanswered questions. Microsoft Learn states: "The Answer rate shows the number of answered and unanswered questions within the selected time period and the percentage change over time." It also breaks down quality categories for poor answers (irrelevant, incomplete, ungrounded). That is the metric for unanswered questions.

Conversation outcomes tracks the four session outcomes: Resolved, Escalated, Abandoned, and Unengaged. Microsoft Learn defines abandoned as "an engaged session that isn't resolved or escalated and is inactive after 60 minutes from the beginning of the session" and exposes the count and percentage on the Conversation outcomes chart. Reactions captures thumbs up/down, and Survey results captures CSAT, neither of which is the abandoned-session count.

Microsoft Learn references

- Analyze conversational agent effectiveness (Generated answer rate and quality, Conversation outcomes) - https://learn.microsoft.com/microsoft-copilot-studio/analytics-improve-agent-effectiveness
- Measure agent outcomes - https://learn.microsoft.com/microsoft-copilot-studio/guidance/measuring-outcomes
- Deflection overview (key metrics including Abandon Rate) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/deflection-overview

Topic 3 - Question 23

Multiple choice | Domain: Deploy

topic-3/q23

You are evaluating a Microsoft Copilot Studio agent that supports Microsoft Dynamics 365 Customer Service representatives.

You need to recommend a testing solution that meets the following requirements:

Evaluates agent effectiveness during active sessions

Validates whether the agent delivers accurate and helpful responses

Provides measurable, actionable insights for continuous improvement

What should you recommend?

A. Track resolution, deflection, and accuracy by using dashboards and use scripts to ensure consistent responses.
B. Perform load testing to validate how the agent scales under a high chat volume.
C. Review historical tickets to find agents that have the shortest resolution times.
D. Measure uptime and page load times.

Answer

A

Correct answer: A. Track resolution, deflection, and accuracy by using dashboards and use scripts to ensure consistent responses.

Why A is correct

The three requirements (evaluate effectiveness during active sessions, validate accuracy and helpfulness, and produce measurable, actionable insights for continuous improvement) align with the Copilot Studio analytics + scripted-test pattern. The Analytics dashboard tracks resolution rate, escalation rate, and deflection in near-real-time, and the Generated answer rate and quality section measures whether responses are accurate. Scripted tests, run through the Test panel or the Copilot Studio Kit's automated testing pipeline, drive consistency by validating the same prompts produce the expected behavior across releases. Together, dashboards plus scripted regression tests deliver measurable signals and a continuous-improvement loop.

Why the other options are wrong

B. Perform load testing to validate how the agent scales under a high chat volume. Load testing measures throughput and latency, not accuracy or helpfulness.

C. Review historical tickets to find agents that have the shortest resolution times. Historical ticket review evaluates human reps, not the AI agent's effectiveness during active sessions.

D. Measure uptime and page load times. These are infrastructure metrics, not agent quality metrics.

Microsoft Learn references

- Analyze conversational agent effectiveness - https://learn.microsoft.com/microsoft-copilot-studio/analytics-improve-agent-effectiveness
- Deflection overview - https://learn.microsoft.com/microsoft-copilot-studio/guidance/deflection-overview
- Automate testing and deployment of agents using pipelines in Power Platform - https://learn.microsoft.com/microsoft-copilot-studio/guidance/kit-automate-test-deploy
- Improve your Copilot Studio projects - https://learn.microsoft.com/microsoft-copilot-studio/guidance/improve-overview

Topic 3 - Question 24

Multiple choice | Domain: Deploy

topic-3/q24

A company uses multiple Microsoft Copilot Studio agents across different channels.

You need to recommend a monitoring solution that provides comprehensive telemetry data and performance insights for the agents.

What should you include in the recommendation?

A. Application Insights
B. Azure Advisor
C. Azure DevOps

D. Microsoft Dynamics 365 Customer Voice

Answer

A

Correct answer: A. Application Insights

Why A is correct

Application Insights is the Azure service Microsoft positions for comprehensive Copilot Studio agent telemetry. Microsoft Learn explains that Copilot Studio sends conversations, events, exceptions, latency, and tool usage to Application Insights and that the Copilot Studio Dashboard workbook surfaces single-pane operational metrics across these signals. For agents running across multiple channels, Application Insights captures the channel dimension on each event, allowing cross-channel performance comparisons through Azure Workbooks, Log Analytics queries, and dashboards.

Why the other options are wrong

B. Azure Advisor issues best-practice recommendations on Azure resources. It is not a telemetry surface for Copilot Studio agents.

C. Azure DevOps provides source control and pipelines for ALM. It does not collect agent telemetry.

D. Microsoft Dynamics 365 Customer Voice is a survey product. It captures user feedback, not comprehensive agent telemetry.

Microsoft Learn references

- Capture telemetry with Application Insights (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/advanced-bot-framework-composer-capture-telemetry
- Configure telemetry to capture bot data - https://learn.microsoft.com/microsoft-copilot-studio/advanced-bot-framework-composer-capture-telemetry
- Combine web and agent analytics - https://learn.microsoft.com/microsoft-copilot-studio/guidance/pass-web-analytics-tracking-id-from-webpage

Topic 3 - Question 25

Multi-select | Domain: Deploy

topic-3/q25

A company has an AI solution built by using Microsoft Copilot Studio and Power Platform. The solution is used by the company's sales, marketing, and customer service teams.

You are performing a return on AI investment (ROAI) analysis to evaluate the impact of the solution.

You need to identify which measurable business drivers to include in the analysis.

Which two business drivers should you identify? Each correct answer presents part of the solution.

A. the reduced average case resolution time
B. market capitalization
C. economic market predictability
D. increased employee productivity
E. brand awareness

Select all that apply.

Answer

A, D

Correct answer: A. the reduced average case resolution time and D. increased employee productivity

Why A and D are correct

Return on AI investment analysis compares the cost of running the AI solution against measurable business outcomes that the solution influences. Microsoft's Copilot Studio guidance and Viva Insights ROI calculations both center on operational outcomes that the agent can directly affect.

Reduced average case resolution time is a direct, measurable outcome of an AI-assisted customer service flow. The Copilot Studio Savings calculator and Viva Insights agent-assisted impact calculations translate time saved into dollar value, which is exactly what ROAI measures.

Increased employee productivity is the second pillar Microsoft emphasizes. Generative AI in sales, marketing, and customer service offloads routine work and lets people spend time on higher-value activities. Productivity gains, measured in hours saved per employee per week or in higher throughput per team, feed directly into ROI calculations.

Why the other options are wrong

B. Market capitalization reflects investor sentiment about the entire company. The AI solution cannot be isolated as the cause.

C. Economic market predictability is a macroeconomic concept. It is not a measurable business driver attributable to a Copilot Studio solution.

E. Brand awareness depends on marketing reach and reputation, which is not a measurable business driver tied to Copilot Studio adoption.

Microsoft Learn references

- Analyze time and cost savings for agents (Copilot Studio Savings calculator) - https://learn.microsoft.com/microsoft-copilot-studio/analytics-cost-savings
- View aggregated analytics in Viva Insights (ROI calculations) - https://learn.microsoft.com/microsoft-copilot-studio/analytics-viva-insights
- Measure agent value using Copilot Studio Kit - https://learn.microsoft.com/microsoft-copilot-studio/guidance/kit-agent-value-dashboard

Topic 3 - Question 26

Multiple choice | Domain: Deploy

topic-3/q26

A company has a Microsoft Foundry agent that summarizes customer feedback and recommends products to customers. The agent references data from multiple knowledge sources.

Users report that the agent response time is slow.

Telemetry data shows that the agent frequently reaches its token usage limit.

You need to recommend a solution to reduce token usage without degrading the quality of the generated responses.

What should you recommend?

A. Chunk documents during indexing.
B. Reduce the number of knowledge sources used by the agent.
C. Reconfigure the prompts to limit the amount of retrieved content from the knowledge sources.
D. Lower the maximum token usage limit for the responses.

Answer

C

Correct answer: C. Reconfigure the prompts to limit the amount of retrieved content from the knowledge sources.

Why C is correct

Token usage on a RAG-style agent is dominated by the retrieved content stuffed into the prompt context for grounding. When telemetry shows the agent is hitting its token cap, the highest-leverage fix is to reduce the volume of retrieved content per prompt while keeping the same knowledge sources in scope. Microsoft's optimization guidance points to refining prompts and retrieval to reduce ambiguity and tighten the context window. Limiting retrieved chunks per call drops token consumption while preserving response quality, because the agent still consults all the knowledge sources, just less of each per turn.

Why the other options are wrong

A. Chunk documents during indexing. Chunking is good practice for retrieval quality, but it does not by itself reduce the number of chunks pulled into the prompt. Without limiting how many chunks the prompt retrieves, token usage stays high.

B. Reduce the number of knowledge sources used by the agent. Cutting knowledge sources risks degrading response quality, which violates the "without degrading quality" requirement.

D. Lower the maximum token usage limit for the responses. Capping the response length might truncate the answer, which degrades quality. It does not address the input-side token consumption from retrieval.

Microsoft Learn references

- Optimize prompts and instructions (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/optimize-prompts-overview
- Knowledge sources summary (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/knowledge-copilot-studio
- Observability in generative AI (token usage and latency) - https://learn.microsoft.com/azure/foundry/concepts/observability

Topic 3 - Question 27

Hotspot | Domain: Deploy

topic-3/q27

HOTSPOT -

A company has Microsoft Power Platform development, staging, and production environments. Each environment has its own Microsoft Dataverse tables and Azure AI Search index.

You are designing an application lifecycle management (AIM) process to deploy a Microsoft Copilot Studio agent between the environments.

The company has a Copilot Studio agent named Agent1 in development. Agent1 uses the following grounding data sources:

- A Dataverse table named CustomerOrders
- An Azure AI Search index named customer-knowledge

You need to deploy Agent1 to production. The solution must ensure that the agent uses the production grounding data sources, minimizes downtime, and handles credentials and endpoints securely.

What should you include in the deployment package solution, and what should you reconfigure after the deployment? To answer, select the appropriate options in the answer area.

Answer area

- Include in the deployment package solution:

[ ] Agent1 only
[ ] The data sources only
[ ] Agent1 and the data source connections
[ ] Agent1 and references to the data sources
[ ] Agent1, the data sources, and the data source connections

- Reconfigure after the deployment:

[ ] The Dataverse connection only
[ ] The Azure AI Search connection only
[ ] The Dataverse and Azure AI Search connections
[ ] The Agent1 configuration
[ ] The environment variables

Answer

Include = Agent1 and references to the data sources; Reconfigure = The environment variables

Correct answer: - **Include in the deployment package solution: Agent1 and references to the data sources** - **Reconfigure after the deployment: The environment variables**

Why these selections are correct

Agent1 and references to the data sources: Each environment has its own Dataverse tables and its own Azure AI Search index, so the production data sources are not the same artifacts as the development data sources. The deployment package should contain the agent and references to the data sources (typically through environment variables), not the data themselves and not the data source connections, which are environment-specific. Microsoft's ALM golden rules recommend exporting agents and components as managed solutions and using environment variables for settings that change across environments.

Reconfigure the environment variables: After import, the environment variables that point at the Dataverse table and the Azure AI Search endpoint need to be set to the production values. The Dataverse and Azure AI Search connections also need to authenticate, but Microsoft's guidance is to drive that through environment variables and connection references that are bound after import. Reconfiguring environment variables minimizes downtime, handles credentials securely, and keeps the package portable across environments.

Microsoft Learn references

- Establish an application lifecycle management strategy (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/alm
- Export and import agents using solutions - https://learn.microsoft.com/microsoft-copilot-studio/authoring-solutions-import-export
- Environment variables in solutions - https://learn.microsoft.com/power-apps/maker/data-platform/environmentvariables
- Connection references -

https://learn.microsoft.com/power-apps/maker/data-platform/create-connection-reference

Topic 3 - Question 28

Multiple choice | Domain: Deploy

topic-3/q28

A company has a Microsoft Copilot Studio agent that uses generative AI to assist Microsoft Dynamics 365 Customer Service representatives.

The agent currently exhibits a low resolution rate and a high escalation rate.

You need to identify the issue.

What should you use?

A. the Agent dashboard of Dynamics 365 Customer Service historical analytics
B. the Insights tab from the Search & intelligence settings of the Microsoft 365 admin center
C. the Copilot hub in the Power Platform admin center
D. the Analytics tab in Copilot Studio

Answer

D

Correct answer: D. the Analytics tab in Copilot Studio

Why D is correct

Resolution rate and escalation rate are core Copilot Studio analytics metrics, surfaced on the agent's Analytics tab. Microsoft Learn details the metrics: Resolution Rate is "the percentage of engaged sessions that are resolved," Escalation Rate is "the percentage of engaged sessions that are escalated," and the Analytics page exposes both alongside Escalation Rate Drivers, which highlights which topics generate the most escalations. That is the right place to identify the issue behind a low resolution rate and a high escalation rate.

Why the other options are wrong

A. The Agent dashboard of Dynamics 365 Customer Service historical analytics focuses on human representative performance and overall conversation handling, not the Copilot Studio agent's topic-level resolution and escalation drivers.

B. The Insights tab from Search & intelligence settings of Microsoft 365 admin center covers Microsoft 365 Search analytics. It is unrelated to Copilot Studio agent performance.

C. The Copilot hub in the Power Platform admin center provides tenant-level Copilot governance views. It does not surface per-agent topic-level resolution or escalation drivers the way the Analytics tab does.

Microsoft Learn references

- Analyze conversational agent effectiveness - https://learn.microsoft.com/microsoft-copilot-studio/analytics-improve-agent-effectiveness
- Deflection overview (Resolution and Escalation rates) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/deflection-overview
- Topic escalation analysis - https://learn.microsoft.com/microsoft-copilot-studio/guidance/deflection-topic-escalation-analysis
- Measure and improve agent performance with KPIs and analytics - https://learn.microsoft.com/microsoft-copilot-studio/guidance/analytics

Topic 3 - Question 29

Multiple choice | Domain: Deploy

topic-3/q29

A company has a Microsoft Foundry generative AI model.

You need to evaluate the model's output to measure the overall quality and coherence of generated responses. The evaluation must use GPT-4o as a judge and return a numeric score for each output.

Which type of metric should you use?

A. AI quality (NLP)
B. risk and safety
C. Groundedness
D. AI quality (AI assisted)

Answer

D

Correct answer: D. AI quality (AI assisted)

Why D is correct

The question describes overall quality and coherence of generated responses, evaluated with GPT-4o as a judge, returning a numeric score per output. That is the AI-assisted quality category in Microsoft Foundry's evaluator framework. Microsoft Learn lists the AI-assisted quality evaluators as Coherence, Fluency, Relevance, Groundedness, Similarity, Retrieval, Response Completeness, and similar measures, all of which use an LLM judge (GPT-4.1, GPT-4o, or an o-series reasoning model) to score outputs on a numeric scale. Coherence specifically "measures the logical and orderly presentation of ideas in a response."

Why the other options are wrong

A. AI quality (NLP) uses traditional NLP overlap metrics such as F1, BLEU, GLEU, ROUGE, METEOR. They compare model output to a ground-truth reference using token overlap and do not require an LLM judge.

B. Risk and safety evaluates harmful content (violence, hate, sexual, self-harm, jailbreak susceptibility, protected materials). It is not a measure of overall quality and coherence.

C. Groundedness is one specific AI-assisted metric (whether the response is supported by the source data). It does not by itself measure overall quality and coherence; it is a component of the broader AI-assisted quality category.

Microsoft Learn references

- Built-in evaluators reference (Microsoft Foundry) - https://learn.microsoft.com/azure/foundry/concepts/built-in-evaluators
- General purpose evaluators (Coherence and Fluency) - https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/general-purpose-evaluators
- Evaluate your AI agents (model support for AI-assisted evaluators) - https://learn.microsoft.com/azure/foundry-classic/how-to/develop/agent-evaluate-sdk
- Observability in generative AI - https://learn.microsoft.com/azure/foundry/concepts/observability

Topic 3 - Question 30

Hotspot | Domain: Deploy

topic-3/q30

HOTSPOT -

You use Microsoft Copilot Studio analytics to analyze the performance of a deployed Copilot Studio agent.

You need to identify which performance metrics to use to measure the following:

- The percentage of engaged sessions that are escalated to a live customer service representative
- The number of agent queries that cause a knowledge source error

What should you identify for each requirement? To answer, select the appropriate options in the answer area.

Answer area

- The percentage of engaged sessions that are escalated to a live customer service representative:

[ ] Answer quality
[ ] Customer Satisfaction (CSAT) score
[ ] Engagement rate
[ ] Escalation rate

- The number of agent queries that cause a knowledge source error:

[ ] Answer quality
[ ] Engagement rate
[ ] Escalation rate

Answer

Escalated sessions = Escalation rate; Knowledge source error = Answer quality

Correct answer: - **Percentage of engaged sessions escalated to a live representative: Escalation rate** - **Number of agent queries that cause a knowledge source error: Answer quality**

Why these selections are correct

Escalation rate is defined in Microsoft Learn as "the percentage of engaged sessions that are escalated. An escalated session is an engaged session escalated to a human representative." That is the literal metric for the first requirement.

Answer quality (part of the Generated answer rate and quality section) is where Copilot Studio surfaces knowledge-source errors. Microsoft Learn states the side panel shows "Errors... the percentage of queries that resulted in a knowledge-related error for each knowledge source type (for example, SharePoint)." The Answer quality breakdown is the metric that exposes knowledge source errors. CSAT score measures customer satisfaction, and engagement rate measures whether sessions engaged a topic, neither of which is the knowledge-error count.

Microsoft Learn references

- Deflection overview (Escalation Rate definition) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/deflection-overview
- Analyze conversational agent effectiveness (Generated answer rate and quality, knowledge source errors) - https://learn.microsoft.com/microsoft-copilot-studio/analytics-improve-agent-effectiveness
- Analytics overview (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/analytics-overview

Topic 3 - Question 31

Multiple choice | Domain: Deploy

topic-3/q31

A company plans to deploy a Microsoft Foundry agent.

You need to recommend an application lifecycle management (ALM) process to ensure that the agent evaluates against baseline accuracy metrics before being deployed.

What should you recommend?

A. Configure GitHub Actions for new agent versions.
B. Deploy each new agent version directly to production.
C. Use Observability in Foundry Control Plane with evaluation and drift monitoring.
D. Enable Application Insights and use Azure Monitor.

Answer

C

Correct answer: C. Use Observability in Foundry Control Plane with evaluation and drift monitoring.

Why C is correct

The question requires baseline accuracy evaluation before each agent version is deployed. Microsoft Foundry Control Plane brings observability, evaluation, and drift monitoring into a single ALM-aligned surface for Foundry agents. Microsoft Learn describes the Control Plane: "Continuously evaluate agent performance, quality, and risk dimensions. Risk dimensions might include task adherence, intent resolution, tool call success, groundedness, sensitive data leakage, and exposure to jailbreak and cross-domain prompt injection attacks." The Agent Monitoring Dashboard supports continuous evaluation and integrates results into the deployment flow, which is the prescribed path for evaluating agents against baseline metrics before promotion.

Why the other options are wrong

A. Configure GitHub Actions for new agent versions. GitHub Actions can orchestrate deployments, but it does not by itself perform Foundry-specific evaluation against baseline accuracy metrics.

B. Deploy each new agent version directly to production. This skips the evaluation gate entirely.

D. Enable Application Insights and use Azure Monitor. Application Insights captures runtime telemetry. It does not run pre-deployment evaluations against baseline accuracy metrics the way Foundry observability and evaluation do.

Microsoft Learn references

- What is Microsoft Foundry Control Plane? - https://learn.microsoft.com/azure/foundry/control-plane/overview
- Monitor agents with the Agent Monitoring Dashboard - https://learn.microsoft.com/azure/foundry/observability/how-to/how-to-monitor-agents-dashboard
- Observability in generative AI (evaluation, monitoring, tracing) - https://learn.microsoft.com/azure/foundry/concepts/observability
- Built-in evaluators reference - https://learn.microsoft.com/azure/foundry/concepts/built-in-evaluators

Topic 3 - Question 32

Multiple choice | Domain: Deploy

topic-3/q32

You need to recommend a security solution for agents in a Microsoft Power Platform environment.

The agents must use only approved connectors and services. The solution must prevent the agents from accessing sensitive data.

What should you recommend?

A. Configure Azure Monitor to capture connector activity logs.
B. Enable a Microsoft Dataverse audit.
C. Deploy data loss prevention (DLP) policies in Power Platform.
D. Enable customer-managed keys in Microsoft Dataverse.

Answer

C

Correct answer: C. Deploy data loss prevention (DLP) policies in Power Platform.

Why C is correct

DLP policies in Power Platform classify connectors as Business, Non-business, or Blocked, and prevent agents and apps from combining connectors across categories. Microsoft Learn states: "Data policies, including data loss prevention (DLP) policies, act as guardrails to help prevent users from unintentionally exposing organizational data and to protect information security in the tenant. Data policies enforce rules for which connectors are enabled for each environment, and which connectors can be used together." That maps directly to "agents must use only approved connectors" and "prevent the agents from accessing sensitive data."

Why the other options are wrong

A. Configure Azure Monitor to capture connector activity logs. Logging captures evidence after the fact. It does not enforce which connectors agents can use.

B. Enable a Microsoft Dataverse audit. Auditing tracks data access in Dataverse but does not control which connectors an agent can call or block sensitive data flow.

D. Enable customer-managed keys in Microsoft Dataverse. CMK encrypts data at rest with a customer-managed key. It does not constrain connector usage.

Microsoft Learn references

- Data policies (Power Platform) - https://learn.microsoft.com/power-platform/admin/wp-data-loss-prevention
- Implement a data policy strategy - https://learn.microsoft.com/power-platform/guidance/adoption/dlp-strategy
- Power Platform DLP policies on Day 1 - https://learn.microsoft.com/microsoft-365/community/power-platform-dlp-policies-you-should-be-considering-on-day-1
- Advanced connector policies (preview) - https://learn.microsoft.com/power-platform/admin/advanced-connector-policies

Topic 3 - Question 33

Multi-select | Domain: Plan

topic-3/q33

A company has an AI business solution that uses Microsoft Copilot Studio agents.

You need to recommend prompt best practices to improve the effectiveness of agent interactions.

Which two actions should you include in the recommendation? Each correct answer presents part of the solution.

A. Track the duration of the average user session.
B. Analyze the prompt length distribution.
C. Regularly test and refine the prompts based on user input.
D. Use clear and specific instructions in the prompts.
E. Measure system resource usage during prompt processing.

Select all that apply.

Answer

C, D

Correct answer: C. Regularly test and refine the prompts based on user input and D. Use clear and specific instructions in the prompts.

Why C and D are correct

Microsoft's prompt and instruction guidance for Copilot Studio centers on two practices: write clear, specific instructions, and iterate based on real interactions. Microsoft Learn's "Optimize prompts and instructions" guide says to "refine prompts, tool descriptions, and instructions to improve clarity, reduce ambiguity, and guide the orchestrator toward more accurate and consistent responses." That covers both clarity (D) and the iterative refinement loop (C). Real user input and feedback signals (reactions, low-quality answers, themes) drive continuous prompt improvement.

Why the other options are wrong

A. Track the duration of the average user session. Session duration is an engagement metric. It does not improve prompt effectiveness directly.

B. Analyze the prompt length distribution. Length analysis is descriptive. It does not produce a better prompt without the clarity and iteration practices.

E. Measure system resource usage during prompt processing. Resource usage is an operational concern, not a prompt-quality practice.

Microsoft Learn references

- Optimize prompts and instructions (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/optimize-prompts-overview
- Improve your Copilot Studio projects (iterative improvement loop) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/improve-overview
- Review the improve checklist - https://learn.microsoft.com/microsoft-copilot-studio/guidance/improve-checklist

Topic 3 - Question 34

Drag and drop | Domain: Deploy

topic-3/q34

DRAG DROP -

A company has a cloud-based AI solution that uses Azure OpenAI models.

You need to design a monitoring solution that meets the following requirements:

- Monitors performance metrics and operational health for the models
- Monitors AI apps and agents for compliance
- Uses Azure-native capabilities
- Minimizes development effort

What should you use for each requirement? To answer, drag the appropriate options to the correct requirements. Each option may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

Drag sources

- Azure API Management
- Azure Policy
- Microsoft Defender
- Azure Monitor
- Azure Stream Analytics
- Microsoft Purview

Drop targets

1. Monitors AI apps and agents for compliance: - drop here
2. Monitors performance metrics and operational health: - drop here

Answer

Compliance = Microsoft Purview; Performance/operational health = Azure Monitor

Correct answer: - **Monitors AI apps and agents for compliance: Microsoft Purview** - **Monitors performance metrics and operational health: Azure Monitor**

Why these selections are correct

Microsoft Purview is Microsoft's compliance and governance platform for AI. Microsoft Learn describes Purview's native integration with Microsoft Foundry and Azure OpenAI: subscription-level toggling sends prompts and responses to Purview for audit, communication compliance, DLP, DSPM for AI, and insider risk monitoring. That matches "monitor AI apps and agents for compliance" and uses Azure-native capabilities with minimal development effort.

Azure Monitor collects metrics, logs, and traces from Azure OpenAI and Foundry resources. It is the Azure-native observability platform for performance and operational health: latency, throughput, request volume, errors, dependencies. Application Insights (a feature of Azure Monitor) feeds the Foundry Agent Monitoring Dashboard.

Defender protects against threats and discovers AI workloads but is not the operational performance monitor. Azure Policy enforces configuration, API Management front-doors APIs, and Stream Analytics processes streaming data, none of which fits either requirement.

Microsoft Learn references

- Use Microsoft Purview to manage data security and compliance for Microsoft Foundry - https://learn.microsoft.com/purview/ai-azure-foundry
- Microsoft Purview data security and compliance protections for generative AI apps - https://learn.microsoft.com/purview/ai-microsoft-purview
- Monitor agents with the Agent Monitoring Dashboard (Azure Monitor / App Insights) - https://learn.microsoft.com/azure/foundry/observability/how-to/how-to-monitor-agents-dashboard
- Azure Monitor overview - https://learn.microsoft.com/azure/azure-monitor/overview

Topic 3 - Question 35

Multi-select | Domain: Deploy

topic-3/q35

A company has Microsoft Copilot Studio agents.

The company plans to deploy custom connectors across development, test and production environments.

You need to design an application lifecycle management (ALM) process to ensure consistency and prevent direct editing in production.

Which two actions should you include in the design? Each correct answer presents part of the solution.

A. Deploy managed solutions to production.
B. Deploy unmanaged solutions to production.
C. Manually rebuild the agents in each environment.
D. Move the agents between the environments by using data export and import.
E. Include agents and connectors in a solution.

Select all that apply.

Answer

A, E

Correct answer: A. Deploy managed solutions to production and E. Include agents and connectors in a solution.

Why A and E are correct

Power Platform's ALM golden rules call for packaging components in a solution and deploying as managed in non-development environments. Microsoft Learn states: "Export and deploy solutions as managed, unless setting up a development environment." Managed solutions in production prevent direct edits, enforce consistency across environments, and provide clean uninstall and upgrade behavior. Including the agents and custom connectors in a single solution treats them as one ALM unit so the connector and agent versions stay in lockstep.

Why the other options are wrong

B. Deploy unmanaged solutions to production. Unmanaged solutions allow direct edits in production, which is exactly what the requirement forbids.

C. Manually rebuild the agents in each environment. Manual rebuilds break consistency and traceability and make connector versions drift across environments.

D. Move the agents between the environments by using data export and import. Data export and import bypasses solution-aware ALM, prevents dependency tracking, and undermines version control.

Microsoft Learn references

- Establish an application lifecycle management strategy (Copilot Studio, ALM golden rules) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/alm
- Organize your solutions - https://learn.microsoft.com/power-platform/alm/organize-solutions
- ALM basics with Microsoft Power Platform - https://learn.microsoft.com/power-platform/alm/basics-alm
- Pipelines in Power Platform - https://learn.microsoft.com/power-platform/alm/pipelines

Topic 3 - Question 36

Multiple choice | Domain: Deploy

topic-3/q36

A company uses multiple Microsoft Copilot Studio agents across different channels.

You need to recommend a monitoring solution that provides comprehensive telemetry data and performance insights for the agents.

What should you include in the recommendation?

A. Application Insights
B. Microsoft Dynamics 365 Customer Voice
C. Log Analytics
D. Microsoft Purview

Answer

A

Correct answer: A. Application Insights

Why A is correct

Application Insights is the integrated telemetry and monitoring service Microsoft positions for Copilot Studio agents. Microsoft Learn shows that Copilot Studio sends conversations, exceptions, latency, tool usage, and custom events into Application Insights. The Copilot Studio Dashboard workbook in Application Insights surfaces total conversations, latency, exceptions, tool usage, and topic analytics in a single view, which provides comprehensive telemetry data and performance insights across multiple channels.

Why the other options are wrong

B. Microsoft Dynamics 365 Customer Voice is a survey product. It captures user feedback, not comprehensive agent telemetry.

C. Log Analytics is the query layer that sits behind Application Insights. The Application Insights workspace is the entry point Microsoft documents for Copilot Studio telemetry. Log Analytics alone is not the labeled telemetry destination for Copilot Studio agents.

D. Microsoft Purview governs and protects data and audits activity for compliance. It is not the comprehensive performance telemetry surface for Copilot Studio agents.

Microsoft Learn references

- Capture telemetry with Application Insights (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/advanced-bot-framework-composer-capture-telemetry
- Configure telemetry to capture bot data - https://learn.microsoft.com/microsoft-copilot-studio/advanced-bot-framework-composer-capture-telemetry
- Combine web and agent analytics -

https://learn.microsoft.com/microsoft-copilot-studio/guidance/pass-web-analytics-tracking-id-from-web page

Topic 3 - Question 37

Multiple choice | Domain: Deploy

topic-3/q37

A company has an AI solution named Solution1 that is deployed to the production environment. Solution1 uses an Azure OpenAI model to generate marketing emails for existing customers.

During an internal review, you identify that Solution1 creates different emails depending on the customers' traits.

You need to recommend a strategy to mitigate the bias. The strategy must adhere to Microsoft responsible AI principles.

What should you recommend?

A. Modify Solution1 to randomly generate emails for different traits.
B. Modify the system instructions of Solution1.
C. Retrain the model by using a larger dataset.
D. Modify the contents of the training dataset.

Answer

D

Correct answer: D. Modify the contents of the training dataset.

Why D is correct

The bias is producing different emails based on customer traits, which is a fairness problem rooted in the data the model learned from. Microsoft's Responsible AI principles include fairness, and the practical mitigation guidance for biased outputs starts with the dataset: identify under-represented or skewed groups in the training data, then rebalance, augment, or curate the dataset to reduce bias before retraining. Microsoft's Responsible AI dashboard, Fairlearn, and the Responsible AI scorecard all support this data-centered fairness assessment and mitigation flow.

Why the other options are wrong

A. Modify Solution1 to randomly generate emails for different traits. Randomization masks the issue but does not address the underlying bias and could degrade quality for everyone.

B. Modify the system instructions of Solution1. Prompt instructions can influence behavior, but they do not fix bias baked into the underlying model. The model still reflects the training data distribution.

C. Retrain the model by using a larger dataset. Larger does not equal less biased. Adding more skewed data can amplify existing bias. The correct fairness intervention is to modify the dataset's contents to address the imbalance, not just its size.

Microsoft Learn references

- Microsoft Responsible AI standard and principles - https://www.microsoft.com/ai/responsible-ai
- Assess AI systems by using the Responsible AI dashboard - https://learn.microsoft.com/azure/machine-learning/concept-responsible-ai-dashboard
- What is Responsible AI? - https://learn.microsoft.com/azure/machine-learning/concept-responsible-ai
- Fairness assessment in Responsible AI - https://learn.microsoft.com/azure/machine-learning/concept-fairness-ml

Topic 3 - Question 38

Multiple choice | Domain: Design

topic-3/q38

A company has a canvas app named App1 in a Microsoft Power Platform environment named Env1.

Env1 uses a customer-managed key for data encryption. App1 connects to multiple data sources to retrieve and update customer and order information.

You need to recommend a solution to add Microsoft Copilot components to App1. The solution must NOT modify the current security or encryption configurations of Env1.

What should you include in the recommendation?

A. Modify the data sources of App1 to make them compatible with Copilot.
B. Duplicate App1 and republish the app in Env1.
C. Enable Copilot features for Env1.
D. Move App1 to a new environment that uses Microsoft-managed keys.

Answer

B

Correct answer: B. Duplicate App1 and republish the app in Env1.

Why B is correct

Microsoft Learn states that environments using customer-managed keys (CMK) and Customer Lockbox have specific Copilot behavior: "Questions and answers about enterprise data environments that have a customer-managed key or Customer Lockbox won't use your data stored in Dataverse." To add Copilot components to a canvas app in such an environment, you typically duplicate the app so the new copy picks up the latest Copilot capabilities and is republished with the components enabled, without changing the environment-level CMK configuration.

The specific issue is that some Copilot capabilities only attach to apps that are saved or republished with the feature available. Duplicating App1 and republishing keeps the existing security and encryption configuration of Env1 intact while letting the new copy include the Copilot components.

Why the other options are wrong

A. Modify the data sources of App1 to make them compatible with Copilot. The data sources are not the problem. CMK encryption applies at the Dataverse level and is governed by the environment configuration, not the data sources.

C. Enable Copilot features for Env1. Toggling tenant or environment-level Copilot features may not give app-level Copilot components without republishing the app. It also risks broader configuration changes than the question allows.

D. Move App1 to a new environment that uses Microsoft-managed keys. This explicitly modifies the security and encryption configuration the question forbids changing.

Microsoft Learn references

- Configure customer-managed encryption keys (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/admin-customer-managed-keys
- Manage your customer-managed encryption key (Power Platform) - https://learn.microsoft.com/power-platform/admin/customer-managed-key
- Add Copilot for app users in model-driven apps - https://learn.microsoft.com/power-apps/maker/model-driven-apps/add-ai-copilot
- Add Copilot to a canvas app -

https://learn.microsoft.com/power-apps/maker/canvas-apps/add-ai-copilot

Topic 3 - Question 39

Multi-select | Domain: Deploy | Case study: Contoso

topic-3/q39

Which two components for the custom AI agent should you include in the application lifecycle management (AIM) process? Each correct answer presents part of the solution.

A. an Azure package

B. a ZIP package

C. a Microsoft Power Platform solution

D. a Cloud Scale Unit (CSU) package

E. an X++model

Select all that apply.

Answer

B, C

Correct answer: B. a ZIP package and C. a Microsoft Power Platform solution

Why B and C are correct

The Contoso custom AI agent is a low-code Copilot Studio agent built on Dynamics 365 Supply Chain Management data. Microsoft Learn defines Power Platform solutions as the ALM container for Copilot Studio agents and connectors: "When you create an agent in Copilot Studio, the agent is created within a Power Platform solution." Solutions are exported as compressed ZIP packages for transport between environments. That is the supported ALM artifact for the Contoso agent.

In Microsoft Learn: "When you export a solution from Dataverse, it's exported as a single compressed file." That compressed file is the ZIP package, and the underlying construct is the Power Platform solution. Both items together are the components included in the ALM process.

Why the other options are wrong

A. an Azure package is not a Power Platform ALM artifact. Azure deployments use ARM, Bicep, or container artifacts, none of which apply to a Copilot Studio agent.

D. a Cloud Scale Unit (CSU) package belongs to Dynamics 365 Finance and Operations apps, not Copilot Studio agents.

E. an X++ model is a Finance and Operations developer artifact. It is not used to package Copilot Studio agents.

Microsoft Learn references

- Solution management in Microsoft Copilot Studio - https://learn.microsoft.com/microsoft-copilot-studio/authoring-solutions-overview
- Establish an application lifecycle management strategy (Copilot Studio) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/alm
- Modernize applications with Power Platform (solutions exported as compressed file) - https://learn.microsoft.com/power-platform/guidance/adoption/application-modernization
- Export and import agents using solutions - https://learn.microsoft.com/microsoft-copilot-studio/authoring-solutions-import-export

Topic 3 - Question 40

Drag and drop | Domain: Deploy | Case study: Contoso
topic-3/q40

DRAG DROP -

Which tools should you recommend to assist the CISO and the CIO with their specific responsibilities? To answer, drag the appropriate tools to the correct executives. Each tool may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

Drag sources

- Azure Blob Storage
- Azure Resource Graph Explorer
- Copilot Studio
- Microsoft Purview

Drop targets

1. CISO: - drop here
2. CIO: - drop here

Answer

CISO = Azure Resource Graph Explorer; CIO = Microsoft Purview

Correct answer: - **CISO: Azure Resource Graph Explorer** - **CIO: Microsoft Purview**

Why these selections are correct

CISO is responsible for "discover and inventory AI resources for auditing." Azure Resource Graph Explorer queries every Azure resource in the tenant across subscriptions, including Foundry projects, Azure OpenAI accounts, AI Search, Storage, and related grounding resources. Microsoft Learn positions Resource Graph as the inventory and exploration tool for resources at scale. Microsoft Purview, Copilot Studio, and Azure Blob Storage do not surface a tenant-wide inventory of AI resources the way Resource Graph does, which makes it the right tool for the CISO's audit mandate.

CIO is responsible for ensuring "appropriate security labels are assigned to the data used by the AI agents." Microsoft Purview is the platform that defines, applies, and audits sensitivity labels across Microsoft 365 and Azure data sources, including data that grounds AI agents. Microsoft Learn details Purview's information protection, sensitivity labels, and DSPM for AI as the controls for labeling and governing AI-relevant data.

Microsoft Learn references

- What is Azure Resource Graph? - https://learn.microsoft.com/azure/governance/resource-graph/overview
- Microsoft Purview overview - https://learn.microsoft.com/purview/purview
- Sensitivity labels in Microsoft Purview - https://learn.microsoft.com/purview/sensitivity-labels
- Microsoft Purview data security and compliance protections for generative AI apps - https://learn.microsoft.com/purview/ai-microsoft-purview

Topic 3 - Question 41

Multiple choice | Domain: Deploy | Case study: Contoso

topic-3/q41

What should you recommend to assist the CEO with their specific responsibilities?

A. the Microsoft Service Trust Portal
B. Microsoft Foundry Tools
C. Microsoft Purview
D. the Responsible AI dashboard
E. Compliance Center

Answer

D

Correct answer: D. the Responsible AI dashboard

Why D is correct

The CEO's responsibility in the Contoso case study is to ensure all solutions adhere to industry-standard responsible AI practices, with quarterly assessment that "must verify reliability, interpretability, fairness, and compliance." Microsoft Learn describes the Responsible AI dashboard as a single interface that brings together model performance and fairness assessment, data exploration, model interpretability, error analysis, counterfactual analysis, and causal inference. The dashboard explicitly supports "model reliability, interpretability, fairness, and compliance," which matches every term in the requirement, and the Responsible AI scorecard packages these results into a shareable PDF for stakeholders and audit reviews.

Why the other options are wrong

A. The Microsoft Service Trust Portal publishes Microsoft's audit reports, certifications, and compliance documents about Microsoft cloud services. It does not assess Contoso's own AI solutions for reliability, interpretability, and fairness.

B. Microsoft Foundry Tools is a generic label and is not the named tool that produces a quarterly responsible AI assessment for the CEO.

C. Microsoft Purview governs data and audits AI interactions for compliance. It does not provide model interpretability and fairness assessments the way the Responsible AI dashboard does.

E. Compliance Center (the legacy Microsoft 365 Compliance Center) is a tenant compliance management portal. It does not perform AI model assessments for reliability and fairness.

Microsoft Learn references

- Assess AI systems by using the Responsible AI dashboard - https://learn.microsoft.com/azure/machine-learning/concept-responsible-ai-dashboard
- What is Responsible AI? - https://learn.microsoft.com/azure/machine-learning/concept-responsible-ai
- Generate Responsible AI insights in the studio UI - https://learn.microsoft.com/azure/machine-learning/how-to-responsible-ai-insights-ui
- Responsible AI scorecard - https://learn.microsoft.com/azure/machine-learning/concept-responsible-ai-scorecard

Topic 3 - Question 42

Drag and drop | Domain: Deploy | Case study: Contoso

topic-3/q42

DRAG DROP -

Which Copilot Studio analytics metrics should you recommend to assist the executives with their specific responsibilities? To answer, drag the appropriate metrics to the correct executives. Each metric may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

Drag sources

- Effectiveness
- Tool use
- Satisfaction
- Use

Drop targets

1. CFO concerns about Copilot Studio credit usage: - drop here
2. CTO concerns about poor feedback on AI agent responses: - drop here

Answer

CFO = Use; CTO = Satisfaction

Correct answer: - **CFO concerns about Copilot Studio credit usage: Use** - **CTO concerns about poor feedback on AI agent responses: Satisfaction**

Why these selections are correct

CFO and the Use metric: The case study states the CFO "wants to identify how many interactions with the AI agents are abandoned on a given day as compared to resolved conversations. Too many abandoned sessions might indicate that Copilot Studio credits are being used inefficiently by end users." The Copilot Studio Analytics page exposes a Use section that surfaces session and run volume and outcomes, and the Conversation outcomes view inside Use shows resolved versus abandoned counts. That is the metric category for tracking credit usage efficiency tied to session outcomes.

CTO and the Satisfaction metric: The case study says the CTO "wants to track user feedback on the quality of the AI agent responses during user interactions with the agents. Consistently poor feedback will trigger an escalated reengineering discussion." The Analytics page's Satisfaction (Effectiveness) section captures CSAT, sentiment, reactions (thumbs up/down), and user comments. That is the surface that reflects user feedback on response quality.

Effectiveness in Analytics aggregates user reactions and survey results, so it overlaps with Satisfaction, but the named Satisfaction metric is the cleanest match for "track user feedback on the quality of agent responses." Tool use measures tool invocations and is unrelated to either question.

Microsoft Learn references

- Analytics overview (Use, Satisfaction, Effectiveness sections) - https://learn.microsoft.com/microsoft-copilot-studio/analytics-overview
- Analyze conversational agent effectiveness (Conversation outcomes, Effectiveness, Satisfaction) - https://learn.microsoft.com/microsoft-copilot-studio/analytics-improve-agent-effectiveness
- Deflection overview (Resolution, Escalation, and Abandon rates) - https://learn.microsoft.com/microsoft-copilot-studio/guidance/deflection-overview
- Analyze time and cost savings for agents - https://learn.microsoft.com/microsoft-copilot-studio/analytics-cost-savings